

Ramsey Alignment and the Postwar Tax System

Jackson Mejia*

August 11, 2026

How far is the tax system from optimal, and how much does it matter? I develop a sufficient statistic that answers both questions locally: the cosine of the angle between the welfare and revenue gradients. I show that this sufficient statistic applies across a wide range of economies, including those featuring dynamics, cross-base spillovers, and nonlinear tax instruments. In an application to federal tax policy in the United States, I estimate a general equilibrium revenue gradient for a planner with access to a personal income tax, a corporate tax, a tariff, and a nonlinear income tax instrument. I find that the tax system became substantially more aligned with the local Ramsey optimum over time across these margins, a result largely driven by the long-run decline in the corporate tax rate.

JEL-Classification: H21, H25, E62, C32

Keywords: Optimal taxation, tax reform, sufficient statistics, marginal value of public funds, corporate taxation, narrative identification, fiscal policy

*Massachusetts Institute of Technology, jpmejia@mit.edu. First Draft: January, 2026. I am grateful to Nathan Hendren, Andrew Johnston, Yahe Li, Pedro Martinez-Bruera, Jim Poterba, and Christian Wolf for useful comments and suggestions, as well as to Tamar den Besten and Diego Känzig for providing narrative tariff shocks.

1 Introduction

The case for tax reform rests on how far the system is from whatever policymakers perceive as optimal. And yet, without a structural model, it is difficult to assess the scope for reform or the direction reform should take. However, structural models also tend to disagree on both of those questions, which creates a difficulty for both policymakers and economists in deciding what ought to be reformed and what the gains would be from reform. In this paper, I develop a set of sufficient statistics that answers both questions locally with minimal structure.

The sufficient statistics emerge from a simple insight: at a local Ramsey optimum, no revenue-neutral tax reform can improve welfare. From the Ramsey planner's first-order condition, the marginal welfare cost per dollar of revenue raised must be equal across tax instruments. Geometrically, that condition holds exactly when the revenue and welfare gradients are parallel. The angle θ between them measures how much the ratio of marginal welfare cost to marginal revenue varies across instruments, so θ is zero at the optimum and positive otherwise. While the classic tax reform literature used this geometry to identify the optimal direction of reform (Dixit 1979; Tirole and Guesnerie 1981), I show how to use the angle θ to form an aggregate diagnostic of the tax system's proximity to the Ramsey optimum and the remaining welfare gains from revenue-neutral tax reform. When $\cos \theta = 1$, the welfare and revenue gradients point in the same direction and the system is perfectly aligned. When $\cos \theta = 0$, the two gradients are orthogonal, so the unconstrained welfare-improving direction is already revenue-neutral. A simple transformation then shows that, for a reform of a given size, $\sin \theta$ measures the share of the largest local welfare gain that remains available under revenue neutrality.

The sufficient statistics are closely related to the literature on the marginal value of public funds (MVPF), which is also exactly what makes them measurable (Mayshar 1990; Bastani 2025). For any particular tax change, the MVPF is the marginal welfare cost relative to the marginal revenue it raises, which just means that equality across MVPFs is another representation of the planner's first-order condition. The cosine provides an aggregate diagnostic of whether that equality holds across the included tax changes, while remaining well-defined when a marginal revenue effect is zero or negative. The same local logic makes the required objects measurable. As Hendren (2016) and Hendren and Sprung-Keyser (2020) show, the envelope theorem allows the welfare cost of a marginal policy change to be recovered from the mechanical burden it imposes, once its incidence is specified. Under the equal incidence benchmark, the welfare gradient therefore depends only on tax bases. On the other hand, the revenue gradient can be recovered from causal estimates of the elasticities of taxable income across estimates, which may include dynamics and spillovers across tax bases. Those dynamic general equilibrium effects are especially important in macroeconomic applications, where a tax change may affect both future revenue from its own

base and revenue collected through other instruments. The next example illustrates this logic in a familiar Ramsey setting.


To get at the economics behind the geometry, consider the canonical neoclassical model in which a government chooses a broad labor tax and a broad capital tax for a representative household with time-separable isoelastic preferences. In this representation, reform concerns the division of the tax burden between two broad income bases; it does not change the distribution of rates within either one. A permanent capital tax distorts saving in every subsequent period, affecting revenue not only from capital income but also from the future labor and capital income generated by accumulated capital. These dynamic fiscal effects make capital taxation especially costly in the Chamley (1986) and Chari, Nicolini, and Teles (2020) models. At the long-run Ramsey allocation, the capital tax is zero, and the welfare and revenue gradients are aligned across the two broad instruments. Away from that allocation, a positive capital tax creates scope for a revenue-neutral shift of the burden toward labor income more broadly.

However, policymakers and economists are typically just as interested in the distribution of a tax burden within an instrument—like the progressivity of the income tax system—as they are in its average burden, and I show that the same alignment statistics apply to both margins, and that adding progressivity must weakly reduce measured alignment. Alongside a broad reform, which changes how much revenue is raised through a tax base, a policymaker can shift tax payments within a schedule while holding total payments fixed before behavioral responses. That shift can nevertheless change total revenue. For example, increasing the top tax rate may induce passthrough businesses to become corporations, thereby increasing corporate tax revenue. These spillovers affect the revenue gradient, while any distributional weight the planner places on the burden affects the welfare gradient. Under equal incidence welfare weights, under which the planner equally values all sources of revenue, the first channel remains even when the second is absent: the planner may prefer a more progressive schedule not because redistribution has direct social value, but because its revenue effects permit a welfare-improving change in broad tax rates. Thus, the alignment metrics can measure alignment for systems with both linear and nonlinear instruments using the same types of welfare and revenue gradients for each.

The postwar transformation of the federal tax system provides a natural setting for this framework. At the end of World War II, corporate and personal income taxes together accounted for nearly all of federal tax revenue, while the tariff share was negligible. Until 2018, almost all reform happened on the income tax margin, with the corporate share falling from nearly 50 percent to below 15 percent and personal taxes absorbing the difference. Since then, tariff increases in 2018 and 2025 brought the customs share of tax revenue to nearly ten percent—well above the typical postwar share of 2-3 percent. At the same time, the personal tax schedule became less progressive, as the marginal tax rate faced by the top 1 percent fell relative to the lower 99 per-

cent. Those developments prompt two questions. First, was the shift from corporate to personal taxation welfare-improving? Second, how large a role should tariffs play in the federal tax system? Third, can changing the progressivity of the personal tax schedule create additional scope for welfare-improving, revenue-neutral reform?

Figure 1: Changes in Postwar Sources of Revenue and Tax Rates



Note: The left panel plots the average tax rates for personal income, corporate income, and imports. The middle panel plots the corresponding revenue shares. The right panel, using Mertens and Montiel Olea (2018), plots the difference between the top income percentile’s marginal tax rate and the marginal tax rate for the bottom 99%.

I study the alignment of the federal tax system by constructing welfare and revenue gradients for four margins: the broad personal, corporate, and tariff rates, as well as the progressivity of the personal tax schedule. Under an equal incidence benchmark, an application of the envelope theorem yields a welfare gradient for each of the three broad instruments equal to its incidence-adjusted observed tax base, while a change in progressivity has no direct first-order welfare effect. I therefore use equal incidence weights as my welfare benchmark for transparency but also conduct extensive robustness with other formulations. The bigger empirical challenge is the revenue gradient. In this setting, the central models and questions pertaining to the broad structure of the tax system imply that reforms may have important spillovers effects on other tax bases and instruments, as well as effects over time. I therefore use time series methods, together with exogenous variation in tax policy, to estimate the general equilibrium responses of tax bases to each instrument and construct the corresponding revenue semi-elasticities.

I use standard methods and narrative instruments to construct the revenue gradient, which can be decomposed into the three broad components and the progressivity entry. I estimate these components separately because the broad analysis uses quarterly variation in average tax rates,

while the progressivity analysis uses annual variation in tax rates across the income distribution. First, I use narrative instruments for personal and corporate tax changes from Mertens and Ravn (2013a), as well as tariff instruments from Besten et al. (2026), to estimate the impulse responses of each tax base to each broad tax instrument using Bayesian local projections (Cloyne et al. 2022). Cumulating those impulse responses and using the complete matrix of discounted tax-rate responses to re-express the matching base responses as responses to a change in one rate while holding the others fixed yields a 3×3 matrix of present-value elasticities of the three tax bases with respect to the three broad rates. These nine elasticities are a key input to the broad revenue gradient. Evaluating progressivity is similar, but requires annual data. I split the income distribution into the top 1% and the lower 99% using data from Mertens and Montiel Olea (2018). I then use the joint variation in those two rates across exogenous tax reforms to construct the semi-elasticities of the two personal income bases, as well as the corporate and tariff bases, with respect to a change that raises the top-1 tax rate relative to the lower-99 rate while holding personal tax revenue fixed. It is then straightforward to construct the revenue semi-elasticity for each of the four reform margins by combining the observed tax rates and tax bases with the estimated broad elasticities and progressivity semi-elasticities.

The resulting semi-elasticities of revenue for each instrument change substantially over the postwar era. The revenue semi-elasticity of personal taxation is by far the largest of the four margins and remains significantly positive throughout the sample. It starts around 2 and grows to roughly 3.5 by the 1980s, staying stable until after then. By contrast, the corporate tax and tariff margins have modestly negative point estimates in the early postwar era, though neither is precisely estimated then. Both rise over time to just below one: by the 1970s, the corporate tax margin is positive, and the tariff margin is positive at the 68-percent credible level. The progressivity margin moves in the opposite direction. Its revenue semi-elasticity is around one initially, but falls to zero from 1980 onward. All four movements are largely attributable to the long-run postwar decline in the corporate tax rate, most of which had occurred by the 1980s. Cuts in all three broad instruments increase corporate profits, so the fiscal externality of higher taxes on any of those three declines when the corporate tax rate itself declines. By contrast, higher progressivity increases corporate profits, likely in part through income shifting, so a lower corporate tax rate reduces the revenue-raising potential of marginally higher progressivity.

Finally, I evaluate the evolution of the federal tax system's local alignment with the Ramsey optimum over time. Under the equal incidence benchmark, I show that when we only consider personal and corporate taxes, the tax system was initially somewhat aligned in the 1950s, but the gains from reform on that margin fell to almost nothing by the 1980s. If we add in tariffs, then there is considerably more scope for alignment early in the period, but again, it becomes almost fully aligned by the 1980s, in part because tariff revenue is so small that it barely affects the

system locally. However, progressivity considerably changes the picture. In the full four-margin system, the median cosine draw is only 0.25 in 1950 and rises to roughly 0.8 by 1990, where it has stayed since, albeit with considerable posterior uncertainty. Thus, the system is largely aligned on the broad margins of how revenue is raised, but not necessarily when accounting for progressivity. That is reflected in the locally recommended reforms: throughout the postwar era, the optimal reform is to modestly cut personal taxes, while there is currently no indication that raising corporate taxes or tariffs is optimal.

In short, the application shows how the framework developed in the first part of the paper converts observable tax bases and causal responses to tax policy into an aggregate diagnostic for the proximity of a tax system to its local Ramsey condition, and for the corresponding direction of reform, without committing to a complete structural model. Nevertheless, there may be powerful ways to reform the system that lie outside the paper’s scope. There are two prominent examples. First, there is no federal variation in broad-based wealth or consumption taxes that would allow me to add them to the instrument mix. Second, the framework holds the government’s revenue requirement fixed in order to study the composition of revenue rather than its level. Changing that requirement would itself have a first-order effect on welfare. Both sorts of reforms—introducing new instruments and changing the government’s spending habits—are central parts of any debate on tax policy.

Related Literature. This paper primarily contributes to a classic literature on tax reform. The central insight of that literature is that the welfare consequences of small revenue-neutral reforms can be characterized using local welfare and revenue derivatives (Dixit 1975; Guesnerie 1977). The projection theorem of Tirole and Guesnerie (1981) shows that projecting the welfare gradient onto the revenue-neutral policy space identifies the best small reform under a specified policy metric.¹ I build on that projection theorem—and the geometric approach to tax reform—to show that the angle between the gradients is sufficient to identify the proximity of the tax system to the planner’s local optimality condition and the welfare gains from reforming in its direction. Hence, the cosine and sine of that angle are sufficient statistics for assessing tax systems as a whole. That remains true across many economies, not just in the static variants occasionally used by the classic literature to study commodity taxation (Dixit 1975). The same local sufficient-statistics logic extends to tax systems in which the distribution of the burden matters. In a static Diamond–Mirrlees setting, Saez and Zucman (2026) characterize the welfare effects of distributional tax reforms by separating mechanically distributed changes in tax payments from aggregate revenue effects of behavioral responses, while Ferey, Lockwood, and Taubinsky (2024) derive empirically implementable formulas for optimal nonlinear tax systems with across-income heterogeneity. I

1. Diewert (1978) and Dixit (1979) provide related many-consumer treatments.

show that the alignment representation accommodates these reforms as well, so that the sufficient statistics defined here apply broadly.

My approach to estimating the sufficient statistics builds on the marginal value of public funds literature’s insight that fiscal externalities can be estimated from causal evidence and that the welfare gradient is observable via the envelope theorem (Hendren 2016; Finkelstein and Hendren 2020). I estimate these fiscal externalities using narrative identification, constructing present-value revenue gradients from impulse responses to tax shocks.² In seeking apply the reform literature empirically, I build most closely on Ahmad and Stern (1984), which applied the commodity tax reform framework to Indian commodity taxes with a static demand system. Whereas their estimates required estimates of the full matrix of own- and cross-price demand responses, I estimate macro revenue semi-elasticities. This builds on McKay and Wolf (2023) and Caravello, McKay, and Wolf (2024), who show that impulse responses are sufficient to construct Lucas-critique-robust counterfactuals under alternative policy rules. My application is narrower, evaluating whether actual tax changes moved toward or away from the locally optimal direction rather than simulating counterfactual policies. The geometric approach—testing whether welfare and revenue gradients are aligned—also parallels Barnichon and Mesters (2023), who evaluate monetary policy by testing whether the gradient of the loss function with respect to policy shocks is zero; both frameworks detect suboptimal policies from sufficient statistics without requiring a fully specified structural model.

Finally, this paper brings together the sufficient statistics literature in public finance and the more parametric literature in macroeconomics. Existing approaches to evaluating tax reform face a trade-off between adequately capturing general equilibrium forces and avoiding parametric disagreement. Structural general equilibrium models capture both cross-base spillovers and dynamic accumulation, but require parametric commitments that can lead to fundamentally different conclusions. For example, a subclass of neoclassical models would conclude the postwar shift toward personal taxation was optimal (Chari, Nicolini, and Teles 2020), while overlapping generations models reach the opposite conclusion (Conesa, Kitao, and Krueger 2009). Even within the neoclassical framework, the optimal tax on capital—and hence the optimal direction of reform—depends sensitively on preference specifications and heterogeneity (Straub and Werning 2020). Sufficient statistics approaches avoid much of this model dependence by working directly with causal elasticities (Chetty 2009; Saez and Stantcheva 2018; Hendren 2016; Kleven 2021). Many empirical applications, however, rely on static own-base elasticities estimated from micro data (Saez 2001; Hendren and Sprung-Keyser 2020). Such estimates do not by themselves recover the cross-base

2. I build on the narrative approach to identification by modestly extending Mertens and Montiel Olea (2018) through the Tax Cuts and Jobs Act, as well as refining the Besten et al. (2026) tariff instruments at the quarterly level.

and dynamic spillovers that may be important for evaluating tax reform in certain contexts.³ Both approaches are generally used to characterize optimal tax rules or the direction of a particular reform rather than to summarize how much scope for local reform remains across an entire tax system.⁴ Rather than infer cross-base and dynamic responses from a structural model or approximate them with micro elasticities, I estimate them directly from causal variation in aggregate tax policy. The resulting revenue gradient, together with an incidence-adjusted welfare gradient, yields an aggregate diagnostic of local Ramsey alignment in a specified reform space without requiring a complete model of preferences or production.

2 The Geometry of Reform

When a tax system is not optimal, two questions follow: which direction to reform it, and how far it stands from the optimum. This section answers both with two sufficient statistics, the welfare gradient and the revenue gradient, which can be estimated from reduced-form evidence without specifying preferences, production, or market structure. The answer is geometric: the optimal direction of revenue-neutral reform is the orthogonal projection of the welfare gradient onto the revenue-neutral hyperplane, and the system's alignment with the Ramsey optimum is the cosine of the angle between the two gradients.

2.1 Setup: Taxes, Bases, and Gradients

Consider a government evaluating small changes to an existing tax system. The system may contain any number of taxes, and a tax change at one date may affect tax bases and revenue at other dates. Let $\{\tau_{i,t}\}_{i=1,\dots,n}^{t=0,\dots,T}$ denote n tax rates over $T + 1$ periods. The index i may refer either to a broad tax or to an individual rate within a nonlinear tax schedule. Stack these rates into the policy vector

$$\boldsymbol{\tau} := (\tau_{1,0}, \dots, \tau_{n,0}, \tau_{1,1}, \dots, \tau_{n,1}, \dots, \tau_{1,T}, \dots, \tau_{n,T})^\top \in \mathbb{R}^{n(T+1)}.$$

Let $B_{i,t}(\boldsymbol{\tau})$ denote the tax base associated with rate $\tau_{i,t}$. If instrument i indexes a tax rate within a nonlinear tax schedule, then $B_{i,t}$ is the amount of taxable income mechanically exposed to that particular rate. Importantly, I allow for tax bases to potentially depend on the entire policy sequence, reflecting the fact that tax reforms have spillovers across bases and over time. Consequently, a change to $\tau_{i,t}$ may affect not only base i at date t , but also other bases at other dates.

3. Saez, Slemrod, and Giertz (2012) likewise emphasize that tax responses may spill across instruments, in their case between capital gains and labor income, and Kleven et al. (2025) directly estimates dynamic elasticities.

4. Vergara and Swonder (2025), for example, incorporates cross-base fiscal externalities from corporate taxes on the labor base in a static optimal corporate tax formula.

For example, a corporate tax cut today may raise investment, increasing the capital stock and expanding future labor tax bases through production complementarity. If that cut was anticipated, then it would also affect tax bases before the reform even happens. Aggregating over bases and tax rates, it is convenient to define the present value of consolidated revenue as

$$R(\boldsymbol{\tau}) := \sum_{t=0}^T \beta^t \sum_{i=1}^n \tau_{i,t} B_{i,t}(\boldsymbol{\tau}), \quad (1)$$

where $0 < \beta \leq 1$ is the discount factor. I assume that tax bases respond smoothly to small changes in policy.

Assumption 1. *Each tax base $B_{i,t}(\boldsymbol{\tau})$ is continuously differentiable in $\boldsymbol{\tau}$ in a neighborhood of $\bar{\boldsymbol{\tau}}$.*

This assumption is satisfied in standard models. As long as the equilibrium is unique and fairly well behaved in the neighborhood of policy, it will hold. For example, standard neoclassical and overlapping generations models typically satisfy this condition. The key exclusion is an economy with lumpy behavior in the aggregate, such as a model with multiple equilibria in which a small tax change triggers a discrete jump. Because present value revenue in equation (1) is constructed from these tax bases, the assumption also implies that $R(\boldsymbol{\tau})$ is continuously differentiable. Under that assumption, we can characterize how present value revenue changes directly from the response of each tax base.

The revenue gradient. Given differentiability, define the revenue gradient $\mathbf{r} := \nabla_{\boldsymbol{\tau}} R \in \mathbb{R}^{n(T+1)}$, which measures how present value revenue responds to marginal tax changes. Conceptually, it is useful to consider two separate decompositions of this response because they illustrate different dimensions of general equilibrium. The first decomposes the revenue response into three components. Suppose $\tau_{i,t}$ rises. It has a mechanical effect from taxing the existing base, a behavioral effect through the same tax base, and spillovers to other tax bases. Differentiating equation (1) gives

$$r_{i,t} = \underbrace{\beta^t B_{i,t}}_{\text{Mechanical effect}} + \underbrace{\sum_{s=0}^T \beta^s \tau_{i,s} \frac{\partial B_{i,s}}{\partial \tau_{i,t}}}_{\text{Own-base responses}} + \underbrace{\sum_{s=0}^T \beta^s \sum_{k \neq i} \tau_{k,s} \frac{\partial B_{k,s}}{\partial \tau_{i,t}}}_{\text{Cross-base spillovers}}. \quad (2)$$

The mechanical effect is the static revenue gain, while the second term includes all dynamic own-base responses, and the final term includes spillovers to other tax bases across dates. Second, the same revenue response can instead be decomposed according to when it occurs. Let $R_s(\boldsymbol{\tau}) :=$

$\sum_{k=1}^n \tau_{k,s} B_{k,s}(\boldsymbol{\tau})$ denote revenue collected at date s . Then

$$r_{i,t} = \underbrace{\beta^t \frac{\partial R_t}{\partial \tau_{i,t}}}_{\text{Contemporaneous}} + \underbrace{\sum_{s>t} \beta^s \frac{\partial R_s}{\partial \tau_{i,t}}}_{\text{Persistence}} + \underbrace{\sum_{s<t} \beta^s \frac{\partial R_s}{\partial \tau_{i,t}}}_{\text{Anticipation}}. \quad (3)$$

Equations (2) and (3) provide two complementary decompositions of the same revenue gradient. The first distinguishes how revenue changes across tax bases, while the second distinguishes when those effects occur. Existing empirical work often focuses on a particular own-base response, usually in a static partial equilibrium setting, rather than the full revenue gradient. The following two examples illustrate what the additional terms capture.

Example 1: Human capital accumulation (persistence, own base). As a first example, consider an increase in the labor tax at $t = 0$. In standard models, that may reduce human capital investment because it would lower the return to skills, lowering workers' future productivity and taxable earnings. As a result, the tax base would decline persistently beyond the reform:

$$\frac{\partial B_{L,t}}{\partial \tau_{L,0}} < 0 \quad \text{for all } t > 0. \quad (4)$$

These effects enter the own-base term in equation (2) and the persistence term in equation (3). In present value, revenue declines, and the gradient $r_{L,0}$ is smaller than the contemporaneous revenue response alone would suggest. Indeed, theoretical and quantitative work from Badel, Huggett, and Luo (2020) and empirical work from Kleven et al. (2025) suggest that accounting for this channel is particularly important, so the revenue gradient must incorporate dynamic tax base responses rather than only contemporaneous responses.

Example 2: Capital accumulation (persistence, cross base). As a second example, consider an increase in the corporate tax rate. In most models—and in canonical models for scoring tax reform (Hall and Jorgenson 1967)—that will reduce investment and therefore shrink the capital stock over time. If capital and labor are complements in production ($F_{KL} > 0$), the lower capital stock reduces the marginal product of labor and may thereby erode the future labor tax base:

$$\sum_{t=1}^T \beta^t \tau_{L,t} \frac{\partial B_{L,t}}{\partial \tau_{K,0}} < 0.$$

As long as future labor tax rates are positive, the channel contributes to the present-value revenue gradient $r_{K,0}$ the quantity

$$\sum_{t=1}^T \beta^t \tau_{L,t} \frac{\partial B_{L,t}}{\partial \tau_{K,0}} < 0.$$

Once more, this contribution appears in the cross-base term of equation (2) and in the persistence term of equation (3). It is distinct from both the mechanical revenue effect of the corporate tax change and its effects on the corporate tax base itself. Indeed, production complementarity of this kind provides a central supply-side rationale for reductions in capital and corporate income taxes (Lucas 1990).⁵

The revenue gradient describes what happens to revenue when taxes change on the margin. But the purpose of reform is also to change welfare. Thus, the next piece of the setup is what happens to welfare when taxes change on the margin—the welfare gradient.

The welfare gradient. Let $W(\boldsymbol{\tau})$ denote the money-metric social welfare cost of the tax system, measured in present-value dollars, with higher values indicating lower welfare. We construct W by aggregating present-value compensating variation across groups using social welfare weights $\{\omega_h\}$, where groups are a potentially arbitrary population. For example, they could be domestic residents versus foreigners, or the top-1% versus the bottom 99% of the income distribution. Regardless, we will eventually want to incorporate how welfare changes when taxes change, which requires an assumption on the differentiability of welfare.

Assumption 2 (Welfare). $W(\boldsymbol{\tau})$ is continuously differentiable in $\boldsymbol{\tau}$ in a neighborhood of $\bar{\boldsymbol{\tau}}$.

Mirroring the revenue gradient, the welfare gradient captures how welfare changes when taxes change on the margin:

$$\mathbf{w} := \nabla_{\boldsymbol{\tau}} W \in \mathbb{R}^{n(T+1)}.$$

Consider a small increase $d\tau_{i,t}$. Holding the tax base fixed, the resulting increase in tax payments at date t is $B_{i,t}d\tau_{i,t}$, or $\beta^t B_{i,t}d\tau_{i,t}$ in present value dollars. Locally, agents are already choosing across all margins optimally given the tax system. The envelope theorem therefore implies that their behavioral adjustments do not generate a separate first-order term in welfare. Those adjustments can still change tax bases and government revenue, which is why they enter the revenue gradient (Finkelstein and Hendren 2020).

The remaining issue is how the mechanical burden is distributed across groups and how the planner values those groups. Let $s_{h,i,t}$ denote group h 's present value burden per dollar of mechanical burden. That burden may operate through direct tax payments or through changes in wages, prices, and profits. A negative value of $s_{h,i,t}$ means that the group benefits from the tax

5. Anticipation provides another intertemporal channel. If agents learn at date 0 that the capital tax will change at date 5, they may adjust beforehand, so that $\partial B_{K,t}/\partial \tau_{K,5} \neq 0$ for $t < 5$. These responses enter the anticipation term in equation (3) and can make the present-value revenue effect differ from its contemporaneous counterpart. See Mertens and Ravn (2012), Ramey (2016), and Leeper, Richter, and Walker (2012) for the relevant theoretical discussion and empirics.

change. Given social welfare weights $\{\omega_h\}$, define

$$\tilde{\gamma}_{i,t} := \sum_h \omega_h s_{h,i,t}.$$

Aggregating the weighted burdens across groups gives

$$w_{i,t}(\bar{\tau}) = \beta^t \tilde{\gamma}_{i,t} B_{i,t}(\bar{\tau}). \quad (5)$$

Thus, $\beta^t B_{i,t}$ determines the size of the mechanical burden, while $\tilde{\gamma}_{i,t}$ captures how that burden is distributed and socially valued.⁶ When the groups represented in W collectively bear the entire burden, $\sum_h s_{h,i,t} = 1$. That need not be the case when, for example, part of a tariff is borne abroad and foreign welfare is excluded from W . A useful benchmark, and one which I will rely on later in the empirics, assigns equal social welfare weights, $\omega_h = 1$, to all groups whose welfare is included. If those groups bear the full marginal burden of tax $\tau_{i,t}$, then

$$\sum_h s_{h,i,t} = 1,$$

so that $\tilde{\gamma}_{i,t} = 1$ and the welfare gradient simplifies to

$$w_{i,t} = \beta^t B_{i,t}. \quad (6)$$

This representation is important for the empirical implementation in this paper, which I discuss later.

The welfare representation here comes with a major caveat. Like other derivations of the marginal value of public funds, it requires all first-order consequences of a tax change to be included either in the measured burdens or in consolidated government revenue (Hendren 2016). However, if tax changes interact with an externality, market power, or another distortion that creates additional welfare consequences, those consequences must be incorporated into the welfare gradient.⁷ In the empirical application, I will maintain that these additional welfare effects are either absent or small enough to ignore. This is an important restriction, but it is the same one maintained in the MVPF literature.

Under the maintained assumptions, the distinction between the welfare gradient and the rev-

6. Alternative approaches use observed tax policy to infer the social welfare weights that would rationalize it; see Lockwood and Weinzierl (2016) and Vergara and Swonder (2025) for examples. I instead evaluate local alignment under specified welfare weights.

7. For example, the effects of federal tax changes on state and local budgets should be included in consolidated revenue; pollution caused by an expansion in production creates an additional welfare effect; and changes in foreign prices following a tariff reform affect domestic welfare through the terms of trade.

enue gradient is exact to first order: $\mathbf{w}$ measures the socially weighted burden on the groups represented in W , while $\mathbf{r}$ measures the change in government resources. Together, they suffice to measure the welfare cost of a given tax instrument through the marginal value of public funds (Mayshar 1990; Bastani 2025).

The marginal value of public funds. For any tax rate with a nonzero marginal revenue effect, the welfare and revenue gradients define a welfare-weighted marginal value of public funds.

$$\text{MVPF}_{i,t} := \frac{w_{i,t}}{r_{i,t}} = \frac{\beta^t \tilde{y}_{i,t} B_{i,t}}{r_{i,t}}, \quad r_{i,t} \neq 0. \quad (7)$$

When $r_{i,t} > 0$, this ratio measures the marginal social burden created per dollar of present-value revenue raised by increasing $\tau_{i,t}$. Because the social welfare weights are already incorporated into $w_{i,t}$ through $\tilde{y}_{i,t}$, no additional welfare weight is applied to the ratio.

At an interior Ramsey optimum, the welfare and revenue gradients are positively proportional. Consequently, the MVPFs in equation (7) are equal across all tax rates for which they are defined (Appendix A.3 derives this condition formally). When the gradients are not proportional, there is scope for a welfare-improving revenue-neutral reform. The next subsection states this condition directly in terms of the two gradients.

2.2 Ramsey Alignment

Having characterized the welfare and revenue effects of marginal tax changes, I now proceed in two steps. First, I ask how those changes should be combined to generate an optimal local reform, which is an old question in public finance (Feldstein 1976). Second, I use the solution to that problem to provide a natural way to summarize how closely the existing system satisfies the Ramsey condition.

I start by considering local reforms. Let $d\boldsymbol{\tau}$ denote a small reform around the baseline policy vector $\bar{\boldsymbol{\tau}}$. To first order, the welfare gain associated with that change is $-\mathbf{w}(\bar{\boldsymbol{\tau}})^\top d\boldsymbol{\tau}$, while its effect on present-value revenue is $\mathbf{r}(\bar{\boldsymbol{\tau}})^\top d\boldsymbol{\tau}$. The reform problem holds revenue fixed:

$$\mathbf{r}(\bar{\boldsymbol{\tau}})^\top d\boldsymbol{\tau} = 0.$$

We are primarily interested in identifying the direction of reform, but will not be able to pin down how far the government should move in that direction. Consequently, I compare reforms of a common maximum size $\alpha > 0$ and restrict attention to reforms satisfying $\|d\boldsymbol{\tau}\| \leq \alpha$, where each included reform direction is measured in percentage point units and $\|\cdot\|$ denotes the Euclidean norm. For a change in a single tax rate, this is simply the percentage point change in that rate. The

next subsection defines the corresponding measure when a reform changes several rates within the same instrument. The planner's local reform problem is

$$\max_{d\boldsymbol{\tau}} \quad -\mathbf{w}(\bar{\boldsymbol{\tau}})^\top d\boldsymbol{\tau} \quad \text{subject to} \quad \mathbf{r}(\bar{\boldsymbol{\tau}})^\top d\boldsymbol{\tau} = 0, \quad \|d\boldsymbol{\tau}\| \leq \alpha. \quad (8)$$

Equation (8) separates the two parts of the reform problem. The revenue gradient determines which combinations of tax changes hold revenue fixed, while the welfare gradient ranks those combinations by their welfare effects. The following lemma characterizes the best one.

Lemma 1 (Best local revenue-neutral reform). *Let $\bar{\boldsymbol{\tau}}$ be an interior baseline, and write $\mathbf{w} := \mathbf{w}(\bar{\boldsymbol{\tau}})$ and $\mathbf{r} := \mathbf{r}(\bar{\boldsymbol{\tau}})$, with $\mathbf{r} \neq 0$. Define*

$$P_{\perp \mathbf{r}} := I - \mathbf{r}(\mathbf{r}^\top \mathbf{r})^{-1} \mathbf{r}^\top.$$

If $P_{\perp \mathbf{r}} \mathbf{w} \neq 0$, the unique solution to equation (8) is

$$d\boldsymbol{\tau}^* = -\alpha \frac{P_{\perp \mathbf{r}} \mathbf{w}}{\|P_{\perp \mathbf{r}} \mathbf{w}\|}. \quad (9)$$

The resulting first-order welfare gain is

$$-\mathbf{w}^\top d\boldsymbol{\tau}^* = \alpha \|P_{\perp \mathbf{r}} \mathbf{w}\|. \quad (10)$$

If $P_{\perp \mathbf{r}} \mathbf{w} = 0$, every first-order revenue-neutral reform has zero first-order welfare effect. In that case, $d\boldsymbol{\tau}^ = 0$ may be selected.*

To see the result, begin with the unrestricted reform direction $-\mathbf{w}$, which produces the largest first-order reduction in welfare cost. In general, that reform also changes revenue, which eliminates it from the reform set. However, if we project $\mathbf{w}$ onto the hyperplane perpendicular to $\mathbf{r}$, then we remove the component associated with that change in revenue. That leaves $P_{\perp \mathbf{r}} \mathbf{w}$, the component of the welfare gradient that can be changed while holding revenue fixed to first order. Reversing and normalizing this component gives the reform in the lemma, while its length determines the available welfare gain for a given reform size α . Appendix A.4 provides the proof.

The solution to this problem is not new. Variations on it exist in Dixit (1979), which studies a similar reform problem in a many-consumer economy, while Tirole and Guesnerie (1981) finds a similar result in developing a gradient-projection process for revenue-neutral reform of linear commodity taxes. However, one difference, which comes down to labeling, is that the object projected here is the gradient of present-value consolidated revenue, which includes behavioral effects across tax bases and across dates. The resulting reform direction therefore incorporates intertemporal and cross-base revenue effects, something which is not apparent in the early reform literature.

One of the theoretical contributions of this paper is using the projection result to show how to assess the tax system itself. At an interior Ramsey optimum, the welfare and revenue gradients point in the same direction. That is akin to saying that welfare and revenue are perfectly aligned, which requires that the revenue-neutral projection is zero. Geometrically, however, it means that the angle between the gradients summarizes how aligned they are, and that is a measurable quantity.

Proposition 1 (Ramsey alignment). *Suppose $\mathbf{w} \neq 0$ and $\mathbf{r} \neq 0$, and let $\theta \in [0, \pi]$ denote the angle between them. Then*

$$\cos \theta = \frac{\mathbf{w}^\top \mathbf{r}}{\|\mathbf{w}\| \|\mathbf{r}\|}, \quad \sin \theta = \frac{\|P_{\perp \mathbf{r}} \mathbf{w}\|}{\|\mathbf{w}\|}. \quad (11)$$

For the best revenue-neutral reform satisfying $\|d\boldsymbol{\tau}\| \leq \alpha$,

$$-\mathbf{w}^\top d\boldsymbol{\tau}^* = \alpha \|\mathbf{w}\| \sin \theta. \quad (12)$$

Moreover, $\cos \theta = 1$ if and only if $\mathbf{w}$ and $\mathbf{r}$ are positively proportional.

The proof is straightforward.

Proof. The expression for $\cos \theta$ is the inner-product definition of the angle. Because $P_{\perp \mathbf{r}} \mathbf{w}$ is orthogonal to the projection of $\mathbf{w}$ onto $\mathbf{r}$,

$$\|P_{\perp \mathbf{r}} \mathbf{w}\|^2 = \|\mathbf{w}\|^2 - \frac{(\mathbf{w}^\top \mathbf{r})^2}{\|\mathbf{r}\|^2} = \|\mathbf{w}\|^2 (1 - \cos^2 \theta) = \|\mathbf{w}\|^2 \sin^2 \theta.$$

Equation (12) then follows from Lemma 1. The final statement follows from the equality condition for the Cauchy–Schwarz inequality. $\square$

Crucially, as I discuss later in more detail, θ is a measurable object and does not require a complete structural model. The welfare gradient can be constructed from baseline tax bases and the weights placed on their burdens, while the revenue gradient can be constructed from estimated responses of tax bases. Once both gradients have been constructed for the same tax changes, substituting them into equation (11) gives a direct measure of local Ramsey alignment. Beyond measurement, however, there is a clear economic interpretation. Consider $\sin \theta$. If revenue were allowed to change, then the best reform of size α would produce a first-order welfare gain of $\alpha \|\mathbf{w}\|$. Requiring the reform to be revenue neutral reduces that gain to $\alpha \|\mathbf{w}\| \sin \theta$. Thus, $\sin \theta$ measures how much of the best local welfare improvement remains available when revenue must be held fixed.

On the other hand, the cosine of that angle is required to interpret cases in which there may be no revenue-neutral gain. To see that, note that $\cos \theta = 1$ and $\cos \theta = -1$ produce the same

measured $\sin \theta$. If $\cos \theta = 1$, then the welfare and revenue gradients point in the same direction and the tax system satisfies the local Ramsey first-order condition. For example, Chamley (1986) shows that the Ramsey allocation sets the capital tax to zero in the long run because a permanent capital tax compounds distortions to saving over time. In the gradient representation, the welfare and revenue gradients are aligned at that allocation.

On the other hand, if $\cos \theta = -1$, then the welfare and revenue gradients point in opposite directions and there is no revenue-neutral gain from reform, so $\sin \theta = 0$. Under the equal-incidence benchmark, that corresponds to being on the wrong side of the Laffer curve: reducing tax rates raises both revenue and welfare.⁸ In the intermediate case, in which $\cos \theta = 0$, the welfare and revenue gradients are orthogonal. That means that the unconstrained welfare-improving direction $-\mathbf{w}$ is already revenue neutral.

Next, I use the same geometric approach to distinguish reforms across broad tax instruments from reforms that change rates within an instrument.

2.3 Alignment Across and Within Tax Instruments

Broadly, policymakers may be interested in two types of reform. The first concerns the burden imposed through a broad tax instrument. The second concerns the distribution of that burden within the instrument. In the first type of reform, the policymaker would be interested in changing, for example, how much revenue is raised via income taxation. The second type of reform would hold fixed the total income tax payments that would be collected before behavioral responses and alter the progressivity of the income tax. Many important policy questions concern the second type of reform. Beyond the progressivity of income taxes, one may worry about the composition of tariffs across product groups or the tariffs levied against certain countries, as well as the distribution of business taxes across different legal entities, industries, and so on. Importantly, although the preceding subsection did not distinguish between these two sets of questions, the degree of alignment with the local Ramsey optimum can be meaningfully measured for both using the same types of inputs.

The alignment criterion from the previous subsection applies whether a reform changes broad tax rates or rates within an instrument. What remains is to separate these two types of change and measure them together. Throughout this subsection, I treat reform as a bundled problem. Rather than require revenue-neutrality within a change in broad tax rates or within a distributional reform, I require neutrality across the reform as a whole. Thus, for example, the optimal reform may cut corporate taxes relative to personal taxes but finance the cut with greater progressivity.

8. This is one of the key advantages over using raw MVPFs; they become undefined on the wrong side of the Laffer curve, whereas tax instruments remain meaningful even on the wrong side of the Laffer curve.

Let $d\tau_A$ denote the broad component of a reform and $d\tau_S$ the within-instrument component, and let dR_A and dR_S denote their respective first-order revenue effects. Overall revenue neutrality requires

$$dR_A + dR_S = 0. \quad (13)$$

Neither component must be revenue-neutral separately; only their combined revenue effect must be zero.

Within each instrument, the distinction between the two components turns on whether a reform changes total tax payments at the existing tax bases or changes the distribution of those payments across rates. Consider an instrument containing J rates at date t . Any set of rate changes can be separated into a common change and changes in each rate relative to that common change. Choose the common change to reproduce the same change in total tax payments at the existing tax bases:

$$d\tau_{A,t} = \frac{\sum_{j=1}^J B_{j,t} d\tau_{j,t}}{\sum_{j=1}^J B_{j,t}}, \quad d\tau_{S,j,t} = d\tau_{j,t} - d\tau_{A,t}. \quad (14)$$

The first expression defines the broad component, while the second assigns the remaining changes to the within-instrument component. Multiplying the second expression by the existing tax bases and summing across rates gives

$$\sum_{j=1}^J \beta^t B_{j,t} d\tau_{S,j,t} = 0. \quad (15)$$

Thus, the within-instrument component changes the distribution of tax payments across rates without changing their total at the existing tax bases. This condition does not require a reform which changes the within-instrument distribution of tax burden to be revenue-neutral, but any changes the reform yields must be offset by other instruments. For example, making taxes more progressive could raise revenue as a whole through cross-base effects, but that would have to be compensated by cuts elsewhere.

The base weights in equation (15) also enter the welfare gradient. Applying the definition from the setup, the first-order welfare effect of a within-instrument change is

$$dW_S = \sum_{j=1}^J \beta^t \tilde{\gamma}_{j,t} B_{j,t} d\tau_{S,j,t}. \quad (16)$$

When the incidence-adjusted welfare weights $\tilde{\gamma}_{j,t}$ differ across rates, changes that offset in dollars need not offset in socially weighted welfare. When the planner applies equal weights and the groups represented in W collectively bear the full marginal burden, however, $\tilde{\gamma}_{j,t} = 1$ for every rate. In that case, equation (15) implies $dW_S = 0$ for every within-instrument direction. This

special case is an important benchmark because it implies that reforming toward progressivity may be optimal even without distributional motives. A distributional reform that raises revenue can finance a tax cut elsewhere along a margin with a positive welfare cost. For example, a planner may place no direct value on progressivity but still prefer a more progressive schedule because it allows taxes on another base to be reduced. Nevertheless, this benchmark is only an illustrative case.

To combine the two types of reform in the same alignment metric, I need to measure their size in the same units. I measure the size of a within-instrument reform by the square root of the average squared percentage-point changes in its rates, weighting each rate by its share of the existing tax base. This preserves the percentage-point interpretation from the previous subsection: if all the rates move together, the measure is simply their common change. It also ensures that the measured size of a reform does not change simply because the same instrument is divided into more rates. Equation (15) says that the weighted average of the within-instrument changes is zero. Because the broad component changes every rate by the same amount, the two components are orthogonal under this measure.

Appendix A.6 proves this result and defines the corresponding coordinates formally. In those coordinates, write the complete welfare and revenue gradients as

$$\mathbf{w} = \begin{pmatrix} \mathbf{w}_A \\ \mathbf{w}_S \end{pmatrix}, \quad \mathbf{r} = \begin{pmatrix} \mathbf{r}_A \\ \mathbf{r}_S \end{pmatrix}.$$

The equal incidence result above then implies $\mathbf{w}_S = \mathbf{0}$. Suppose $\mathbf{w} \neq \mathbf{0}$ and $\mathbf{r} \neq \mathbf{0}$. At a regular interior Ramsey optimum, provided that a larger revenue requirement increases welfare cost, there must exist some $\mu > 0$ such that

$$\mathbf{w}_A = \mu \mathbf{r}_A, \quad \mathbf{w}_S = \mu \mathbf{r}_S.$$

Following the notation of the previous subsection, let θ denote the angle between the complete welfare and revenue gradients. When $\mathbf{w}_A$ and $\mathbf{r}_A$ are nonzero, let θ_A denote the angle between their broad components.

Proposition 2 (Alignment with within-instrument reforms). *The alignment metric over the complete reform space is*

$$\cos \theta = \frac{\mathbf{w}_A^\top \mathbf{r}_A + \mathbf{w}_S^\top \mathbf{r}_S}{\sqrt{\|\mathbf{w}_A\|^2 + \|\mathbf{w}_S\|^2} \sqrt{\|\mathbf{r}_A\|^2 + \|\mathbf{r}_S\|^2}}. \quad (17)$$

Under the equal incidence benchmark, $\mathbf{w}_S = \mathbf{0}$. If $\mathbf{w}_A$ and $\mathbf{r}_A$ are nonzero, this specializes to

$$\cos \theta = \cos \theta_A \frac{\|\mathbf{r}_A\|}{\sqrt{\|\mathbf{r}_A\|^2 + \|\mathbf{r}_S\|^2}}. \quad (18)$$

Consequently, under this benchmark, $\cos \theta = 1$ if and only if $\cos \theta_A = 1$ and $\mathbf{r}_S = \mathbf{0}$.

Appendix A.6 provides the proof and derives the corresponding expression for the scope of revenue-neutral reform below.

Compared to the broad notion of alignment in the previous subsection, adding within-instrument reforms can change alignment through both gradients. With general welfare weights, a distributional change may have a direct welfare effect through $\mathbf{w}_S$ as well as a revenue effect through $\mathbf{r}_S$. The result is sharpest under the equal incidence benchmark, when $\mathbf{w}_S = \mathbf{0}$. In that case, a nonzero $\mathbf{r}_S$ adds a component to the revenue gradient without adding one to the welfare gradient. Consequently, when broad alignment is positive, full alignment equals broad alignment multiplied by the fraction of the revenue gradient's magnitude that comes from the broad tax directions. Intuitively, even though the planner does not value progressivity directly, it may still prefer a more progressive schedule if shifting the burden toward higher rates raises revenue because rates at different points in the schedule generate different behavioral responses and fiscal spillovers. This creates a revenue margin that is absent from the broad tax calculation $\cos \theta_A$. Even if $\cos \theta_A = 1$, the system may be misaligned with the Ramsey optimum because there are gains from changing the burden within an instrument.

The same logic applies to the scope for revenue-neutral reform. Under the equal incidence benchmark, adding within-instrument reforms cannot reduce that scope:

$$\sin^2 \theta = \sin^2 \theta_A + \cos^2 \theta_A \frac{\|\mathbf{r}_S\|^2}{\|\mathbf{r}_A\|^2 + \|\mathbf{r}_S\|^2}. \quad (19)$$

The additional term is positive whenever $\mathbf{r}_S \neq \mathbf{0}$ and $\cos \theta_A \neq 0$. This term is larger when within-instrument revenue effects are large relative to broad revenue effects and when the broad reform space leaves less scope for revenue-neutral reform. As a result, a within-instrument reform can finance a welfare-improving change in broad tax rates without adding a direct first-order welfare cost of its own.

Taken together, the framework provides a local criterion for assessing alignment with a generic Ramsey planner's first-order condition across rich instrument sets. The next section shows how to apply the theory empirically.

3 Estimation of the Revenue Gradient

As an application of the geometric framework, I evaluate the alignment of the U.S. federal tax system with the Ramsey optimum. Throughout the postwar era, personal and corporate income taxes accounted for over 90 percent of federal tax revenue, with tariffs contributing a negligible share. Figure 1 shows that until quite recently, most reform happened on the income tax mix, with corporate rates and revenue share falling dramatically and personal taxes rising modestly to pick up the burden. At the same time, tax policy appears to have become less progressive over time, with the marginal tax rate of the top-1% falling relative to the bottom-99%. That motivates a retrospective look at whether the federal tax system became more or less aligned over time, both across broad tax instruments, as well as within the personal tax mix.

3.1 From Theory to Empirics

I consider a version of the planner's problem in which he has access to four instruments at each time t : a broad corporate tax rate $\tau_{C,t}$, a personal income tax rate $\tau_{P,t}$, a tariff $\tau_{T,t}$, and a progressivity instrument d_t . Each of the first three instruments is an average tax rate, while the progressivity instrument allows the planner to alter the tax rates paid by the top 1% relative to the lower 99% of the income distribution, holding initial personal tax payments fixed. There are then two empirical objects: the welfare gradient

$$\mathbf{w}_t = \begin{pmatrix} w_{C,t} \\ w_{P,t} \\ w_{T,t} \\ w_{d,t} \end{pmatrix}$$

and the revenue gradient

$$\mathbf{r}_t = \begin{pmatrix} r_{C,t} \\ r_{P,t} \\ r_{T,t} \\ r_{d,t} \end{pmatrix}.$$

Following the decomposition in Section 2.3, write

$$\mathbf{w}_t = \begin{pmatrix} \mathbf{w}_{A,t} \\ w_{S,t} \end{pmatrix}, \quad \mathbf{r}_t = \begin{pmatrix} \mathbf{r}_{A,t} \\ r_{S,t} \end{pmatrix},$$

where

$$\mathbf{w}_{A,t} = \begin{pmatrix} w_{C,t} \\ w_{P,t} \\ w_{T,t} \end{pmatrix}, \quad \mathbf{r}_{A,t} = \begin{pmatrix} r_{C,t} \\ r_{P,t} \\ r_{T,t} \end{pmatrix},$$

and $w_{S,t} = w_{d,t}$ and $r_{S,t} = r_{d,t}$.

As a baseline, I take advantage of the envelope theorem and apply equal incidence weights across the groups bearing each of the three broad tax burdens. That means the mechanical welfare cost of raising each tax rate is proportional to its observable tax base in the national accounts data, i.e., $w_{i,t} = \tilde{B}_{i,t}$ for $i = C, P, T$.⁹ Under equal incidence, the progressivity instrument has no direct first-order welfare effect, so $w_{d,t} = 0$. I discuss these assumptions later in the paper, as well as in the appendix. I also vary the equal incidence assumption to allow for preferences over progressivity.

The more difficult object is the revenue gradient. Its construction requires present-value semi-elasticities of the tax bases with respect to each instrument, accounting for dynamic adjustments and spillovers across bases. In this context, that is particularly important for three reasons. First, the instruments themselves have economically important effects that unfold over time. For example, a corporate tax cut should increase investment, which raises the corporate tax base dynamically (Hall and Jorgenson 1967). Second, the instruments probably have large cross-base effects. For instance, corporate tax cuts may affect the personal tax base by causing wages to rise over time, as in the canonical model, or they may cause the personal tax base to fall due to organizational form shifting (Barro and Wheaton 2019). Third, I am studying a national tax system, and that system naturally has important aggregate effects. Thus, the revenue gradient should account for all of these spillovers. As an example, the corporate tax's entry in the revenue gradient is, after suppressing time subscripts,

$$r_C = \tilde{B}_C \left[1 + \tau_C \varepsilon_{CC} + \tau_P \frac{\tilde{B}_P}{\tilde{B}_C} \varepsilon_{PC} + \tau_T \frac{\tilde{B}_T}{\tilde{B}_C} \varepsilon_{TC} \right]. \quad (20)$$

Here, ε_{kC} is the present-value semi-elasticity of base k with respect to the corporate tax rate.¹⁰

9. Equal incidence is a potentially important assumption for tariffs and corporate taxes. For the former, the burden should in principle be split between domestic consumers and foreign exporters, with the division depending on terms-of-trade effects. Recent evidence on the 2018–19 trade war with China indicates that the United States bore around 60–80 percent of the incidence (Ganapati and Hottman 2026), while a significant number of studies show full passthrough (Amiti, Redding, and Weinstein 2019, 2020; Cavallo et al. 2021; Fajgelbaum et al. 2020). At baseline, I assume that domestic consumers bear the full burden of tariffs and show the effect of varying that assumption. For corporate taxes, many corporations are owned in part by foreigners, and under a nationalist planner the welfare cost borne by foreign owners should be excluded from the baseline under equal incidence.

10. I report semi-elasticities as the percent change in the tax base produced by a one-percentage-point increase in the relevant tax instrument.

The first term in brackets is the mechanical revenue gain. The second captures the present-value response of the corporate tax base. The remaining terms are the cross-semi-elasticities of the personal and tariff bases, scaled by the tax rate applying to each base and its size relative to the corporate base. That matters since tariff revenue is small, which means that the tariff cross-term receives little weight even if the semi-elasticity is large. The same basic formula holds for the entries of tariffs and personal taxes in the revenue gradient. The broad entries hold d fixed by definition, since the reform space separates the level of each instrument from the distribution of rates within it.

While the broad entries in $\mathbf{r}_{A,t}$ evaluate changes in the overall level of each tax instrument while holding fixed the distribution of rates within the personal tax schedule, the distributive entry $r_{d,t}$ asks whether a change in that distribution affects revenue. I define d to raise the tax rate paid by the top 1% relative to the lower 99%, while reducing rates elsewhere so that initial personal tax payments remain fixed. I implement this direction using two income groups—the lower 99% and the top 1%—which enter separately because their tax rates, tax bases, and behavioral responses can differ. This yields

$$r_d = \tilde{B}_P \left[\tau_L \frac{\tilde{B}_L}{\tilde{B}_P} \varepsilon_{Ld} + \tau_H \frac{\tilde{B}_H}{\tilde{B}_P} \varepsilon_{Hd} + \tau_C \frac{\tilde{B}_C}{\tilde{B}_P} \varepsilon_{Cd} + \tau_T \frac{\tilde{B}_T}{\tilde{B}_P} \varepsilon_{Td} \right]. \quad (21)$$

Here, τ_L and τ_H are the tax rates paid by the lower 99% and top 1%, $\tilde{B}_L$ and $\tilde{B}_H$ are their discounted bases, and ε_{kd} denotes the present-value semi-elasticity of base k with respect to d . There is no mechanical revenue term because the change in initial personal tax payments offsets across the two groups by construction. Consequently, every term in equation (21) reflects a behavioral response. The first two terms capture the within-base behavioral response to a progressivity shock, while the final two terms capture spillovers to the corporate and tariff bases.

Given equations (20) and (21), the difficult empirical task is obtaining the relevant semi-elasticities. The next subsection presents the common identification and estimation procedure used to recover them. I then apply that procedure to the broad tax terms using inputs from Mertens and Ravn (2013a) and Besten et al. (2026), and to the progressivity reform using inputs from Mertens and Montiel Olea (2018).

3.2 Identification and Estimation

The formulas in the previous subsection reduce the empirical task to estimating the present-value response of each tax base to each relevant tax rate change. However, theory and data constraints require separate estimation of the broad tax and progressivity semi-elasticities. Theoretically, however, they are different reforms. The broad tax reforms describe changes in the overall level

of each tax instrument, and the progressivity reform changes the within-instrument distribution of the tax burden. Consequently, they cannot be identified simultaneously because one requires holding the income tax fixed. As discussed in the following sections, the data requirements are also different. The broad component of the revenue gradient can leverage more powerful quarterly variation in the three average tax rates and their corresponding bases, while distributional personal income tax data are available only at an annual frequency. Nevertheless, I use the same approach to recover the present-value semi-elasticities required by the revenue gradient.

As my baseline approach, I recover the semi-elasticities in three steps. In the first step, I estimate a Bayesian vector autoregression (BVAR) and use narrative tax measures to recover a tax policy impact matrix A_0 , which I propagate forward with Bayesian local projections (BLP). This two-stage approach was developed by Cloyne et al. (2022) and has three distinct advantages in this setting. First, the BVAR uses the joint covariance structure of the system, together with the narrative measures, to separate contemporaneous tax shocks when several tax rates move at the same time. Second, after these impact responses have been identified, BLP estimates each subsequent response directly rather than propagating a BVAR with a finite number of lags, reducing sensitivity to lag truncation bias. Finally, the postwar window is short and reform episodes are sparse, so Bayesian estimation both regularizes the estimates and produces joint posterior draws that carry uncertainty from identification through the response paths and resulting semi-elasticities.¹¹ In the second step, I discount and sum the resulting impulse responses of tax rates and tax bases. Finally, I use the complete matrix of tax rate responses to re-express the corresponding tax base responses as present-value semi-elasticities with respect to each tax rate while holding the others fixed. In the progressivity subsection, I describe one additional step that gives the semi-elasticities of each tax base with respect to the progressivity reform. The remainder of this section describes the identification and estimation approach in detail.

To recover A_0 , I estimate a BVAR with hierarchical priors following Giannone, Lenza, and Primiceri (2015), details of which are in Appendix C. Letting $\mathbf{Z}_t = (\boldsymbol{\tau}_t, \mathbf{B}_t)^\top$ denote the state vector, where $\boldsymbol{\tau}_t$ is a vector of tax rates and $\mathbf{B}_t$ is the corresponding vector of tax bases, the BVAR is

$$\mathbf{Z}_t = \mathbf{c} + \sum_{\ell=1}^p \Pi_\ell \mathbf{Z}_{t-\ell} + \mathbf{u}_t, \quad (22)$$

where $\mathbf{u}_t$ is the vector of reduced-form innovations. For simplicity, I suppress controls and any additional variables in the state vector; the application subsections describe these where relevant. I use narrative tax measures $\mathbf{m}_t$ to identify the tax policy component of the reduced-form inno-

11. The BVAR and BLP can also be replaced by estimating a proxy SVAR, as in Mertens and Ravn (2013b), and then propagating forward with frequentist local projection approaches. Alternatively, one can simply estimate a BVAR or proxy SVAR. The estimation section contains robustness with all such approaches, along with direct LP and LP-IV to guard against noninvertibility.

vations, so the identification argument is similar to Mertens and Ravn (2013b). The relationship between the reduced-form and structural shocks is

$$\mathbf{u}_t = A_0 \boldsymbol{\epsilon}_t + \mathbf{u}_t^{\text{other}}, \quad (23)$$

where $\boldsymbol{\epsilon}_t$ contains the K mutually orthogonal structural tax shocks, A_0 maps those shocks into the contemporaneous reduced-form innovations, and $\mathbf{u}_t^{\text{other}}$ collects the effects of the remaining structural shocks, $\boldsymbol{\epsilon}_t^{\text{other}}$. With K relevant tax shocks, the identifying restrictions are

$$E[\mathbf{m}_t \boldsymbol{\epsilon}_t^\top] = \Phi, \quad E[\mathbf{m}_t (\boldsymbol{\epsilon}_t^{\text{other}})^\top] = 0, \quad \text{rank}(\Phi) = K. \quad (24)$$

The first restriction requires the narrative measures to be relevant for the tax shocks, while the second requires them to be unrelated to any remaining structural disturbances, and the third requires Φ to have full rank. It follows that

$$E[\mathbf{u}_t \mathbf{m}_t^\top] = A_0 \Phi^\top,$$

where A_0 is the contemporaneous impact matrix for the K tax shocks. Together with the covariance matrix of $\mathbf{u}_t$, these restrictions recover this impact block up to an orthogonal rotation. For each BVAR draw, I select a particular full rank basis using a Cholesky decomposition. This basis is needed to construct the intermediate impulse responses, but its ordering has no effect on the final semi-elasticity matrix constructed below.

Having recovered A_0 , I estimate local projections at each horizon $h \geq 1$:

$$\mathbf{Z}_{t+h} = \mathbf{c}^{(h)} + \Psi_h \mathbf{Z}_t + \sum_{\ell=1}^{p-1} B_\ell^{(h)} \mathbf{Z}_{t-\ell} + \boldsymbol{\eta}_{t+h}^{(h)}. \quad (25)$$

Here, Ψ_h is the horizon-specific coefficient matrix on the impact-period state $\mathbf{Z}_t$, $B_\ell^{(h)}$ collects the coefficients on its lags, and $\boldsymbol{\eta}_{t+h}^{(h)}$ is the error term. The impact response is given by A_0 , while the response at each subsequent horizon is obtained by applying the corresponding local projection coefficient matrix to A_0 :

$$IRF_0 = A_0, \quad IRF_h = \Psi_h A_0 \quad \text{for } h \geq 1.$$

Each column of the impact matrix A_0 gives the contemporaneous shock of the endogenous system to one tax shock, while Ψ_h maps that response directly into horizon h . In each application, I trace the response paths over a five-year window. The resulting medium-run impulse responses are not yet the semi-elasticities required by the revenue gradient because each tax shock may move

several tax rates at once. That is the final step.

Before constructing the response matrices, I convert responses of the transformed tax variables into ordinary percentage point changes in tax rates using the corresponding local chain rule. Let D_h denote the resulting matrix of tax rate responses at horizon h , and let Y_h denote the corresponding matrix of log tax base responses. I construct their discounted sums over the five-year window,

$$D = \sum_{h=0}^H \beta^h D_h, \quad Y = \sum_{h=0}^H \beta^h Y_h. \quad (26)$$

If D has full rank, the matrix of present-value semi-elasticities is

$$E = YD^{-1}. \quad (27)$$

The matrix E contains the full set of semi-elasticities. Although the individual shock responses depend on the ordering of the shocks in the Cholesky decomposition, the resulting semi-elasticities do not. A common change of basis rotates D and Y in the same way and therefore cancels from equation (27).

With the common estimation and identification procedure in place, I now apply it to the two components of the revenue gradient. I begin with the quarterly system of broad personal, corporate, and tariff rates, and then use annual distributional data to recover the progressivity entry.

3.3 Broad Tax Instruments

I start by assessing the broad tax system, which is comprised of average tax rates on personal income, corporate profits, and imports.

Data and Shocks

For the broad tax application, the state vector is

$$\mathbf{Z}_t = \left[\log(1 - \tau_{P,t}), \log(1 - \tau_{C,t}), -\log(1 + \tau_{T,t}), \log B_{P,t}, \log B_{C,t}, \log B_{T,t} \right]^\top,$$

where $\tau_{i,t}$ is the average tax rate, $B_{i,t}$ is the corresponding tax base, and t indexes quarters. The variables in the state vector are constructed using the National Income and Product Accounts as described in Appendix B. The personal and corporate tax variables exactly correspond to those used in Mertens and Ravn (2013a) and Cloyne et al. (2022), while the tariff rate is customs duties divided by total imports of goods, where the latter is the base. Although these definitions are

standard, there is one important shortcoming given the nature of the exercise. In the national accounts, corporate profits include both C-corporations and S-corporations, and they are not typically taxed the same. As a result, some cross-base income shifting, which is a natural response to tax changes, may not show up if S- and C-corps are treated as being in the same base. I discuss the implications of this choice in more detail later in the section. As a robustness check, I also later use the Statistics of Income to reallocate income from the corporate tax base to the personal tax base, with corresponding details about construction in Appendix B.1. The sample covers 1947Q1–2019Q4 with four lags.¹²

As described in the previous subsection, I use narrative shocks to recover the broad tax semi-elasticities. I rely on two separate sources to supply the narrative shocks.

First, I use the tax shock series for personal and corporate shocks which was initially developed by Mertens and Ravn (2013a) and later extended by Cloyne et al. (2022) through the 2017 Tax Cuts and Jobs Act.¹³ The shocks come from reading the narrative record around the time of each major tax reform to understand the motivations of policymakers. If the reforms seem to be in pursuit of objectives unrelated to the business cycle, like deficit reduction, ideological commitments, or growth, then they get classified as exogenous. To ensure that there are no anticipation effects, I exclude reforms implemented more than 90 days after announcement.¹⁴ Each reform receives a quantitative measure equal to the projected revenue impact scaled by the lagged tax base, as announced at the time of enactment. This scaling delivers an approximation to the change in the average tax rate on that instrument. The resulting series are denoted m_t^P for personal income taxes and m_t^C for corporate income taxes. Cloyne et al. (2022) argues that these instruments are strong and I provide additional diagnostics in the appendix in Table F.1. I find that the personal instrument is strong, but both the corporate and the tariff instruments are borderline relative to the critical threshold.

To get exogenous variation in tariffs, I rely on the narrative series developed by Besten et al. (2026). They classify postwar tariff changes as exogenous when the legislative record indicates motivations unrelated to contemporaneous economic conditions. These are somewhat sparse and concentrated around the GATT negotiating rounds and the 2018–2019 tariff increases. Because of that, I make two modifications to improve instrument strength. First, I assign each shock

12. Note that the average tax rate is the object required by theory since revenue from each base is $R_i = \tau_i B_i$.

13. Mertens and Ravn (2013a) decompose the Romer and Romer (2010) aggregate tax shock series into corporate and personal tax shocks.

14. This mitigates anticipation effects that operate between enactment and implementation, but does not address anticipation that precedes enactment itself. For example, it was widely expected that President Reagan would pursue tax reform well before specific legislation was introduced, and similarly that President Trump would push for a large corporate tax cut in 2017. If agents adjusted behavior in advance of these reforms, the narrative shock—which is dated at enactment—would be preceded by movements in the tax bases. Figure F.1 provides a direct test by estimating impulse responses.

to the implementation quarter rather than distributing it across adjacent years. Second, I replace their sign-only coding with magnitudes, using tariff rate changes from Bown and Irwin (2015) for the pre-Trump period and from Amiti, Redding, and Weinstein (2020) for the Trump-era tariffs. Appendix Figure B.1 plots the resulting shock series alongside the Mertens-Ravn shocks. I denote this instrument as m_t^T .

Present-Value Semi-Elasticities

The revenue gradient for each instrument depends on the full matrix of present-value semi-elasticities. Table 1 reports the corresponding 3×3 matrix under the Bayesian local projection estimator. Each column gives the response to a one percentage-point increase in the named tax rate, holding the other two rates fixed. The diagonal tells a familiar story: the median own semi-elasticity of the personal base is small, at roughly -0.4 , though imprecisely estimated; the corporate base falls by roughly 3.4 percent following a one percentage-point corporate tax increase; and the tariff base falls by roughly 14 percent following a one percentage-point tariff increase. The most directly comparable estimate comes from Besten et al. (2026). In their postwar sample, a one percentage-point tariff increase reduces real imports by approximately 10 percent on impact, with the response persisting at 5–6 percent through year five. That is a similar order of magnitude as semi-elasticity reported here.

However, the off-diagonal pattern is striking. The corporate base is highly responsive to both personal tax rate increases and tariff increases; the former reduces the corporate base by roughly 11 percent, while the latter reduces it by roughly 5 percent. By contrast, a corporate tax increase raises the personal tax base by roughly 0.4 percent, and the three remaining cross-semi-elasticities are imprecisely estimated. The clearest cross-base pattern is therefore that the corporate base responds to changes elsewhere in the tax system. Whether those responses create large fiscal externalities depends on the tax rate and size of the responding base, which is what the revenue gradient incorporates.

Table 1: Present-Value Semi-Elasticities of Tax Bases

	τ_P	τ_C	τ_T
ε_P . (Personal Base)	-0.37 (-0.99, 0.28)	0.41 (0.14, 0.71)	0.83 (-0.09, 1.86)
ε_C . (Corporate Base)	-10.55 (-13.36, -8.20)	-3.40 (-4.45, -2.45)	-4.99 (-8.75, -1.72)
ε_T . (Tariff Base)	0.16 (-1.86, 2.07)	-0.58 (-1.41, 0.19)	-14.26 (-17.32, -11.91)

Note: Each semi-elasticity ε_{ij} is the cumulative discounted percent change in base i per percentage-point increase in rate j , holding the other two rates fixed. Within each draw the rate-response matrix D and base-response matrix Y are mapped into $E = YD^{-1}$, an elasticity with respect to the net-of-tax coordinate, and converted to the tax rate by $-E_{ij}/(1 - \tau_j)$ for the income taxes and $-E_{ij}/(1 + \tau_T)$ for tariffs, evaluated at sample-average tax rates. Parentheses report 90% intervals. The baseline is estimated by Bayesian local projection; Table D.1 reports alternative methods.

The large cross-semi-elasticity from personal tax rates to corporate income survives across three different robustness checks. First, Appendix Table D.1 shows that when I maintain the same six-variable system, the large cross-semi-elasticity from personal tax rates to corporate income remains negative in the point estimates across methods, although its magnitude varies. Some methods, including the proxy SVAR, feature wide confidence intervals. Second, Appendix Table D.3 shows that this large cross-semi-elasticity is largely unchanged when reallocating S-corporate profits to the personal tax base. That is surprising. Since S-corporations are passthrough entities, they are taxed at the personal tax rate. When personal taxes fall, S-corporate activity may expand and appear in the BEA’s measured corporate tax base even though it does not generate corporate revenue. In principle, this channel should be quantitatively important because S-corporations grew from a negligible share of business income to roughly half between the Tax Reform Act of 1986 (TRA86) and the 2010s (Dyrda and Pugsley 2024; Smith et al. 2022). That growth stemmed in part from TRA86 lowering the top personal rate below the corporate rate for the first time. Nevertheless, this particular source of reclassification appears unimportant for the large cross-semi-elasticity. Next, Appendix Table D.2 shows that the large cross-semi-elasticity also survives when the system expands to nine variables including federal purchases, federal debt, and GDP. Finally, Appendix Figure D.2 shows that few reforms are influential for the magnitudes in the resulting semi-elasticity matrix, but the Tax Reform Act of 1986 is particularly pivotal for corporate taxes.

Rationalizing the Cross-Base Response. Given the survival of the large response of the corporate base to personal taxes across various robustness checks, it is important to rationalize. One potential explanation, which I explore in greater detail in Appendix D.1, originates in operating leverage. Corporate profits are a residual claim on revenue minus costs. Some of these costs, including capital payments, debt service, leases, and overhead, adjust slowly or are entirely predetermined. When the cost structure of firms is slow to adjust to shocks, profits absorb more of the change in revenue. Consider a firm with \$100 in revenue, \$85 in costs, and \$15 in profits. If revenue rises 1 percent to \$101 but costs do not immediately adjust, profits rise to \$16. In other words, with sluggish adjustment, profits rise by $1/0.15 \approx 6.7$ percent in response to a small shock.

For this calculation, I return to the corresponding net-of-tax elasticities, since operating leverage maps proportional changes in output into proportional changes in profits. The same arithmetic applies to the aggregate corporate sector. In the national accounts, corporate value added decomposes into wages wL , profits π , and other costs F that are plausibly predetermined. If output raises by dY and labor incomes goes up proportionally to its labor share (γdY , where γ is the labor share), then the remainder $1 - \gamma$ must accrue entirely to corporate profits. Because profits are only a share s_π of output, this translates into an elasticity of $(1 - \gamma)/s_\pi$. This ratio is directly observable from the National Income and Product Accounts and ranges from approximately 1.9 to 3.4 over the postwar period (Appendix Figure D.3). Theoretically, assuming a frictionless labor market, the cross-elasticity is then the product of this leverage and the output response to personal tax cuts: $\varepsilon_{CP} \approx [(1 - \gamma)/s_\pi] \times \varepsilon_{YP}$. When I add GDP to the specification, the output elasticity is $\varepsilon_{YP} \approx 1.0$. Combining a postwar average leverage of 2.59 with this output elasticity predicts $\varepsilon_{CP} \approx 2.5$, roughly a third of the baseline estimate of 8.89. This prediction assumes frictionless labor, and it is a lower bound if wages adjust more slowly than output. When wages are also slow to adjust, a larger share of output gains accrues to profits, raising effective leverage above the frictionless-labor benchmark. Table 2 extends the decomposition to the PIT–CIT elasticity matrix. The frictionless-labor benchmark tends to underpredict corporate entries and overpredict personal entries, which is consistent with what one would expect when wages are slow to adjust.

Table 2: Frictionless-Labor Predictions vs. Data

	Frictionless Labor	Preferred BLP	Macro Controls
ε_{CC}	$2.59 \times 0.9 \approx 2.2$	2.50	4.13
ε_{CP}	$2.59 \times 1.0 \approx 2.5$	8.80	8.96
ε_{PC}	$1 \times 0.9 \approx 0.9$	-0.30	0.52
ε_{PP}	$1 \times 1.0 \approx 1.0$	0.31	0.10

Note: The first column applies operating leverage $(1 - \gamma)/s_\pi$, whose NIPA mean is 2.59, to the draw-level GDP elasticities of the Macro Controls specification; personal-base pass-through is one under frictionless labor. The two data columns are posterior medians of five-year present-value pure-rate responses from BLPs that jointly estimate the personal, corporate, and tariff rates and bases. Macro Controls adds GDP, government spending, and federal debt. This table is stated in net-of-tax elasticities rather than the tax-rate semi-elasticities reported elsewhere, so that the columns line up with the primitive BLP estimates they are built from. The leverage identity itself is unaffected by the choice of coordinate: it multiplies the output response by $(1 - \gamma)/s_\pi$, and rewriting both sides as semi-elasticities divides each by the same factor.

3.4 The Personal Tax Progressivity Gradient

The previous subsection estimates the broad entries of the revenue gradient by treating the personal income tax as a single instrument. I now estimate the remaining entry of the revenue gradient, which captures how tax revenue responds to a change in the progressivity of personal taxes. In particular, I consider a reform that raises the marginal tax rate paid by the top one percent of the income distribution by one percentage point relative to the lower 99 percent. This approach to nonlinearity and progressivity is intentionally stylized.

Data and Shocks

This specification begins with a six-variable annual BVAR:

$$\mathbf{Z}_t = [\log(1 - \tau_{L,t}), \log(1 - \tau_{H,t}), \log B_{L,t}, \log B_{H,t}, \log B_{C,t}, \log B_{T,t}]^\top,$$

where L denotes the lower 99 percent and H denotes the top one percent. The rates $\tau_{L,t}$ and $\tau_{H,t}$ are the combined marginal income and payroll tax rates faced by each group. $B_{L,t}$ and $B_{H,t}$ are the corresponding real average market income measures excluding capital gains. The corporate and tariff bases are the same as in the quarterly analysis in the previous section, but aggregated to an annual frequency. Including those bases allows me to trace the effects of changes in personal tax rates on the corporate and tariff bases, which is necessary to recover the semi-elasticities entering

the revenue gradient.

The group-specific personal tax rates and income bases come from Mertens and Montiel Olea (2018). Their original series end in 2012, which I extend through 2023 using Statistics of Income data. Appendix B.2 describes the extension and the construction of the lower-99 and top-1 series. The annual sample covers 1947–2023. I use one annual lag and include annual corporate and tariff narrative measures as contemporaneous controls in both stages of the estimation. I use the annual discount factor $\beta_A = \beta^4$ and trace responses from impact through year four, yielding the same five-year window as in the quarterly analysis. The personal narrative measures also come from Mertens and Montiel Olea (2018). For each tax reform they classify as exogenous, they calculate the change in the average marginal tax rate faced by each income group. I use their measures for the lower 99 percent and the top 1 percent, which Figure B.4 displays together with my extension to the 2017 Tax Cuts and Jobs Act.

Because reforms almost always change both personal tax rates at once, I use the two narrative measures jointly rather than treating them as separate experiments in which one rate changes while the other remains fixed. I first purge from each personal narrative measure the part associated with contemporaneous corporate and tariff narrative measures. The procedure in the previous subsection then yields a 4×2 matrix of semi-elasticities E . The first column of E gives the response of every tax base to a one-percentage-point increase in the lower-99 rate while holding the top-1 rate fixed. The second gives the corresponding response to the top-1 rate.

However, there is one additional step required to isolate a pure progressivity reform because a change in one rate also changes the aggregate personal tax burden. I therefore combine the two rate changes such that the top-1-minus-lower-99 rate gap rises by one percentage point while mechanical personal tax payments remain fixed. To do that, let $b_{L,t}$ and $b_{H,t}$ denote the amount of income exposed to a marginal change in each tax rate. I construct these exposures from group income and the official ex ante liability scores in the MMO reform data, as described in Appendix B.2. Define the top-1 share of total exposure as

$$s_{H,t} = \frac{b_{H,t}}{b_{L,t} + b_{H,t}}.$$

In that case, a one-percentage-point increase in progressivity changes the two rates in the direction

$$\mathbf{h}_t = \begin{pmatrix} -s_{H,t} \\ 1 - s_{H,t} \end{pmatrix}. \quad (28)$$

The group with the larger initial exposure receives the smaller rate change. The rate gap therefore

risers by one percentage point. By construction, this satisfies

$$b_{L,t}h_{L,t} + b_{H,t}h_{H,t} = 0,$$

so initial personal tax payments remain fixed. Applying this reform to the two pure-rate responses gives

$$\boldsymbol{\varepsilon}_{d,t} = E\mathbf{h}_t = \begin{pmatrix} \varepsilon_{Ld,t} & \varepsilon_{Hd,t} & \varepsilon_{Cd,t} & \varepsilon_{Td,t} \end{pmatrix}^\top. \quad (29)$$

These are the present-value semi-elasticities of the four tax bases with respect to a one-percentage-point increase in progressivity.

Present-Value Semi-Elasticities

Table 3 reports the present-value semi-elasticities of each tax base with respect to progressivity. While each semi-elasticity is computed date by date, I report its 1948–2023 average within each posterior draw before reporting the posterior distribution across draws. I find that a one-percentage-point increase in progressivity reduces the tax base of the top income percentile by 2.2 percent, with a credible interval well clear of zero. The response of the bottom-99 percent, by contrast, is imprecisely estimated. However, the reform does increase the corporate tax base by roughly 4.0%, likely reflecting income shifting, and it also increases the tariff base by roughly 7.6%. Both of those cross-base responses are significantly different from zero.

Table 3: Five-Year Present-Value Semi-Elasticities with Respect to Personal-Tax Progressivity

	Bottom 99	Top 1	Corporate base	Tariff base
BLP	-0.15	-2.28	4.12	7.97
	(-0.65, 0.42)	(-3.80, -1.08)	(1.83, 12.63)	(4.64, 21.37)

Note: Posterior medians with 90-percent credible intervals in parentheses. Entries are the percentage change in the discounted five-year response of each tax base to a one-percentage-point increase in the top-1-minus-bottom-99 average marginal tax-rate gap, holding initial personal tax payments fixed, evaluated draw by draw at each year’s observed rates and exposure weights and averaged across accounting years. The personal block has two groups, identified by two narrative instruments over nine reform years, conditioning on the annual corporate and tariff narrative measures.

The appendix addresses three concrete concerns about this interpretation. First, the estimated responses could depend on the BLP propagation procedure or on the recovery of tax shocks from reduced-form innovations. To address that, Appendix Table E.1 repeats the exercise in the main table across different methods. Standard, smooth, and bias-corrected local projections, which propagate a frequentist impact matrix estimated from a proxy SVAR, produce a similar pattern,

but are less precise. The BVAR yields a corporate response near zero and a substantially smaller tariff response, while the proxy SVAR estimates are too imprecise to adjudicate. Meanwhile, adding direct LP–IV, which does not rely on the same invertibility condition, yields intervals too wide to independently confirm the baseline pattern. Second, the spillovers could reflect the accounting of passthrough business income or of omitted aggregate conditions. Appendix Table E.3 shows that adding GDP, government spending, and debt leaves both spillovers positive, although smaller than in the baseline, while Appendix Table E.4 shows that reallocating S-corporation income strengthens the estimated corporate response and leaves the tariff response similar, while mildly changing the estimated responses within the personal tax base. Third, there are not many reforms in the Mertens and Montiel Olea (2018) shocks. Appendix Figure E.1 re-estimates the semi-elasticities after omitting each reform in turn. The resulting semi-elasticities are similar when leaving reforms out one at a time.


The evidence here is plainly less robust than for the quarterly broad tax estimates against methodological choices, but it nevertheless holds up against changing the specification, passthrough accounting, and the leave-one-out estimation.

4 The Revenue Gradient over the Postwar Era

The previous section estimated the inputs to the components of the revenue gradient $\mathbf{r}_t = (\mathbf{r}_{A,t}^\top, r_{d,t})^\top$. This subsection illustrates their combination in Figure 2, which shows the semi-elasticity of total tax revenue with respect to each instrument. Each panel shows the consolidated revenue semi-elasticity along with 68% and 90% credible intervals.

Positive values mean that increasing the tax rate—or increasing progressivity—raises revenue, while a negative value means that the instrument is on the wrong side of the Laffer curve. The personal tax revenue semi-elasticity is the highest of any instrument by a considerable degree and it remains significantly positive throughout the sample. By contrast, the corporate tax and tariff instruments initially have negative point estimates, indicating that in the early postwar era higher taxes of either kind would have reduced revenue, though not significantly so. However, by the 1970s, the corporate tax instrument would have raised revenue. That is similarly true for the tariff instrument, but only for the 68% credible interval. In contrast to the broad instruments, which all saw their semi-elasticities rise over time, the progressivity semi-elasticity of consolidated revenue declined over time and the point estimate has been roughly zero for several decades. In other words, raising the top-1% tax rate relative to the bottom-99 by one percentage point would not raise revenue.


Figure 2: Revenue Semi-Elasticities of the Four Reform Margins



Note: Each panel plots the semi-elasticity of consolidated revenue with respect to a one percentage point reform. Broad-rate changes hold the other broad rates fixed. The progressivity coordinate widens the top-1-minus-lower-99 rate gap while holding initial personal tax payments fixed. Solid lines are posterior medians; dark and light shaded regions are 68-percent and 90-percent credible intervals.

All four trends in Figure 2 are largely explained by the decline in the corporate tax rate. Figure 3 shows that if the corporate tax rate were held at its 1948 level—roughly 40%—then the personal tax semi-elasticity would have been stagnant until the 1990s and then declined, the increases in the corporate tax and tariff semi-elasticities would have been considerably less steep, and the progressivity semi-elasticity would have hardly budged. Hence, much of the movement in the components of the revenue gradient are explained by the large cross-partial between each of the instruments with the corporate tax base, which the previous subsections have shown is robust. Indeed, they suggest that raising revenue via broad tax instruments is considerably easier now on the margin because corporate tax rates have declined so much since the 1940s.

Figure 3: The Corporate Tax Rate Revalues the Revenue Margins



Note: Solid lines show posterior-median revenue semi-elasticities. Dashed lines recompute each margin after holding the average corporate tax rate at its 1948 level. All other tax rates, tax bases, and estimated behavioral responses follow their historical values. Each series is expressed relative to actual total revenue. Panels use different vertical scales.

While this section can sign the revenue gradient, it has nothing to say about alignment with the local Ramsey optimum. The next section combines these objects to evaluate Ramsey alignment in the postwar era for progressively larger reform spaces.

5 Ramsey Alignment and Local Reform


The previous section constructed four components of the revenue gradient for the federal tax system. I now combine them with the welfare gradient to evaluate the proximity of the system to the local Ramsey optimum. Since the annual and quarterly blocks are estimated separately, I treat their posteriors as independent. Consequently, I evaluate every alignment statistic on the full Cartesian product of the two sets of draws, so its posterior reflects uncertainty in both. It is likely that this is conservative because the two blocks share data and are similar policy experiments, so the errors are likely positively correlated.¹⁵

Figure 4 shows how Ramsey alignment (captured by $\cos \theta$) and the gains to reform (captured by $\sin \theta$) change as more ingredients are added to the tax system. Under the equal incidence benchmark, which I use throughout the main results, the personal–corporate tax system is closely aligned with the Ramsey optimum throughout the sample, showing only modest misalignment before the 1980s. On its own, that would suggest little room for reform post-1980. However, adding in tariffs substantially raises the degree of misalignment early in the period, and the right panel indicates that there are significant gains to reform left when only evaluating the broad tax mix. Finally, adding in progressivity, which is the red line and associated 90% credible interval, significantly raises both the degree of misalignment and the uncertainty around it throughout the postwar era. Between 1950 and 1990, the median posterior draw increases of $\cos \theta$ goes from 0.2 to roughly 0.6, but stays around that point from 1990 onward.¹⁶ That contrasts strongly with the broad instrument alignment, which is close to one with a tight credible interval. Importantly, the contrast also emerges despite equal incidence, which assigns zero weight to progressivity, so the entirety of the effect comes through the effect of progressivity on revenue.

15. Because the cosine rises in the broad revenue gradient and falls in the progressivity gradient, positively dependent errors would narrow this posterior rather than widen it, so the reported intervals are conservative under the plausible violation; re-coupling the two posteriors at the extreme admissible dependence moves the 90 percent width by less than ten percent after 1970 and by eighteen percent in 1948.

16. Since the interval is wide, Appendix Figure G.1 plots the probability that alignment has improved compared to the beginning of the period, finding that it exceeds 0.9 by 1990 and remains stable.


Figure 4: Ramsey Alignment in Nested Reform Spaces



Note: The left panel plots $\cos \theta$ and the right plots $\sin \theta$. The black series permits changes in personal and corporate tax rates; the blue series additionally permits changes in tariffs; and the red series additionally admits progressivity. The blue series is the broad tax space. Shaded regions are 90-percent posterior credible intervals.

Figure 4 establishes that revenue-neutral reforms to the tax system existed throughout the postwar era, although the degree to which they would matter declines over time. Next, I apply the projection result from equation (9) to show which ones would have been most welfare-improving over time in Figure 5. Since the magnitude of reform is normalized to the unit vector, each panel should be interpreted as its contribution to that along with associated posterior uncertainty. Importantly, the uncertainty in each panel should be understood as increasing as alignment goes up because the space for obvious reforms goes away. From left, the planner would optimally have significantly cut taxes on personal income throughout the postwar era. That cut was largest in the early period and would become more mild by the 1980s. That was similarly true for both tariffs and corporate taxes until the 1980s, after which both become indistinguishable from zero. By contrast, the rightmost panel indicates that, within personal taxes, the tax system should have become more progressive. Pointwise, that held true until the 1980s. Since the planner places no weight on progressivity under equal incidence, the progressivity coordinate of the optimal reform is proportional to the progressivity semi-elasticity of consolidated revenue. That coordinate therefore collapses toward zero, and its sign stops being identified, once the semi-elasticity becomes indistinguishable from zero. Following the discussion in the previous section, all of the trends in Figure 5 are largely explicable with the falling corporate tax rate, which is the most elastic of all of the tax bases.

Figure 5: Coordinates of the Locally Improving Revenue-Neutral Reform



Note: The figure plots the coordinates of the best local revenue-neutral reform, $-P_{\perp r} \mathbf{w} / \|P_{\perp r} \mathbf{w}\|$, in the stated policy metric. A negative broad-rate coordinate denotes a tax cut; a positive progressivity coordinate raises the top-1-minus-lower-99 rate gap. At each posterior draw, the four coordinates have unit Euclidean norm. Solid lines are coordinatewise posterior medians; dark and light shaded regions are 68-percent and 90-percent credible intervals. Consequently, the four median series need not themselves form a unit-norm reform vector.


Figures 4 and 5 do not establish that every feature of the tax schedule is misaligned. Rather, the takeaway, particularly since the approach to progressivity is stylized, is that alignment among broad tax rates is not sufficient to establish alignment in a policy space that also allows for progressivity. The remaining question is why admitting progressivity changes the alignment statistic even when the broad tax rates are closely aligned. Under equal incidence, and whenever $\cos \theta_A \neq 0$, Proposition 2 implies that the ratio of alignments between the full system and the broad tax system is

$$\frac{\cos \theta}{\cos \theta_A} = \frac{\|\mathbf{r}_A\|}{\sqrt{\|\mathbf{r}_A\|^2 + r_S^2}},$$

where $|r_S| = |r_d|/\sqrt{m_d}$ is the progressivity component of the gradient in our application, expressed in the same broad personal-rate units as the entries of $\mathbf{r}_A$. The metric m_d is the squared length of the progressivity direction, given in equation (A.34). The attenuation factor on the right is well defined whenever $\mathbf{r}_A \neq 0$, so it is the object plotted below. The ratio between the two is the share of broad alignment that remains after admitting the within-personal revenue margin. Moreover, as long as r_S is nonzero, then the revenue gradient moves away from the revenue gradient without adding any tilt to the welfare gradient. Since, in practice, r_S is distinguishable from zero before the 1980s, that substantially increased the degree of misalignment relative to the broad tax system. Figure 6 reports the degree of attenuation draw by draw. Early in the post-war period, relatively little broad alignment survives one progressivity becomes another margin for the planner. Later on, it rises, although with a wide confidence interval. The uncertainty is large because the progressivity semi-elasticity of revenue goes to zero as the corporate tax rate

declines.

Figure 6: The Attenuation of Broad Alignment from Progressivity



Note: The figure plots $\cos \theta / \cos \theta_A$, where θ_A is alignment in the three-instrument personal, corporate, and tariff reform space and θ additionally admits progressivity. The solid line is the posterior median. Shaded regions are 68% and 90% credible intervals.

Taken together, the results point toward a system closer to the Ramsey optimum than what we observed in the early postwar period. Nevertheless, this conclusion is firmer for the broad tax mix than it is for the progressivity component. The next section examines the robustness of this conclusion.

5.1 Robustness

This section studies the robustness of the alignment conclusion to methodological concerns, specification choices, welfare assumptions, and other potential threats.

I first vary how the revenue gradient is estimated methodologically. Appendix Figures G.2, G.3, and G.6 compare the baseline Bayesian local projection with alternative estimators, direct IV methods, and different horizons for the present-value elasticities. These exercises study how other methodological choices influence the main result, as well as if the result is specific to a particular choice of horizon. I find that these exercises do not produce identical levels of alignment. In particular, the full-system cosine is sensitive to both the method used to estimate the response matrix and the horizon over which those responses are valued. While the LP methods and the proxy SVAR agree on the rise in alignment, the BVAR and LP-IV methods are considerably flatter

and lie beneath the posterior median BLP. Switching from a five- to a three-year elasticity flattens the alignment profile, while switching to seven-year elasticities produces a less dramatic rise in alignment.

I also investigate whether my stylized approach to measuring progressivity matters. Mertens and Montiel Olea (2018) include shocks for four income groups in their analysis, so I can split the tax schedule into up to three kinks. I therefore redo the analysis for two, three, and four income groups. For each partition, I reconstruct r_d from the corresponding semi-elasticities and recompute alignment in the resulting complete reform space. Appendix Figure G.8 shows that moving from two to three and then four income groups progressively lowers the cosine and raises the scope for revenue-neutral reform. A finer partition admits additional changes in the shape of the personal tax schedule, and those directions reveal revenue effects that are hidden by the top-1-versus-lower-99 comparison. However, splitting the tax schedule up into beyond one kink substantially reduces the strength of the first stage, which makes it difficult to interpret the estimates for additional partitions.

Separately, one may worry about if the broad personal tax variation is contaminated by changes in progressivity. The quarterly system does not hold it fixed, so it is possible that there is some contamination. Figure G.4 re-estimates the broad block using the component of the personal narrative measure orthogonal to the progressivity direction. Alignment rises at every date, indicating that this contamination biases the baseline alignment estimates downward rather than upward, but the difference is minimal.

Next, I check how important it is that the corporate income measured by the BEA includes S-corporations. Appendix Figure G.7 shows the path of alignment when moving S-corp income to the personal bucket, and I find that the path of alignment is unchanged, but the level is shifted down.

Additionally, I investigate whether the assumed equal incidence welfare weights in the benchmark alter the resulting conclusions about the path of alignment. Appendix Figure G.9 varies the social value assigned to burdens borne by the top 1 percent relative to the lower 99 percent and, separately, the value assigned to the corporate tax burden relative to the personal tax burden. Alignment is especially sensitive to the within-personal comparison: changing the relative weight on the top 1 percent materially changes both the level and the historical path of the cosine. However, this sensitivity is generally only relevant if the planner values the 1-percent above the 99-percent, which may not be empirically relevant given what is typically estimated in inverse welfare exercises (Hendren 2020). If personal income is, on average, valued 10 times more as corporate income, then the alignment path does not change. However, if corporate income is valued 10 times more than personal income, then the measured alignment path does not change, but the level shifts down moderately. In similar exercises, I check whether altering assumptions

on foreign ownership of corporations (Figure I.2) and the incidence of tariffs (Figure I.3) matter. They do not.

Finally, Appendix H considers tax reform when the responses of debt and government spending must also remain fixed. Empirically, reforms like the 2017 Tax Cuts and Jobs Act were deficit-financed. Given that, it is empirically true that taxes must rise at some point to satisfy the government budget constraint. I show in the Appendix that we can expand the reform problem to include constraints on debt and government spending and, by using their responses over time to reforms, can constrain the set of reforms to only span those which hold fixed debt and government spending. Imposing these additional restrictions substantially reduces the scope for welfare-improving reform, although the postwar improvement remains visible. Nevertheless, because this expanded scope mechanically reduces the feasible reform space, it should not be considered an alternative measure of unconstrained Ramsey alignment.

The robustness exercises show that the baseline is sensitive to a number of assumptions. That is, in part, a consequence of estimating macro elasticities. As a result, the baseline should be interpreted as a five-year, medium-run benchmark, rather than as a specification-free estimate of the distance between the tax system and the Ramsey optimum. Nevertheless, taken together, the exercises support treating the baseline as a transparent empirical benchmark for how to implement Ramsey alignment and evaluate the tax system as a whole across both linear and nonlinear dimensions.

6 Conclusion

This paper uses tools from the literature on the marginal value of public funds to answer old questions about the direction of optimal reform and develops new sufficient statistics for the alignment of a tax system with the local Ramsey optimum. The projection of the welfare gradient onto the revenue-neutral reform space gives the best local reform direction, while the angle between the welfare and revenue gradients measures alignment. These statistics are valid in many economic settings, including in general equilibrium. I show how to apply the framework both to alignment across broad tax instruments, which are relevant in the macro literature on optimal taxation (Straub and Werning 2020), and to nonlinear instruments. Under equal incidence, a change in the distribution of payments within an instrument may have no direct first-order welfare effect and still matter through its effect on revenue. The necessary ingredients are a welfare gradient, constructed from tax bases and incidence-adjusted welfare weights, and a revenue gradient.

Empirically, I construct the revenue gradient for four margins of the federal tax system in the United States. Because it is appropriate in this context to account for spillovers across bases

and over time, I estimate it using time series methods with established instruments from the literature on tax and tariff shocks (Besten et al. 2026; Mertens and Ravn 2013a; Mertens and Montiel Olea 2018). The resulting estimates yield what can be interpreted as general equilibrium semi-elasticities of consolidated revenue for corporate taxes, personal taxes, tariffs, and a top-1-versus-lower-99 measure of progressivity. There are large cross-base effects, particularly from the other instruments to corporate profits. Those cross-base effects drive much of the variation in both the time series of revenue gradients and the resulting measures of alignment. That's because the corporate tax rate has declined considerably in the postwar era. Combining these objects, I find that the broad tax instruments are now close to alignment under the baseline specification. Adding progressivity lowers measured alignment and introduces considerably more uncertainty, although the posterior-median revenue effect of progressivity itself has been close to zero in recent decades.

While the application is specific to a particular problem, the framework is more general. The same gradients could be constructed for other settings and larger sets of taxes. For example, one could estimate the revenue gradient within tariffs using a full matrix of product-specific elasticities, including spillovers to other tax bases. Together with the corresponding welfare gradient, this would yield a first-order notion of the alignment of the tariff schedule with the local Ramsey condition. One could also incorporate state and local taxes into the system, which would likely change my own estimates of alignment. The framework is therefore broadly applicable provided that the welfare and revenue gradients can be reliably constructed for the same policy space.

References

- Ahmad, Ehtisham, and Nicholas Stern. 1984. "The theory of reform and indian indirect taxes." *Journal of Public Economics* 25, no. 3 (December): 259–298. ISSN: 00472727. [https://doi.org/10.1016/0047-2727\(84\)90057-4](https://doi.org/10.1016/0047-2727(84)90057-4).
- Amiti, Mary, Stephen J. Redding, and David E. Weinstein. 2019. "The Impact of the 2018 Tariffs on Prices and Welfare." *Journal of Economic Perspectives* 33, no. 4 (November): 187–210. ISSN: 0895-3309. <https://doi.org/10.1257/jep.33.4.187>.
- . 2020. "Who's Paying for the US Tariffs? A Longer-Term Perspective." *AEA Papers and Proceedings* 110. ISSN: 2574-0768. <https://doi.org/10.1257/pandp.20201018>.
- Badel, Alejandro, Mark Huggett, and Wenlan Luo. 2020. "Taxing Top Earners: a Human Capital Perspective." *The Economic Journal* 130, no. 629 (July): 1200–1225. ISSN: 0013-0133. <https://doi.org/10.1093/ej/ueaa021>.
- Barnichon, Régis, and Christian Brownlees. 2019. "Impulse Response Estimation by Smooth Local Projections." *The Review of Economics and Statistics* 101, no. 3 (July): 522–530. ISSN: 0034-6535. <https://doi.org/10.1162/rest{\-}a{\-}00778>.
- Barnichon, Régis, and Geert Mesters. 2023. "A Sufficient Statistics Approach for Macro Policy." *American Economic Review* 113, no. 11 (November): 2809–2845. ISSN: 0002-8282. <https://doi.org/10.1257/aer.20220581>.
- Barro, Robert, and Brian Wheaton. 2019. *Taxes, Incorporation, and Productivity*. Technical report. Cambridge, MA: National Bureau of Economic Research. <https://doi.org/10.3386/w25508>. <http://www.nber.org/papers/w25508.pdf>.
- Bastani, Spencer. 2025. "The marginal value of public funds: a brief guide and application to tax policy." *International Tax and Public Finance* 32, no. 4 (August): 919–956. ISSN: 0927-5940. <https://doi.org/10.1007/s10797-024-09860-x>.
- Besten, Tamar den, Régis Barnichon, Diego Känzig, and Aayush Singh. 2026. *The Macroeconomic Effects of Tariffs: Insights from 180 Years of U.S. Trade Policy*. Technical report. Cambridge, MA: National Bureau of Economic Research, February. <https://doi.org/10.3386/w35102>.
- Bown, Chad, and Douglas Irwin. 2015. *The GATT's Starting Point: Tariff Levels circa 1947*. Technical report. Cambridge, MA: National Bureau of Economic Research, December. <https://doi.org/10.3386/w21782>.
- Caballero, Ricardo J., Eduardo M.R.A. Engel, and John Haltiwanger. 1997. "Aggregate Employment Dynamics: Building from Microeconomic Evidence." *American Economic Review* 87 (1): 115–137.
- Caravello, Tomás, Alisdair McKay, and Christian Wolf. 2024. *Evaluating Monetary Policy Counterfactuals: (When) Do We Need Structural Models?* Technical report. Cambridge, MA: National Bureau of Economic Research, September. <https://doi.org/10.3386/w32988>.
- Cavallo, Alberto, Gita Gopinath, Brent Neiman, and Jenny Tang. 2021. "Tariff Pass-Through at the Border and at the Store: Evidence from US Trade Policy." *American Economic Review: Insights* 3, no. 1 (March): 19–34. ISSN: 2640-205X. <https://doi.org/10.1257/aeri.20190536>.
- Chamley, Christophe. 1986. "Optimal Taxation of Capital Income in General Equilibrium with Infinite Lives." *Econometrica* 54 (3): 607–622.
- Chari, V.V., Juan Pablo Nicolini, and Pedro Teles. 2020. "Optimal capital taxation revisited." *Journal of Monetary Economics* 116 (December): 147–165. ISSN: 03043932. <https://doi.org/10.1016/j.jmoneco.2019.09.015>.
- Chetty, Raj. 2009. "Sufficient Statistics for Welfare Analysis: A Bridge Between Structural and Reduced-Form Methods." *Annual Review of Economics* 1, no. 1 (September): 451–488. ISSN: 1941-1383. <https://doi.org/10.1146/annurev.economics.050708.142910>.
- Christiano, Lawrence J., and Terry J. Fitzgerald. 2003. "The Band Pass Filter." *International Economic Review* 44, no. 2 (May): 435–465. ISSN: 0020-6598. <https://doi.org/10.1111/1468-2354.t01-1-00076>.

- Cloyne, James, Joseba Martinez, Haroon Mumtaz, and Paolo Surico. 2022. *Short-Term Tax Cuts, Long-Term Stimulus*. Technical report. Cambridge, MA: National Bureau of Economic Research, July. <https://doi.org/10.3386/w30246>.
- Conesa, Juan Carlos, Sagiri Kitao, and Dirk Krueger. 2009. "Taxing Capital? Not a Bad Idea After All!" *American Economic Review* 99, no. 1 (February): 25–48. ISSN: 0002-8282. <https://doi.org/10.1257/aer.99.1.25>.
- Diewert, W.E. 1978. "Optimal tax perturbations." *Journal of Public Economics* 10, no. 2 (October): 139–177. ISSN: 00472727. [https://doi.org/10.1016/0047-2727\(78\)90034-8](https://doi.org/10.1016/0047-2727(78)90034-8).
- Dixit, Avinash. 1975. "Welfare effects of tax and price changes." *Journal of Public Economics* 4, no. 2 (February): 103–123. ISSN: 00472727. [https://doi.org/10.1016/0047-2727\(75\)90014-6](https://doi.org/10.1016/0047-2727(75)90014-6).
- . 1979. "Price changes and optimum taxation in a many-consumer economy." *Journal of Public Economics* 11, no. 2 (March): 143–157. ISSN: 00472727. [https://doi.org/10.1016/0047-2727\(79\)90001-X](https://doi.org/10.1016/0047-2727(79)90001-X).
- Doan, Thomas, Robert Litterman, and Christopher Sims. 1984. "Forecasting and conditional projection using realistic prior distributions." *Econometric Reviews* 3, no. 1 (January): 1–100. ISSN: 0747-4938. <https://doi.org/10.1080/07474938408800053>.
- Dyrda, Sebastian, and Benjamin Pugsley. 2024. "The Rise of Pass-Throughs: an Empirical Investigation." *The Economic Journal* 135, no. 665 (December): 387–403. ISSN: 0013-0133. <https://doi.org/10.1093/ej/ueae062>.
- Fajgelbaum, Pablo D, Pinelopi K Goldberg, Patrick J Kennedy, and Amit K Khandelwal. 2020. "The Return to Protectionism*." *The Quarterly Journal of Economics* 135, no. 1 (February): 1–55. ISSN: 0033-5533. <https://doi.org/10.1093/qje/qjz036>.
- Favero, Carlo, and Francesco Giavazzi. 2012. "Measuring Tax Multipliers: The Narrative Method in Fiscal VARs." *American Economic Journal: Economic Policy* 4, no. 2 (May): 69–94. ISSN: 1945-7731. <https://doi.org/10.1257/pol.4.2.69>. <https://pubs.aeaweb.org/doi/10.1257/pol.4.2.69>.
- Feldstein, Martin. 1976. "On the theory of tax reform." *Journal of Public Economics* 6, nos. 1-2 (July): 77–104. ISSN: 00472727. [https://doi.org/10.1016/0047-2727\(76\)90042-6](https://doi.org/10.1016/0047-2727(76)90042-6).
- Ferey, Antoine, Benjamin B. Lockwood, and Dmitry Taubinsky. 2024. "Sufficient Statistics for Nonlinear Tax Systems with General Across-Income Heterogeneity." *American Economic Review* 114, no. 10 (October): 3206–3249. ISSN: 0002-8282. <https://doi.org/10.1257/aer.20221053>.
- Finkelstein, Amy, and Nathaniel Hendren. 2020. "Welfare Analysis Meets Causal Inference." *Journal of Economic Perspectives* 34, no. 4 (November): 146–167. ISSN: 0895-3309. <https://doi.org/10.1257/jep.34.4.146>.
- Ganapati, Sharat, and Colin Hottman. 2026. *Did Foreigners Pay America's Tariffs? Quantity Discounts, Scale Economies and Incomplete Pass-Through*. Technical report. Cambridge, MA: National Bureau of Economic Research, February. <https://doi.org/10.3386/w34901>.
- Giannone, Domenico, Michele Lenza, and Giorgio E. Primiceri. 2015. "Prior Selection for Vector Autoregressions." *Review of Economics and Statistics* 97, no. 2 (May): 436–451. ISSN: 0034-6535. https://doi.org/10.1162/REST{a}_{_}00483.
- Guesnerie, Roger. 1977. "On the direction of tax reform." *Journal of Public Economics* 7, no. 2 (April): 179–202. ISSN: 00472727. [https://doi.org/10.1016/0047-2727\(77\)90045-7](https://doi.org/10.1016/0047-2727(77)90045-7).
- Hall, Robert E., and Dale Jorgenson. 1967. "Tax Policy and Investment Behavior." *American Economic Review* 57:391–414.
- Hendren, Nathaniel. 2016. "The Policy Elasticity." *Tax Policy and the Economy* 30, no. 1 (January): 51–89. ISSN: 0892-8649. <https://doi.org/10.1086/685593>.
- . 2020. "Measuring economic efficiency using inverse-optimum weights." *Journal of Public Economics* 187 (July): 104198. ISSN: 00472727. <https://doi.org/10.1016/j.jpubeco.2020.104198>.

- Hendren, Nathaniel, and Ben Sprung-Keyser. 2020. "A Unified Welfare Analysis of Government Policies." *The Quarterly Journal of Economics* 135, no. 3 (August): 1209–1318. ISSN: 0033-5533. <https://doi.org/10.1093/qje/qjaa006>.
- Herbst, Edward P., and Benjamin K. Johannsen. 2024. "Bias in local projections." *Journal of Econometrics* 240, no. 1 (March): 105655. ISSN: 03044076. <https://doi.org/10.1016/j.jeconom.2024.105655>.
- Kleven, Henrik, Claus Kreiner, Kristian Larsen, and Jakob Søgaaard. 2025. "Micro versus Macro Labor Supply Elasticities: The Role of Dynamic Returns to Effort." *American Economic Review* 115, no. 9 (September): 2849–2890. ISSN: 0002-8282. <https://doi.org/10.1257/aer.20240554>.
- Kleven, Henrik J. 2021. "Sufficient Statistics Revisited." *Annual Review of Economics* 13:515–538. <https://doi.org/10.1146/annurev-economics-060220-023547>.
- Koby, Yann, and Christian K. Wolf. 2020. "Aggregation in Heterogeneous-Firm Models: Theory and Measurement."
- Leeper, Eric M, Alexander W Richter, and Todd B Walker. 2012. "Quantitative Effects of Fiscal Foresight." *American Economic Journal: Economic Policy* 4, no. 2 (May): 115–144. ISSN: 1945-7731. <https://doi.org/10.1257/pol.4.2.115>. <https://pubs.aeaweb.org/doi/10.1257/pol.4.2.115>.
- Lockwood, Benjamin B., and Matthew Weinzierl. 2016. "Positive and normative judgments implicit in U.S. tax policy, and the costs of unequal growth and recessions." *Journal of Monetary Economics* 77 (February): 30–47. ISSN: 03043932. <https://doi.org/10.1016/j.jmoneco.2015.10.006>.
- Lucas, Robert E. 1990. "Supply-Side Economics: An Analytical Review." *Oxford Economic Papers* 42 (2): 293–316.
- Mayshar, Joram. 1990. "On measures of excess burden and their application." *Journal of Public Economics* 43, no. 3 (December): 263–289. ISSN: 00472727. [https://doi.org/10.1016/0047-2727\(90\)90001-X](https://doi.org/10.1016/0047-2727(90)90001-X).
- McKay, Alisdair, and Christian K. Wolf. 2023. "What Can Time-Series Regressions Tell Us About Policy Counterfactuals?" *Econometrica* 91 (5): 1695–1725. ISSN: 0012-9682. <https://doi.org/10.3982/ECTA21045>.
- Mertens, Karel, and José Luis Montiel Olea. 2018. "Marginal Tax Rates and Income: New Time Series Evidence*." *The Quarterly Journal of Economics* 133, no. 4 (November): 1803–1884. ISSN: 0033-5533. <https://doi.org/10.1093/qje/qjy008>.
- Mertens, Karel, and Morten O. Ravn. 2012. "Empirical evidence on the aggregate effects of anticipated and unanticipated US tax policy shocks." *American Economic Journal: Economic Policy* 103 (4): 145–181.
- . 2013a. "The Dynamic Effects of Personal and Corporate Income Tax Changes in the United States." *American Economic Review* 103, no. 4 (June): 1212–1247. ISSN: 0002-8282. <https://doi.org/10.1257/aer.103.4.1212>. <https://pubs.aeaweb.org/doi/10.1257/aer.103.4.1212>.
- . 2013b. "The Dynamic Effects of Personal and Corporate Income Taxes in the United States." *American Economic Review* 103:1212–1247.
- Primiceri, Giorgio E. 2005. "Time Varying Structural Vector Autoregressions and Monetary Policy." *The Review of Economic Studies* 72, no. 3 (July): 821–852. ISSN: 1467-937X. <https://doi.org/10.1111/j.1467-937X.2005.00353.x>.
- Ramey, Valerie A. 2016. "Macroeconomic Shocks and Their Propagation." In *Handbook of Macroeconomics*, 71–162. <https://doi.org/10.1016/bs.hesmac.2016.03.003>. <https://linkinghub.elsevier.com/retrieve/pii/S1574004816000045>.
- Romer, Christina D, and David H Romer. 2010. "The Macroeconomic Effects of Tax Changes: Estimates Based on a New Measure of Fiscal Shocks." *American Economic Review* 100 (3): 763–801.
- Saez, E. 2001. "Using Elasticities to Derive Optimal Income Tax Rates." *The Review of Economic Studies* 68, no. 1 (January): 205–229. ISSN: 0034-6527. <https://doi.org/10.1111/1467-937X.00166>.
- Saez, Emmanuel, Joel Slemrod, and Seth H Giertz. 2012. "The Elasticity of Taxable Income with Respect to Marginal Tax Rates: A Critical Review." *Journal of Economic Literature* 50, no. 1 (March): 3–50. ISSN: 0022-0515. <https://doi.org/10.1257/jel.50.1.3>.

- Saez, Emmanuel, and Stefanie Stantcheva. 2018. "A simpler theory of optimal capital taxation." *Journal of Public Economics* 162 (June): 120–142. ISSN: 00472727. <https://doi.org/10.1016/j.jpubeco.2017.10.004>.
- Saez, Emmanuel, and Gabriel Zucman. 2026. "Distributional tax analysis in theory and practice: Harberger meets Diamond-Mirrlees." *Journal of Public Economics* 259 (July): 105675. ISSN: 00472727. <https://doi.org/10.1016/j.jpubeco.2026.105675>.
- Sims, Christopher A. 1993. "A Nine-Variable Probabilistic Macroeconomic Forecasting Model." In *Business Cycles, Indicators and Forecasting*, edited by James H. Stock and Mark W. Watson. National Bureau of Economic Research.
- Smith, Matthew, Danny Yagan, Owen Zidar, and Eric Zwick. 2022. "The Rise of Pass-Throughs and the Decline of the Labor Share." *American Economic Review: Insights* 4, no. 3 (September): 323–340. ISSN: 2640-205X. <https://doi.org/10.1257/aeri.20210268>.
- Straub, Ludwig, and Iván Werning. 2020. "Positive Long-Run Capital Taxation: Chamley-Judd Revisited." *American Economic Review* 110, no. 1 (January): 86–119. ISSN: 0002-8282. <https://doi.org/10.1257/aer.20150210>.
- Tirole, J., and R. Guesnerie. 1981. "Tax reform from the gradient projection viewpoint." *Journal of Public Economics* 15, no. 3 (June): 275–293. ISSN: 00472727. [https://doi.org/10.1016/0047-2727\(81\)90012-8](https://doi.org/10.1016/0047-2727(81)90012-8).
- Vergara, Damian, and Dustin Swonder. 2025. "A Sufficient Statistics Approach to Optimal Corporate Taxes."
- Winberry, Thomas. 2021. "Lumpy Investment, Business Cycles, and Stimulus Policy." *American Economic Review* 111, no. 1 (January): 364–396. ISSN: 0002-8282. <https://doi.org/10.1257/aer.20161723>.

Appendices

A	Theoretical Derivations	2
A.1	The Revenue Gradient	2
A.2	The Welfare Gradient	3
A.3	The Ramsey First-Order Condition	4
A.4	Proof of the Projection Lemma	5
A.5	Proof of the Alignment Identities	5
A.6	Broad and Within-Instrument Reforms	6
B	Data	10
B.1	Quarterly Tax Data	10
B.2	Progressivity Data and Construction	14
C	Estimation and Posterior Construction	17
C.1	Estimator and Posterior Draws	17
C.2	Quarterly Broad Tax Block	19
C.3	Annual Progressivity Block	20
C.4	Combining the Two Blocks	22
D	Robustness of the Broad Tax Elasticity Estimates	22
D.1	Rationalizing the Magnitude of ε_{CL}	29
E	Robustness of the Progressivity Semi-Elasticities	32
F	Identification Diagnostics	36
G	Alignment Metric Robustness	37
H	Fiscal Constraints over the Observed Horizon	42
H.1	Empirical Implementation	43
I	Additional Robustness Figures	45

Online Appendix

A Theoretical Derivations

This appendix provides the formal derivations underlying Section 2. Unless otherwise stated, all derivatives are evaluated at the baseline policy vector $\bar{\tau}$. Bars are retained where they prevent ambiguity and otherwise suppressed. The main text writes $\mathbf{w} = \nabla_{\tau} W$ and $\mathbf{r} = \nabla_{\tau} R$. When the coordinate system matters, I write these same derivatives as

$$\mathbf{w}^{\tau} := \nabla_{\tau} W, \quad \mathbf{r}^{\tau} := \nabla_{\tau} R.$$

The superscript distinguishes derivatives in the original rate coordinates from their representation after a change in the policy metric. It does not denote a different welfare or revenue object.

A.1 The Revenue Gradient

Recall that present value revenue is

$$R(\tau) = \sum_{s=0}^T \beta^s \sum_{k=1}^n \tau_{k,s} B_{k,s}(\tau).$$

Assumption 1 permits differentiation term by term. For any rate $\tau_{i,t}$,

$$\begin{aligned} r_{i,t} &:= \frac{\partial R}{\partial \tau_{i,t}} = \beta^t B_{i,t} + \sum_{s=0}^T \beta^s \sum_{k=1}^n \tau_{k,s} \frac{\partial B_{k,s}}{\partial \tau_{i,t}} \\ &= \beta^t B_{i,t} + \sum_{s=0}^T \beta^s \tau_{i,s} \frac{\partial B_{i,s}}{\partial \tau_{i,t}} + \sum_{s=0}^T \beta^s \sum_{k \neq i} \tau_{k,s} \frac{\partial B_{k,s}}{\partial \tau_{i,t}}. \end{aligned} \quad (\text{A.1})$$

The first equality follows from the product rule. The second merely separates responses of the same tax base from responses of other tax bases. No sign restriction is imposed on either set of behavioral responses.

The elasticity representation in Section 3.1 follows directly from equation (A.1). Suppose that $B_{k,s} > 0$ and $\bar{\tau}_{i,t} < 1$, and define

$$\varepsilon_{k,s;i,t} := \left. \frac{\partial \ln B_{k,s}(\tau)}{\partial \ln(1 - \tau_{i,t})} \right|_{\bar{\tau}}.$$

By the chain rule,

$$\varepsilon_{k,s;i,t} = \frac{1 - \bar{\tau}_{i,t}}{B_{k,s}} \frac{\partial B_{k,s}}{\partial (1 - \tau_{i,t})} = - \frac{1 - \bar{\tau}_{i,t}}{B_{k,s}} \frac{\partial B_{k,s}}{\partial \tau_{i,t}}, \quad (\text{A.2})$$

and hence

$$\frac{\partial B_{k,s}}{\partial \tau_{i,t}} = - \frac{B_{k,s}}{1 - \bar{\tau}_{i,t}} \varepsilon_{k,s;i,t}. \quad (\text{A.3})$$

A positive net-of-tax elasticity therefore means that the base contracts when the tax rate rises. Substituting equation (A.3) into equation (A.1) gives the exact local formula

$$r_{i,t} = \beta^t B_{i,t} - \frac{1}{1 - \bar{\tau}_{i,t}} \sum_{s=0}^T \beta^s \sum_{k=1}^n \bar{\tau}_{k,s} B_{k,s} \varepsilon_{k,s;i,t}. \quad (\text{A.4})$$

Separating the term $(k, s) = (i, t)$, the remaining responses of base i , and the responses of all other bases yields

$$r_{i,t} = \beta^t B_{i,t} - \frac{1}{1 - \bar{\tau}_{i,t}} \left[\beta^t \bar{\tau}_{i,t} B_{i,t} \varepsilon_{i,t;i,t} + \sum_{s \neq t} \beta^s \bar{\tau}_{i,s} B_{i,s} \varepsilon_{i,s;i,t} + \sum_{s=0}^T \beta^s \sum_{k \neq i} \bar{\tau}_{k,s} B_{k,s} \varepsilon_{k,s;i,t} \right]. \quad (\text{A.5})$$

Thus, each elasticity is weighted by the baseline revenue collected from the responding base, $\beta^s \bar{\tau}_{k,s} B_{k,s}$, and by the conversion from a change in the tax rate to a proportional change in its net-of-tax rate. An unweighted sum of elasticities would not recover the revenue gradient.

The same derivatives also recover the response to a reform that changes several rates at once. For any infinitesimal reform $d\boldsymbol{\tau}$, the chain rule and equation (A.3) give

$$dB_{k,s} = -B_{k,s} \sum_{t=0}^T \sum_{i=1}^n \frac{\varepsilon_{k,s;i,t}}{1 - \bar{\tau}_{i,t}} d\tau_{i,t}. \quad (\text{A.6})$$

Thus, rate-specific elasticities can be combined using the rate changes that define the reform. Alternatively, an empirical design may estimate the left side of equation (A.6) directly for the combined reform. In either case, the reform's first-order fiscal effect is $dR = \mathbf{r}^\top d\boldsymbol{\tau}$.

The decomposition by date is the same derivative regrouped in a different way. Defining $R_s(\boldsymbol{\tau}) = \sum_k \tau_{k,s} B_{k,s}(\boldsymbol{\tau})$ gives

$$r_{i,t} = \sum_{s=0}^T \beta^s \frac{\partial R_s}{\partial \tau_{i,t}}.$$

Partitioning the sum into $s = t$, $s > t$, and $s < t$ gives equation (3). This temporal decomposition and the base decomposition in equation (A.1) are therefore two partitions of the same derivative, not separate assumptions about the revenue response.

A.2 The Welfare Gradient

The formula for the welfare gradient requires a distinction between the envelope argument and the incidence accounting that follows it. Let $C_h(\boldsymbol{\tau})$ denote group h 's present value money-metric welfare cost after its members have optimized, and write

$$W(\boldsymbol{\tau}) = \sum_h \omega_h C_h(\boldsymbol{\tau}).$$

For an interior private choice, the envelope theorem removes the terms that operate solely through the induced change in that choice. For example, if $\mathbf{x}_h^*(\boldsymbol{\tau})$ denotes the vector of choices made by members of group h , the derivative of optimized welfare contains no term proportional to the private first-order condition times $\partial \mathbf{x}_h^* / \partial \tau_{i,t}$. This does not remove changes in wages, prices, profits, transfers, or other equilibrium objects from welfare. Those changes affect optimized welfare directly and must be included in $\partial C_h / \partial \tau_{i,t}$.

When $\beta^t B_{i,t} > 0$, define the incidence share

$$s_{h,i,t} := \frac{1}{\beta^t B_{i,t}} \frac{\partial C_h}{\partial \tau_{i,t}}. \quad (\text{A.7})$$

This definition assigns to group h every first-order effect on its money-metric welfare, including effects operating

through equilibrium wages, prices, and profits. A negative share means that the group benefits from the tax increase. Combining equation (A.7) with the definition of W gives

$$\begin{aligned} w_{i,t} &:= \frac{\partial W}{\partial \tau_{i,t}} = \sum_h \omega_h \frac{\partial C_h}{\partial \tau_{i,t}} \\ &= \beta^t B_{i,t} \sum_h \omega_h s_{h,i,t} = \beta^t B_{i,t} \tilde{y}_{i,t}, \end{aligned} \quad (\text{A.8})$$

where $\tilde{y}_{i,t} := \sum_h \omega_h s_{h,i,t}$. Once the incidence shares are defined to include the complete money-metric burden, the final equality is an accounting identity. The envelope theorem explains why the agents' induced optimizing choices do not need to be added to that burden as separate first-order terms; it does not, by itself, determine the incidence shares.

As in the main text, equation (A.8) requires every first-order consequence of the reform to enter either household money-metric burdens or consolidated government revenue. A fiscal effect on a government omitted from R must be added to consolidated revenue. A nonfiscal externality or a policy interaction that changes social resources must be added to the appropriate C_h or entered as a separate term in W . Otherwise, equation (A.8) omits a first-order consequence of the reform.

The further condition

$$\sum_h s_{h,i,t} = 1 \quad (\text{A.9})$$

is not implied by the envelope theorem. It instead requires the groups represented in W to bear the entire marginal mechanical burden. When equation (A.9) holds and $\omega_h = 1$ for every included group, $\tilde{y}_{i,t} = 1$ and $w_{i,t} = \beta^t B_{i,t}$. When some burden falls on groups excluded from W , or when social welfare weights differ, the general expression in equation (A.8) applies.

A.3 The Ramsey First-Order Condition

Consider the local Ramsey problem

$$\min_{\boldsymbol{\tau}} W(\boldsymbol{\tau}) \quad \text{subject to} \quad R(\boldsymbol{\tau}) = \bar{R}. \quad (\text{A.10})$$

Proposition 3 (Ramsey first-order condition). *Suppose that $\boldsymbol{\tau}^*$ is an interior solution to equation (A.10), that W and R are continuously differentiable in a neighborhood of $\boldsymbol{\tau}^*$, and that the revenue constraint is regular, so $\mathbf{r}(\boldsymbol{\tau}^*) \neq 0$. Then there exists a scalar μ such that*

$$\mathbf{w}(\boldsymbol{\tau}^*) = \mu \mathbf{r}(\boldsymbol{\tau}^*). \quad (\text{A.11})$$

If the minimized welfare cost is nondecreasing in the revenue requirement and $\mathbf{w}(\boldsymbol{\tau}^) \neq 0$, then $\mu > 0$. For every component with $r_{i,t}(\boldsymbol{\tau}^*) \neq 0$,*

$$\frac{w_{i,t}(\boldsymbol{\tau}^*)}{r_{i,t}(\boldsymbol{\tau}^*)} = \mu. \quad (\text{A.12})$$

Proof. Form the Lagrangian

$$\mathcal{L}(\boldsymbol{\tau}, \mu) = W(\boldsymbol{\tau}) - \mu(R(\boldsymbol{\tau}) - \bar{R}).$$

Interior first-order conditions give $\mathbf{w}(\boldsymbol{\tau}^*) - \mu \mathbf{r}(\boldsymbol{\tau}^*) = 0$. Let $V(\bar{R})$ denote the minimized value in equation (A.10). The envelope theorem gives $V'(\bar{R}) = \mu$ wherever the value function is differentiable, so the economically standard condition that an additional revenue requirement does not lower the minimized welfare cost implies $\mu \geq 0$. Monotonicity

alone does not deliver a strict inequality, since a strictly increasing function can have a zero derivative at a point. The sign is pinned down instead by equation (A.11): if μ were zero then $\mathbf{w}(\boldsymbol{\tau}^*)$ would be zero as well, so $\mathbf{w}(\boldsymbol{\tau}^*) \neq 0$ gives $\mu \neq 0$ and hence $\mu > 0$. Dividing the vector condition by any nonzero component of $\mathbf{r}$ gives equation (A.12). $\square$

The vector condition in equation (A.11) is more general than the ratio formulation. If $r_{i,t} = 0$, the corresponding MVPF is undefined, but the first-order condition remains well defined and requires $w_{i,t} = 0$ at an interior optimum. Positive μ is also what distinguishes Ramsey alignment from negative proportionality. These are necessary first-order conditions at a regular interior optimum. They do not by themselves establish that a candidate is a local or global optimum.

A.4 Proof of the Projection Lemma

Proof of Lemma 1. Let

$$\mathcal{H}_{\mathbf{r}} := \{d\boldsymbol{\tau} : \mathbf{r}^\top d\boldsymbol{\tau} = 0\}$$

denote the revenue-neutral hyperplane. Since $\mathbf{r} \neq 0$, $P_{\perp\mathbf{r}} = I - \mathbf{r}(\mathbf{r}^\top \mathbf{r})^{-1} \mathbf{r}^\top$ is the orthogonal projector onto $\mathcal{H}_{\mathbf{r}}$. For every $d\boldsymbol{\tau} \in \mathcal{H}_{\mathbf{r}}$,

$$\mathbf{w}^\top d\boldsymbol{\tau} = (P_{\perp\mathbf{r}} \mathbf{w})^\top d\boldsymbol{\tau},$$

because the component of $\mathbf{w}$ parallel to $\mathbf{r}$ is orthogonal to every feasible reform. The Cauchy–Schwarz inequality therefore gives

$$-\mathbf{w}^\top d\boldsymbol{\tau} = -(P_{\perp\mathbf{r}} \mathbf{w})^\top d\boldsymbol{\tau} \leq \|P_{\perp\mathbf{r}} \mathbf{w}\| \|d\boldsymbol{\tau}\| \leq \alpha \|P_{\perp\mathbf{r}} \mathbf{w}\|.$$

If $P_{\perp\mathbf{r}} \mathbf{w} \neq 0$, both inequalities bind only when

$$d\boldsymbol{\tau} = -\alpha \frac{P_{\perp\mathbf{r}} \mathbf{w}}{\|P_{\perp\mathbf{r}} \mathbf{w}\|}.$$

This vector belongs to $\mathcal{H}_{\mathbf{r}}$, has norm α , and is therefore the unique solution. Its objective value is $\alpha \|P_{\perp\mathbf{r}} \mathbf{w}\|$. If $P_{\perp\mathbf{r}} \mathbf{w} = 0$, then $\mathbf{w}^\top d\boldsymbol{\tau} = 0$ for every feasible reform. Every feasible reform is then a solution to the first-order problem, and the zero reform may be selected. $\square$

A.5 Proof of the Alignment Identities

Proof of Proposition 1. For nonzero $\mathbf{w}$ and $\mathbf{r}$, the inner-product definition of the angle gives

$$\cos \theta = \frac{\mathbf{w}^\top \mathbf{r}}{\|\mathbf{w}\| \|\mathbf{r}\|}.$$

The decomposition of $\mathbf{w}$ into its components parallel and perpendicular to $\mathbf{r}$ is orthogonal, so

$$\begin{aligned} \|P_{\perp\mathbf{r}} \mathbf{w}\|^2 &= \|\mathbf{w}\|^2 - \left\| \mathbf{r} \frac{\mathbf{r}^\top \mathbf{w}}{\mathbf{r}^\top \mathbf{r}} \right\|^2 = \|\mathbf{w}\|^2 - \frac{(\mathbf{w}^\top \mathbf{r})^2}{\|\mathbf{r}\|^2} \\ &= \|\mathbf{w}\|^2 (1 - \cos^2 \theta) = \|\mathbf{w}\|^2 \sin^2 \theta. \end{aligned}$$

Because $\theta \in [0, \pi]$, $\sin \theta \geq 0$, which gives $\sin \theta = \|P_{\perp \mathbf{r}} \mathbf{w}\| / \|\mathbf{w}\|$. Combining this identity with Lemma 1 yields

$$-\mathbf{w}^\top d\boldsymbol{\tau}^* = \alpha \|\mathbf{w}\| \sin \theta.$$

Finally, equality in the Cauchy–Schwarz inequality implies $|\cos \theta| = 1$ if and only if the two gradients are proportional. The sign of their inner product distinguishes the two cases: $\cos \theta = 1$ if and only if $\mathbf{w} = \mu \mathbf{r}$ for some $\mu > 0$. $\square$

A.6 Broad and Within-Instrument Reforms

The projection lemma already applies to a policy vector containing every individual rate. This subsection does not introduce a new geometry. It decomposes that same policy space into broad directions and their within-instrument complements under a normalization that nests the broad reform problem. Restricting attention to broad reforms sets the second block to zero; allowing within-instrument reforms relaxes that restriction. The two components are therefore parts of one complete reform, and revenue neutrality applies only after they have been combined. This subsection makes the construction exact and proves Proposition 2.

Let $f \in \mathcal{F}$ index instrument-date families. Family f contains J_f rates, and t_f denotes its date. A broad instrument represented by a single rate is included by setting $J_f = 1$. Let $\mathbf{1}_f$ denote the J_f -vector of ones. For each family, write

$$B_{f+} := \sum_{j=1}^{J_f} B_{fj}, \quad q_{fj} := \frac{B_{fj}}{B_{f+}}, \quad \Omega_f := \text{diag}(q_{f1}, \dots, q_{fJ_f}),$$

and assume that every rate retained as a separate coordinate has $B_{fj} > 0$. Given a vector of rate changes $d\boldsymbol{\tau}_f$, define

$$da_f := \sum_{j=1}^{J_f} q_{fj} d\tau_{fj}, \quad d\boldsymbol{\tau}_{S,f} := d\boldsymbol{\tau}_f - da_f \mathbf{1}_f. \quad (\text{A.13})$$

This gives a unique decomposition of every change in family f . It also implies

$$\sum_{j=1}^{J_f} \beta^{t_f} B_{fj} d\tau_{S,fj} = \beta^{t_f} B_{f+} \sum_{j=1}^{J_f} q_{fj} d\tau_{S,fj} = 0. \quad (\text{A.14})$$

The common component therefore reproduces the change in tax payments at the existing bases, while the remaining component redistributes those payments across the rates. Equation (A.14) defines the within-instrument component. It does not require that component to be revenue-neutral once tax bases respond.

The base weights also provide a common measure of reform size. Stack the family-specific weight matrices as

$$\Omega := \bigoplus_{f \in \mathcal{F}} \Omega_f$$

and define

$$\|d\boldsymbol{\tau}\|_\Omega^2 := d\boldsymbol{\tau}^\top \Omega d\boldsymbol{\tau} = \sum_{f \in \mathcal{F}} \sum_{j=1}^{J_f} q_{fj} (d\tau_{fj})^2.$$

Within each family, equation (A.14) implies

$$\langle da_f \mathbf{1}_f, d\boldsymbol{\tau}_{S,f} \rangle_{\Omega_f} = da_f \sum_{j=1}^{J_f} q_{fj} d\tau_{S,fj} = 0.$$

Adding these identities across families gives

$$\|d\boldsymbol{\tau}\|_{\Omega}^2 = \sum_{f \in \mathcal{F}} (da_f)^2 + \sum_{f \in \mathcal{F}} \|d\boldsymbol{\tau}_{S,f}\|_{\Omega_f}^2. \quad (\text{A.15})$$

Thus, changes across broad instruments and changes within instruments are orthogonal blocks of the same reform space. A one-percentage-point common change in one family has size one. Dividing that family into additional rate brackets does not change the size of the common change. The main text takes a common movement in all rates as the broad change. Another definition of a broad change can be used in the same way after its direction has been specified and normalized. The personal instrument in the application is such a case. Its broad coordinate is the average personal tax rate rather than the common movement in the two group marginal rates, and the two are worth different amounts at the existing bases, so the within-family direction is converted to the units of the broad coordinate before its length is taken. Equations (A.33) and (A.34) carry out that conversion.

To express the blocks in ordinary Euclidean coordinates, define

$$d\mathbf{z} := \Omega^{1/2} d\boldsymbol{\tau}.$$

For each family, let

$$\mathbf{e}_f := \Omega_f^{1/2} \mathbf{1}_f = (\sqrt{q_{f1}}, \dots, \sqrt{q_{fJ_f}})^{\top},$$

so $\mathbf{e}_f^{\top} \mathbf{e}_f = 1$. The orthogonal projectors for that family are

$$P_{A,f} := \mathbf{e}_f \mathbf{e}_f^{\top}, \quad P_{S,f} := I_{J_f} - \mathbf{e}_f \mathbf{e}_f^{\top}.$$

Thus, if $d\mathbf{z}_f = \Omega_f^{1/2} d\boldsymbol{\tau}_f$, its components are

$$d\mathbf{z}_{A,f} := P_{A,f} d\mathbf{z}_f = \mathbf{e}_f da_f, \quad d\mathbf{z}_{S,f} := P_{S,f} d\mathbf{z}_f.$$

Embed each $\mathbf{e}_f$ in the complete standardized policy space and collect the embedded vectors as the columns of E . Because different families have disjoint coordinates, $E^{\top} E = I$. For each family, choose a matrix U_f whose columns form an orthonormal basis for the range of $P_{S,f}$, and set

$$U := \bigoplus_{f \in \mathcal{F}} U_f.$$

The columns of E and U together form an orthonormal basis. Every complete reform therefore has the unique representation

$$d\mathbf{z} = E d\mathbf{a} + U d\mathbf{s}, \quad d\mathbf{a} := E^{\top} d\mathbf{z}, \quad d\mathbf{s} := U^{\top} d\mathbf{z}. \quad (\text{A.16})$$

The particular matrices U_f used to represent the within-instrument subspaces do not affect the result. A different orthonormal basis rotates the coordinates in $d\mathbf{s}$, $\mathbf{w}_S$, and $\mathbf{r}_S$ together and therefore leaves their inner products, norms, and implied reforms unchanged. The entry da_f is the common change in equation (A.13) and corresponds to $d\tau_{A,t_f}$

in the main text. The vector $d\mathbf{s}$ is an orthonormal, standardized representation of the rate changes $\{d\mathbf{r}_{S,f}\}_{f \in \mathcal{F}}$; it is not the unstandardized rate vector itself. It follows directly that

$$\|d\mathbf{r}\|_{\Omega}^2 = \|d\mathbf{a}\|^2 + \|d\mathbf{s}\|^2.$$

Return now to the original rate-coordinate derivatives $\mathbf{w}^r$ and $\mathbf{r}^r$. Their gradients in the standardized coordinates are

$$\tilde{\mathbf{w}} := \Omega^{-1/2} \mathbf{w}^r, \quad \tilde{\mathbf{r}} := \Omega^{-1/2} \mathbf{r}^r, \quad (\text{A.17})$$

because $(\mathbf{w}^r)^\top d\mathbf{r} = \tilde{\mathbf{w}}^\top d\mathbf{z}$ and similarly for revenue. Define

$$\mathbf{w}_A := E^\top \tilde{\mathbf{w}}, \quad \mathbf{w}_S := U^\top \tilde{\mathbf{w}}, \quad \mathbf{r}_A := E^\top \tilde{\mathbf{r}}, \quad \mathbf{r}_S := U^\top \tilde{\mathbf{r}}. \quad (\text{A.18})$$

Then

$$dW = \mathbf{w}_A^\top d\mathbf{a} + \mathbf{w}_S^\top d\mathbf{s}, \quad dR = \mathbf{r}_A^\top d\mathbf{a} + \mathbf{r}_S^\top d\mathbf{s}.$$

Thus, $\mathbf{w}_S$ and $\mathbf{r}_S$ are the welfare and revenue derivatives with respect to the standardized within-instrument coordinates. Their ordinary Euclidean norms are the norms used for the within-instrument block in the main text. For the remainder of this subsection, and in Proposition 2, write

$$\mathbf{w} := \begin{pmatrix} \mathbf{w}_A \\ \mathbf{w}_S \end{pmatrix}, \quad \mathbf{r} := \begin{pmatrix} \mathbf{r}_A \\ \mathbf{r}_S \end{pmatrix} \quad (\text{A.19})$$

for the complete gradients in these standardized block coordinates. This is an explicit change of coordinates from $\mathbf{w}^r$ and $\mathbf{r}^r$, not a change in the underlying welfare and revenue derivatives. When every family has one rate, $\Omega = I$ and the within-instrument block is empty, so this notation reduces to the broad reform problem.

In these coordinates, that same reform problem becomes

$$\begin{aligned} \max_{d\mathbf{a}, d\mathbf{s}} \quad & -\mathbf{w}_A^\top d\mathbf{a} - \mathbf{w}_S^\top d\mathbf{s} \\ \text{subject to} \quad & \mathbf{r}_A^\top d\mathbf{a} + \mathbf{r}_S^\top d\mathbf{s} = 0, \\ & \|d\mathbf{a}\|^2 + \|d\mathbf{s}\|^2 \leq \alpha^2. \end{aligned} \quad (\text{A.20})$$

Setting $d\mathbf{s} = 0$ recovers the broad reform problem and the corresponding angle θ_A . Allowing $d\mathbf{s}$ to vary enlarges that same reform space; it does not introduce a separate alignment criterion. Neither block must be revenue-neutral on its own. Only their combined revenue effect must be zero. The projection and alignment results therefore apply without modification to the complete gradients in equation (A.19).

This construction also proves the equal incidence result. For rate j in family f , equal incidence gives

$$w_{fj}^r = \beta^{tf} B_{fj} = \beta^{tf} B_{f+} q_{fj}.$$

The standardized welfare gradient for that family is therefore

$$\tilde{\mathbf{w}}_f = \Omega_f^{-1/2} \mathbf{w}_f^r = \beta^{tf} B_{f+} \mathbf{e}_f.$$

It lies entirely in the broad direction, so $U_f^\top \tilde{\mathbf{w}}_f = 0$ for every family and $\mathbf{w}_S = \mathbf{0}$. The same result holds whenever $\tilde{y}_{fj}$ is constant within each family, although the constant may differ across families. The revenue gradient need not lie

entirely in the broad block because rates within a family can generate different behavioral responses. The mechanical component of the revenue gradient is proportional to $(B_{f1}, \dots, B_{fJ_f})$ and therefore lies in the broad direction as well. Consequently, $\mathbf{r}_S$ is generated by behavioral effects on the taxed base, other tax bases, or revenue at other dates.

Because the broad and within-instrument blocks are orthogonal in the standardized coordinates, the alignment result applies to their complete gradients without modification. Their inner product and squared norms are

$$\mathbf{w}^\top \mathbf{r} = \mathbf{w}_A^\top \mathbf{r}_A + \mathbf{w}_S^\top \mathbf{r}_S, \quad \|\mathbf{w}\|^2 = \|\mathbf{w}_A\|^2 + \|\mathbf{w}_S\|^2, \quad \|\mathbf{r}\|^2 = \|\mathbf{r}_A\|^2 + \|\mathbf{r}_S\|^2.$$

Substitution into the definition of the cosine gives

$$\cos \theta = \frac{\mathbf{w}_A^\top \mathbf{r}_A + \mathbf{w}_S^\top \mathbf{r}_S}{\sqrt{\|\mathbf{w}_A\|^2 + \|\mathbf{w}_S\|^2} \sqrt{\|\mathbf{r}_A\|^2 + \|\mathbf{r}_S\|^2}}. \quad (\text{A.21})$$

At a regular interior Ramsey optimum, Proposition 3 requires a common $\mu > 0$ such that

$$\mathbf{w}_A = \mu \mathbf{r}_A, \quad \mathbf{w}_S = \mu \mathbf{r}_S.$$

Thus, alignment in only the broad block is not sufficient for alignment in the complete reform space.

Under equal incidence, $\mathbf{w}_S = \mathbf{0}$. If $\mathbf{w}_A$ and $\mathbf{r}_A$ are nonzero and θ_A is the angle between them, equation (A.21) becomes

$$\begin{aligned} \cos \theta &= \frac{\mathbf{w}_A^\top \mathbf{r}_A}{\|\mathbf{w}_A\| \sqrt{\|\mathbf{r}_A\|^2 + \|\mathbf{r}_S\|^2}} \\ &= \cos \theta_A \frac{\|\mathbf{r}_A\|}{\sqrt{\|\mathbf{r}_A\|^2 + \|\mathbf{r}_S\|^2}}. \end{aligned} \quad (\text{A.22})$$

The multiplicative term is at most one and is strictly positive. It equals one if and only if $\mathbf{r}_S = \mathbf{0}$. Consequently, $\cos \theta = 1$ if and only if $\cos \theta_A = 1$ and $\mathbf{r}_S = \mathbf{0}$.

The bundled reform behind this result can also be constructed directly. Suppose the broad block is positively aligned, $\mathbf{w}_A = \mu \mathbf{r}_A$ for some $\mu > 0$, with $\mathbf{r}_A \neq \mathbf{0}$, while $\mathbf{w}_S = \mathbf{0}$ and $\mathbf{r}_S \neq \mathbf{0}$. Choose $d\mathbf{s}$ such that $\mathbf{r}_S^\top d\mathbf{s} > 0$; reversing any direction with a nonzero revenue effect produces this sign. Because $\mathbf{r}_A \neq \mathbf{0}$, choose

$$d\mathbf{a} = -\frac{\mathbf{r}_S^\top d\mathbf{s}}{\mathbf{r}_A^\top \mathbf{r}_A} \mathbf{r}_A.$$

The complete reform is revenue neutral, since $\mathbf{r}_A^\top d\mathbf{a} + \mathbf{r}_S^\top d\mathbf{s} = 0$, while its first-order effect on welfare cost is

$$\begin{aligned} \mathbf{w}_A^\top d\mathbf{a} + \mathbf{w}_S^\top d\mathbf{s} &= \mu \mathbf{r}_A^\top d\mathbf{a} \\ &= -\mu \mathbf{r}_S^\top d\mathbf{s} < 0. \end{aligned}$$

Both components can be scaled down together to satisfy any prescribed local size bound. A nonzero within-instrument revenue gradient therefore leaves a revenue-neutral welfare improvement even when the broad block by itself satisfies the Ramsey condition.

Finally, squaring equation (A.22) and using $\sin^2 x = 1 - \cos^2 x$ gives

$$\begin{aligned}\sin^2 \theta &= 1 - \cos^2 \theta_A \frac{\|\mathbf{r}_A\|^2}{\|\mathbf{r}_A\|^2 + \|\mathbf{r}_S\|^2} \\ &= \sin^2 \theta_A + \cos^2 \theta_A \frac{\|\mathbf{r}_S\|^2}{\|\mathbf{r}_A\|^2 + \|\mathbf{r}_S\|^2},\end{aligned}\tag{A.23}$$

which is equation (19). The additional term is nonnegative. Under equal incidence, enlarging the reform space to include within-instrument directions therefore cannot reduce the first-order scope for revenue-neutral reform under the stated normalization.

B Data

The empirical implementation combines quarterly data on broad personal, corporate, and tariff rates with annual data on four personal marginal rates and group incomes, together with annual corporate and tariff variables. The quarterly BLP is estimated over 1947Q1–2019Q4. The underlying quarterly rate and base series extend beyond that estimation endpoint and supply annual corporate and tariff states and accounting histories through 2023. The annual BLP spans 1947–2023, and the accounting series used to construct the alignment measures begin in 1948.

B.1 Quarterly Tax Data

The personal and corporate series extend the data in Mertens and Ravn (2013a) and Cloyne et al. (2022). The tariff series are constructed over the same quarterly sample.

- **Personal Income Tax Base.** The personal income tax base is NIPA personal income (Table 2.1 Line 1) plus contributions for government social insurance (Table 3.2 Line 11) less transfers (Table 2.1 Line 17).
- **Personal Income Tax Revenue and Rate.** Personal tax revenue is the sum of NIPA Table 3.2 Lines 3 and 11. The average personal income tax rate is revenue divided by the contemporaneous base.
- **Corporate Income Tax Base.** The corporate tax base is corporate profits with inventory valuation and capital consumption adjustments (NIPA Table 1.12 Line 13), less Federal Reserve Bank profits from Historical Tables 6.16B–D. However, corporations are not domestically owned entirely, and the share of foreign ownership has changed overtime. In some robustness exercises, I remove the foreign owners from the tax base using the annual rest-of-world share of U.S. corporate equity in the Flow of Funds. I measure it as

$$\text{Foreign ownership share}_t = \frac{\text{BOGZ1LM263064105A}_t}{\text{BOGZ1LM893064105A}_t},$$

where both series are from FRED. I match the ratio to the accounting year.

- **Corporate Income Tax Revenue and Rate.** Corporate tax revenue is federal taxes on corporate income excluding Federal Reserve Banks (NIPA Table 3.2 Line 8). The average corporate income tax rate is revenue divided by the contemporaneous base.
- **Tariff Base, Revenue, and Rate.** The tariff base is imports of goods (NIPA Table 4.1 Line 19). Tariff revenue is federal customs duties (NIPA Table 3.2 Line 6), and the average tariff rate is customs duties divided by imports

of goods. However, the quarterly NIPA customs-duty series begins in 1959Q1. Over 1947–1958, I combine the annual NIPA totals with monthly Treasury customs receipts from the NBER Macroeconomy Database (FRED Series M15001USM144NNBR). I seasonally adjust the monthly series using X-13ARIMA-SEATS, aggregate it to quarters, and use proportional Denton–Cholette benchmarking to distribute each annual NIPA total across the four quarters. The quarterly values are expressed at seasonally adjusted annual rates and spliced to the NIPA series in 1959Q1

All base series are deflated by the GDP deflator (NIPA Table 1.1.9 Line 1), divided by the population over age 16, and log-transformed. Population is the Bureau of Labor Statistics series CNP16OV. The series begins in 1948, so I linearly extrapolate it over the four quarters of 1947. The three policy variables enter the quarterly system as


$$\log(1 - \tau_{P,t}), \quad \log(1 - \tau_{C,t}), \quad -\log(1 + \tau_{T,t}),$$

where P , C , and T denote the broad personal, corporate, and tariff coordinates. The last sign convention puts a tariff cut in the same direction as an income-tax cut.

The unfiltered transformed rates and log bases enter the quarterly BLP over its 1947Q1–2019Q4 estimation sample. The same underlying series, extended beyond that endpoint, provide the annual corporate and tariff states described below. For the accounting histories used to evaluate the alignment measures, I instead apply the symmetric filter of Christiano and Fitzgerald (2003) to the personal and corporate rates, the tariff rate, and the log bases, removing components with periods between 2 and 40 quarters. The resulting trends supply the broad quarterly accounting histories. For the annual nonlinear accounting, I average them within each year and denote the annual bases by $B_{P,t}$, $B_{C,t}$, and $B_{T,t}$ for personal income, corporate income, and imports, respectively. Thus the annual corporate and tariff states and the smooth accounting histories are constructed from the same underlying quarterly data but serve different purposes.

Narrative Instruments. The personal and corporate instruments are the narrative series in Mertens and Ravn (2013a), extended through 2019 by Cloyne et al. (2022). I construct the tariff instrument from the narrative episodes in Besten et al. (2026), whose annual series records the sign of each change. I assign a change to its implementation quarter and set the other quarters to zero. Before the Trump tariffs, I scale each change by the prevailing tariff rate and the tariff reduction reported by Bown and Irwin (2015); for the Trump years, I use the timing and magnitudes in Amiti, Redding, and Weinstein (2020). The three instruments enter jointly, and I impose no exclusion that a narrative measure affect only its named tax rate. For each series, I subtract its mean across nonzero reform observations and leave non-reform quarters at zero. These recentered series are the instruments used in the quarterly BLP.

Figure B.1: Narrative Tax Reforms



Note: The top panel reports the quarterly personal income tax, the middle the corporate income tax, and the bottom the tariff narrative measures. Only nonzero observations are shown.

Other Macro Variables. The expanded-information-set specification adds real federal government consumption and gross investment (NIPA Table 1.1.3 Line 23), real GDP (NIPA Table 1.1.3 Line 1), and federal debt held by the public. The debt series splices the measure in Favero and Giavazzi (2012) to FRED Series FYGFDPU. These variables are auxiliary controls; they are not part of the baseline revenue gradient. Federal purchases and GDP are divided by the population over age 16 and logged. I deflate nominal debt with the GDP deflator, divide it by the same population series, and log the resulting real per-capita quantity.

SOI Construction

The baseline specification uses NIPA corporate profits as the corporate tax base, which includes both C-corporation and S-corporation income. Since S-corporations are passthrough entities whose income is taxed at the shareholder level rather than at the entity, one might be concerned that the results are driven by organizational form arbitrage rather than traditional general equilibrium channels.


To address this concern, we construct adjusted tax bases that reallocate S-corporation income from the corporate to the personal base. I use the Annual Corporate Reports from the Statistics of Income (SOI) to compile a time series of the profits of C-corporations and S-corporations. I measure profits for the former using the SOI's measure of "Income Subject to Tax" for form 1120 filers (excluding 1120-REIT, 1120-RIC, and 1120-S filers) and profits for

the latter as “Net Income (Less Deficit)” for form 1120-S filers. Since Subchapter S was enacted in 1958, allowing qualifying corporations to elect passthrough treatment, I do this for every year from 1958–2022 and carry the 2022 log adjustment forward for the 2023 robustness observation. Toward reallocating NIPA corporate income for S-corps, I compute the S-corporation share of total corporate income as:

$$s_t = \frac{\text{S-corp net income}_t}{\text{S-corp net income}_t + \text{C-corp net income}_t} \quad (\text{A.24})$$

using SOI data for both numerator and denominator to maintain consistency. The resulting series is plotted in Figure B.2. As is well-documented elsewhere, s_t is negligible until after the Tax Reform Act of 1986; since then, S-corps have been large relative to C-corps (Dyrda and Pugsley 2024).

Figure B.2: S-Corporate Share of Corporate Profits



Note: The S-corporate share of corporate profits is computed from the Statistics of Income’s Annual Corporate Reports.


NIPA applies a number of conceptual adjustments to corporate profits which are sensible economically, but are not present in the SOI data. To maintain consistency, I assume that these adjustments are proportional, and get quarterly tax bases by applying s_t in year t to NIPA data in every quarter q of that year as:

$$\text{Corporate base}_{q,t}^{adj} = \text{NIPA corporate profits}_{q,t} \times (1 - s_t) \quad (\text{A.25})$$

$$\text{Personal base}_{q,t}^{adj} = \text{NIPA personal income}_{q,t} + \text{NIPA corporate profits}_{q,t} \times s_t \quad (\text{A.26})$$

The reallocation also changes the measured average rates and associated quarterly first-stage inputs. Figure B.3 records those source series before they enter the robustness estimation.

Figure B.3: Comparing Mertens-Ravn and Reallocated S-Corp Tax Rates and Shocks



Note: After reallocating S-corps to the personal tax base, I recompute average tax rates and tax shocks (as shown above). There is a somewhat sizable difference for the average corporate tax rate, but the results are similar otherwise.

B.2 Progressivity Data and Construction

The annual progressivity exercise uses the lower 99 percent and top 1 percent series constructed by Mertens and Montiel Olea (2018). For each group, the tax rate is the average marginal individual income tax rate plus the average marginal payroll tax rate. The outcome is log real average market income excluding capital gains. The narrative measures are selected federal reforms scored as changes in the marginal tax rate faced by each group. I use $\log(1 - \tau_{g,t})$ as the policy variable for each group in $\{L, H\}$, where L denotes the lower 99 percent and H denotes the top 1 percent. The corresponding narrative measure is the negative of the MMO rate change divided by 100, so that a positive value denotes a tax cut.

The annual dataset also includes the corporate and tariff tax bases. I construct these series from their quarterly counterparts. Annualized corporate and tariff narrative measures enter the annual system as contemporaneous controls. The corporate and tariff rates do not enter the annual state vector.

Extension after 2012

The original MMO data end in 2012. I preserve their observations through that year and extend the tax rate and income series through 2023 using Statistics of Income data. Beginning in 2013, I apply year-over-year changes from SOI Tables 3.5, 4.1, and 1.4 to the original 2012 levels. The individual income tax and payroll tax components are constructed separately and add to the combined marginal rate.

The post-2012 top-1 series is extended directly. I construct the lower-99 series by aggregating the bottom 90 percent, 90th–95th percentiles, and 95th–99th percentiles. Rates are aggregated using AGI-dollar band weights; average income is aggregated as total band income excluding capital gains divided by total band returns. Both series are anchored to their MMO values in 2012. The post-2012 income extension uses returns rather than tax units. The replication package provides the complete construction.

I add one post-2012 narrative reform: the Tax Cuts and Jobs Act. I construct an MMO-style direct-schedule score using published SOI cells and the fixed 2017 income distribution. The underlying rate changes for the bottom 90 percent, 90th–95th percentiles, 95th–99th percentiles, and top 1 percent are -2.407 , -3.036 , -4.996 , and -2.819 percentage points, respectively. Aggregating the first three changes with the accounting schedule’s group income shares in the reform year gives a lower-99 narrative measure of 0.03033 , while the top-1 measure is 0.02819 . These values are added in code to the eight original MMO reform years.

The TCJA was enacted in December 2017 and applied in 2018, so I classify it as a long-run reform under MMO’s timing criterion. I also report estimates that omit it. I do not include the 2013 extension of the American Taxpayer Relief Act because its near-term stabilization motivation does not satisfy the same narrative criterion. All other post-2012 observations of the narrative measures are zero.

Mechanical Tax Exposure and the Progressivity Direction

The annual estimates describe how the two group incomes and the corporate and tariff bases respond to changes in the lower-99 and top-1 marginal tax rates. To define the progressivity reform, I use the MMO liability scores to determine how much each rate must change for initial personal tax payments to remain fixed.

Let $Y_{L,t}$ and $Y_{H,t}$ denote current-dollar income, in millions, assigned to the lower 99 percent and top 1 percent. Through 2012, these series are constructed from MMO market income excluding capital gains. Thereafter, I extend the 2012 total using aggregate SOI AGI excluding capital gains and allocate it between the two groups using the extended average-income series and their fixed group masses.

I use the official ex ante liability scores from the eight original MMO reforms to map group income into mechanical tax exposure. These scores include the effects of individual income tax and payroll tax marginal-rate provisions rather than the total revenue effects of every provision in each law. Let $s_{H,e-1}$ denote the lagged top-1 income share in reform episode e , let $\Delta\tau_{H,e}$ and $\Delta\tau_{L,e}$ denote the scored tax rate changes for the top 1 percent and lower 99 percent, and let $\Delta L_e/Y_{e-1}$ denote the associated liability change relative to lagged market income. I estimate

$$\frac{\Delta L_e}{Y_{e-1}} = b_H s_{H,e-1} \Delta\tau_{H,e} + b_L (1 - s_{H,e-1}) \Delta\tau_{L,e} + u_e. \quad (\text{A.27})$$

The estimates are $b_H = 0.191$ and $b_L = 0.599$.

The resulting mechanical exposures are

$$m_{L,t} = b_L Y_{L,t}, \quad m_{H,t} = b_H Y_{H,t}.$$

Define the corresponding exposure shares as

$$s_{L,t} = \frac{m_{L,t}}{m_{L,t} + m_{H,t}}, \quad s_{H,t} = \frac{m_{H,t}}{m_{L,t} + m_{H,t}}. \quad (\text{A.28})$$

Let $B_{P,t}^{\text{nom}}$ denote the nominal NIPA personal tax base, in billions. The scale of total group exposure relative to the

broad personal base is

$$q_t = \frac{(m_{L,t} + m_{H,t})/1000}{B_{P,t}^{\text{nom}}}. \quad (\text{A.29})$$

The accounting bases used to value the resulting income responses are

$$\widehat{B}_{g,t} = B_{P,t} \frac{Y_{g,t}/1000}{B_{P,t}^{\text{nom}}}, \quad g \in \{L, H\}. \quad (\text{A.30})$$

Because the group masses are fixed, the estimated response of log average income is also the response of log total group income.

A change dd_t in the progressivity instrument changes the two marginal rates by $\mathbf{h}_t dd_t$, where

$$\mathbf{h}_t = \begin{pmatrix} -s_{H,t} \\ 1 - s_{H,t} \end{pmatrix}. \quad (\text{A.31})$$

The lower-99 rate therefore falls by $s_{H,t}$ percentage points and the top-1 rate rises by $1 - s_{H,t}$ percentage points. This direction satisfies

$$s_{L,t} h_{L,t} + s_{H,t} h_{H,t} = 0, \quad h_{H,t} - h_{L,t} = 1. \quad (\text{A.32})$$

The first condition holds mapped mechanical personal tax liability fixed to first order. The second makes dd_t a one-percentage-point increase in the gap between the top-1 and lower-99 marginal rates.

I measure its length with the metric of Appendix A.6, which assigns family f the weights $q_{fj} = B_{fj}/B_{f+}$ and gives a within-family direction the squared length $\sum_j q_{fj} (d\tau_{S,fj})^2$. The mechanical-neutrality condition in (A.32) is $s_{L,t} h_{L,t} + s_{H,t} h_{H,t} = 0$, so the bases that define this family are the two exposures $m_{L,t}$ and $m_{H,t}$ and the family total is $B_{f+} = m_{L,t} + m_{H,t}$. Applying the metric to $\mathbf{h}_t$ directly gives


$$s_{L,t} h_{L,t}^2 + s_{H,t} h_{H,t}^2 = s_{L,t} s_{H,t}^2 + s_{H,t} s_{L,t}^2 = s_{L,t} s_{H,t}. \quad (\text{A.33})$$

That length is stated in the units of the family's own common change, which raises both group marginal rates by one percentage point and moves mechanical liability by $m_{L,t} + m_{H,t}$. The broad personal coordinate in this paper is not that reform. It is a one-percentage-point change in the average personal tax rate, which moves mechanical liability by $B_{P,t}^{\text{nom}}$. By (A.29) the two differ by the factor q_t , so one percentage point on the group marginal rates is worth q_t percentage points of the broad personal rate at the existing bases. Putting the progressivity direction in broad personal-rate units therefore replaces $\mathbf{h}_t$ by $q_t \mathbf{h}_t$ before the length is taken, and its squared length is

$$m_{d,t} = q_t^2 \left(s_{L,t} h_{L,t}^2 + s_{H,t} h_{H,t}^2 \right) = q_t^2 s_{L,t} s_{H,t}. \quad (\text{A.34})$$

The conversion preserves the normalization of Appendix A.6. The family's common change has size one in group marginal-rate units and size q_t in broad personal-rate units, which is what it is worth at the existing bases, and the four revenue coordinates are then lengths of reforms measured in one common unit.

Figure B.4: Narrative Personal Tax Reform Measures by Income Group



Note: The panels report the narrative personal tax measures for the lower 99 percent and top 1 percent. A positive value denotes a tax cut. Only the nine nonzero reform years are shown, including the recovered 2018 Tax Cuts and Jobs Act measures.

C Estimation and Posterior Construction

This appendix describes how the two empirical designs are mapped into the revenue gradients used in the paper. The quarterly block estimates the responses of the personal, corporate, and tariff bases to broad tax rate changes. The annual block estimates the additional response to a change in the progressivity of the personal tax schedule. The former produces the three broad entries of the revenue gradient, while the latter produces its progressivity entry. Appendix B describes the broad tax data, and Appendix B.2 describes the annual personal tax data and the construction of the progressivity direction.

C.1 Estimator and Posterior Draws

Both blocks use the two-stage Bayesian local projection procedure of Cloyne et al. (2022). Let $\mathbf{Z}_t$ denote the relevant state vector and let $\mathbf{C}_t$ contain any contemporaneous controls. The first stage estimates

$$\mathbf{Z}_t = \mathbf{c} + \sum_{\ell=1}^p A_{\ell} \mathbf{Z}_{t-\ell} + \Theta \mathbf{C}_t + \mathbf{u}_t, \quad \mathbf{u}_t \sim N(0, \Sigma). \quad (\text{A.35})$$

The quarterly block uses four lags and contains no contemporaneous controls. The annual block uses one lag, and $\mathbf{C}_t$ contains the annualized corporate and tariff narrative measures.

For each outcome v and economic response horizon $h \geq 1$, the second stage estimates a local projection of $Z_{v,t+h}$ on the impact-period state $\mathbf{Z}_t$ and its lags. Suppressing the annual controls for simplicity, write

$$Z_{v,t+h} = c_v^{(h)} + \boldsymbol{\psi}_{v,h}^\top \mathbf{Z}_t + \sum_{\ell=1}^{p-1} \mathbf{b}_{v,\ell}^{(h)\top} \mathbf{Z}_{t-\ell} + \eta_{v,t+h}^{(h)}. \quad (\text{A.36})$$

Stacking $\boldsymbol{\psi}_{v,h}^\top$ across outcomes gives the matrix Ψ_h . For a given posterior draw and impact vector $\mathbf{a}_s$, the response path is

$$\mathbf{IRF}_{0,s} = \mathbf{a}_s, \quad \mathbf{IRF}_{h,s} = \Psi_{h,s} \mathbf{a}_s, \quad h \geq 1. \quad (\text{A.37})$$

Thus, the impact response comes from the first stage and each subsequent response comes from the corresponding local projection. The same local projection coefficient draw is applied to every impact vector within a posterior draw. This preserves the dependence among the response columns and ensures that linear transformations of the impact basis produce the same transformations of the impulse responses.

The quarterly block uses $h = 0, \dots, 19$, while the annual block uses $h = 0, \dots, 4$. These windows contain twenty quarterly observations and five annual observations, respectively. The quarterly discount factor is $\beta = 0.9926$, and the annual discount factor is $\beta_A = \beta^4$. All response matrices below are constructed draw by draw.

Bayesian estimation and computation. I use the hierarchical Minnesota specification of Giannone, Lenza, and Primiceri (2015) in the first-stage BVAR. The prior centers each first own lag at one and all remaining lag coefficients at zero. Lag variances decay quadratically, the prior variance of the constant is 10^7 , and the lag-decay parameter is fixed at two. The gamma hyperprior for overall tightness has mode 0.2, standard deviation 0.4, and support $[10^{-4}, 5]$. The variable-specific innovation scales are calibrated from univariate autoregressions. I also use the sum-of-coefficients prior of Doan, Litterman, and Sims (1984) and the single-unit-root dummy-observation prior of Sims (1993). Their tightness hyperpriors each have mode and standard deviation equal to one and support $[10^{-4}, 50]$. Overall tightness and the two dummy-prior tightness parameters are sampled by Metropolis–Hastings; conditional on them, the VAR coefficients and Σ are drawn from their conjugate posterior. In the annual block, the coefficients on the contemporaneous narrative controls have zero prior mean and variance 10^7 .

Each first-stage chain runs for 6,700 iterations. I discard the first 200 iterations and retain every fifth draw, producing 1,300 reduced-form posterior draws. After constructing the relevant impact basis, I retain 1,000 valid draws without replacement.

The local projections are estimated separately for every outcome and horizon. For a given outcome, the Minnesota prior mean places its preliminary univariate AR(1) coefficient on the first own lag and zero on all other lag coefficients. If s_y and s_j are the sample standard deviations of the outcome and regressor j , respectively, the prior variance on lag ℓ is

$$\left(\frac{s_y}{s_j} \right)^2 \frac{10^2}{\ell^2},$$

and the prior variance on the constant is 100. In the annual block, the coefficients on the contemporaneous narrative controls have zero prior mean and variance 10^7 .

I allow for heavy-tailed projection errors through the scale-mixture representation

$$\eta_{v,t+h}^{(h)} \mid v_{v,t,h} \sim N\left(0, \frac{\sigma_{v,h}^2}{v_{v,t,h}}\right), \quad v_{v,t,h} \sim \text{Gamma}\left(\frac{\nu_{v,h}}{2}, \frac{\nu_{v,h}}{2}\right).$$

Integrating out $v_{v,t,h}$ gives a Student- t error with $v_{v,h}$ degrees of freedom. I use $p(v_{v,h}) \propto \exp(-v_{v,h}/10)$ for $v_{v,h} > 2$. Conditional on the remaining parameters, the latent precisions, regression coefficients, and error variance have standard Gamma, Normal, and inverse-Gamma updates. I update $v_{v,h}$ using the random-walk proposal $v'_{v,h} = v_{v,h} + \xi$, with $\xi \sim N(0, 1)$, rejecting proposals at or below two.

Each second-stage chain runs for 5,200 iterations. I discard the first 200 iterations and retain every fifth draw, leaving 1,000 draws. Within each block, every retained local projection draw is paired with one distinct first-stage impact-basis draw.

C.2 Quarterly Broad Tax Block

The quarterly state vector is

$$\mathbf{z}_t^Q = (\log(1 - \tau_{P,t}), \log(1 - \tau_{C,t}), -\log(1 + \tau_{T,t}), \log B_{P,t}, \log B_{C,t}, \log B_{T,t})^\top.$$

The sample covers 1947Q1–2019Q4. Let $\mathbf{z}_t^Q$ collect the three quarterly narrative measures. Writing $\boldsymbol{\epsilon}_{\text{tax},t}$ for the three tax innovations and $\boldsymbol{\epsilon}_{\text{other},t}$ for all remaining innovations, the proxy restrictions are

$$\text{rank}\left(E[\mathbf{z}_t^Q \boldsymbol{\epsilon}_{\text{tax},t}^\top]\right) = 3, \quad E[\mathbf{z}_t^Q \boldsymbol{\epsilon}_{\text{other},t}^\top] = 0. \quad (\text{A.38})$$

For each BVAR draw, these restrictions and the reduced-form covariance matrix recover the three-dimensional tax policy impact space. I apply the covariance algebra in Mertens and Ravn (2013a), generalized to three narrative measures, and select a basis using one common Cholesky ordering: personal taxes, corporate taxes, and tariffs. Let $A_{0,s}^Q$ denote the resulting 6×3 impact matrix in draw s .

I divide each column by the absolute value of the impact response of its own transformed tax rate. This fixes an otherwise arbitrary scale but does not imply that a named tax shock moves only its corresponding rate. I therefore retain the response of every broad rate and every tax base to all three impact directions. Because the complete rate response matrix is inverted below, the resulting semi-elasticities do not depend on the Cholesky ordering or the column normalizations.

Let $D_{h,s}^Q$ denote the matrix of responses of the three transformed tax rates at horizon h in draw s , and let $Y_{h,s}^Q$ denote the corresponding matrix of log tax base responses, ordered as personal, corporate, and tariff. Their discounted sums are

$$D_s^Q = \sum_{h=0}^{19} \beta^h D_{h,s}^Q, \quad Y_s^Q = \sum_{h=0}^{19} \beta^h Y_{h,s}^Q. \quad (\text{A.39})$$

At accounting date t , the corresponding responses of the ordinary tax rates are

$$P_{s,t}^Q = J_t^Q D_s^Q, \quad J_t^Q = -\text{diag}(1 - \tau_{P,t}, 1 - \tau_{C,t}, 1 + \tau_{T,t}). \quad (\text{A.40})$$

The matrix of present-value semi-elasticities with respect to the ordinary tax rates is

$$E_{s,t}^Q = Y_s^Q \left(P_{s,t}^Q \right)^{-1}. \quad (\text{A.41})$$

The j th column of equation (A.41) gives the response of each tax base to a change in broad rate j while setting the discounted changes in the other two broad rates equal to zero. This is why the calculation uses the complete response matrices rather than treating each narrative measure as an experiment that changes only one rate.

For $i, j \in \{P, C, T\}$, the resulting broad revenue derivative is

$$r_{j,s,t} = \tilde{B}_{j,t} + \sum_{i \in \{P,C,T\}} \tau_{i,t} \tilde{B}_{i,t} \left(E_{s,t}^Q \right)_{ij}. \quad (\text{A.42})$$

The rates and bases in this expression are evaluated at the accounting date. The common present-value factor contained in the discounted bases cancels from the alignment measures. The static semi-elasticity tables evaluate equation (A.41) at the sample-average tax rates, while the revenue gradients and alignment series use the rate prevailing at each accounting date.

C.3 Annual Progressivity Block

The annual state vector is

$$\mathbf{Z}_t^A = (\log(1 - \tau_{L,t}), \log(1 - \tau_{H,t}), \log B_{L,t}, \log B_{H,t}, \log B_{C,t}, \log B_{T,t})^\top,$$

where L denotes the lower 99 percent of the income distribution and H denotes the top 1 percent. The personal tax rates and income measures are those constructed by Mertens and Montiel Olea (2018) and extended through 2023 as described in Appendix B.2. The annualized corporate and tariff narrative measures enter both estimation stages as contemporaneous controls.

Let $\mathbf{z}_t^A = (z_{L,t}, z_{H,t})^\top$ collect the two personal narrative measures. The top-1 measure is the assembled MMO top-1 series. The lower-99 measure aggregates the three lower MMO brackets using the same income weights employed in the data extension. Both measures include the recovered 2018 Tax Cuts and Jobs Act observation, so the annual analysis uses nine reform years.

Because corporate and tariff reforms may occur in the same year as personal tax reforms, I first residualize the two personal narrative measures on a constant and the contemporaneous corporate and tariff narrative measures. Let $\tilde{\mathbf{z}}^A$ denote the resulting $N_A \times 2$ matrix. For BVAR draw s , let U_s^A denote the $N_A \times 6$ matrix of reduced-form innovations after the annual lag. I construct

$$\Gamma_s^A = \frac{1}{N_A} \left(U_s^A \right)^\top \tilde{\mathbf{z}}^A, \quad N_A = 76. \quad (\text{A.43})$$

To see the role of this matrix, let $A_{0,s}^A$ contain the two structural personal tax impact vectors and let

$$\Phi_A = E[\tilde{\mathbf{z}}_t^A \boldsymbol{\epsilon}_{\text{tax},t}^\top].$$

Under the proxy exclusion restriction,

$$E[\mathbf{u}_t \tilde{\mathbf{z}}_t^{A\top}] = A_{0,s}^A \Phi_A^\top.$$

If Φ_A has rank two, the columns of Γ_s^A span the same personal tax policy impact space as the structural impact vectors. I retain draws in which the first two rows of Γ_s^A have rank two.

Together with the reduced-form covariance matrix, these restrictions recover the two-dimensional personal tax impact space up to an orthogonal rotation. For each retained draw, I select a particular basis using a Cholesky decomposition ordered lower-99 first and top-1 second, and normalize each column by the absolute impact response of its own tax rate. I then propagate the two columns of $A_{0,s}^A$ through the annual local projections. This choice is needed to label the intermediate impulse responses, but it has no effect on the final semi-elasticity matrix. Any

common nonsingular change of basis applied to both the tax rate and tax base responses cancels when the complete 2×2 rate response matrix is inverted below.

Let $D_{h,s}^A$ denote the 2×2 matrix of responses of the two log retention rates at horizon h , and let $Y_{h,s}^A$ denote the corresponding 4×2 matrix of responses of the lower-99 income base, the top-1 income base, the corporate base, and the tariff base. Their five-year discounted sums are

$$D_s^A = \sum_{h=0}^4 \beta_A^h D_{h,s}^A, \quad Y_s^A = \sum_{h=0}^4 \beta_A^h Y_{h,s}^A. \quad (\text{A.44})$$

At accounting date t , I convert the log retention-rate responses into ordinary tax rate responses:

$$P_{s,t}^A = J_t^A D_s^A, \quad J_t^A = -\text{diag}(1 - \tau_{L,t}, 1 - \tau_{H,t}). \quad (\text{A.45})$$

The rates in J_t^A are the same lower-99 and top-1 marginal tax rates used to construct the state vector.

The matrix of present-value semi-elasticities with respect to the two ordinary tax rates is

$$E_{s,t}^A = Y_s^A \left(P_{s,t}^A \right)^{-1}. \quad (\text{A.46})$$

Its first column gives the base responses to a change in the lower-99 rate holding the top-1 rate fixed, while its second column gives the responses to a change in the top-1 rate holding the lower-99 rate fixed. For any nonsingular transformation Q of the impact basis,

$$Y_s^A Q \left(P_{s,t}^A Q \right)^{-1} = Y_s^A \left(P_{s,t}^A \right)^{-1}.$$

The complete rate response matrix therefore removes the particular basis and normalization of the two initial policy directions.

Let $b_{L,t}$ and $b_{H,t}$ denote the mechanical tax exposures of the lower 99 percent and top 1 percent, as constructed in Appendix B.2, and define

$$s_{H,t} = \frac{b_{H,t}}{b_{L,t} + b_{H,t}}.$$

A one-percentage-point increase in progressivity changes the two rates in the direction

$$\mathbf{h}_t = \begin{pmatrix} -s_{H,t} \\ 1 - s_{H,t} \end{pmatrix}. \quad (\text{A.47})$$

The top-1-minus-lower-99 rate gap rises by one percentage point, while initial personal tax payments remain fixed:

$$h_{H,t} - h_{L,t} = 1, \quad b_{L,t} h_{L,t} + b_{H,t} h_{H,t} = 0.$$

Rates are stored as fractions, so a one-percentage-point reform moves the two rates by $0.01 \mathbf{h}_t$.

The present-value semi-elasticities with respect to progressivity are

$$\boldsymbol{\varepsilon}_{d,s,t} = E_{s,t}^A \mathbf{h}_t = \begin{pmatrix} \varepsilon_{Ld,s,t} \\ \varepsilon_{Hd,s,t} \\ \varepsilon_{Cd,s,t} \\ \varepsilon_{Td,s,t} \end{pmatrix}. \quad (\text{A.48})$$

These coefficients vary by accounting date because both the conversion from log retention rates and the progressivity direction depend on the contemporaneous tax rates and exposure shares. The headline semi-elasticity table first averages each coefficient over 1948–2023 within a posterior draw,

$$\bar{\epsilon}_{d,s} = \frac{1}{76} \sum_{t=1948}^{2023} \epsilon_{d,s,t}, \quad (\text{A.49})$$

and then reports the median and credible interval of these draw-level averages.

Because the reform leaves initial personal tax payments unchanged, its revenue-gradient entry contains no mechanical term:

$$r_{d,s,t} = \tau_{L,t} \widetilde{B}_{L,t} \epsilon_{Ld,s,t} + \tau_{H,t} \widetilde{B}_{H,t} \epsilon_{Hd,s,t} + \tau_{C,t} \widetilde{B}_{C,t} \epsilon_{Cd,s,t} + \tau_{T,t} \widetilde{B}_{T,t} \epsilon_{Td,s,t}. \quad (\text{A.50})$$

Let M_t denote the policy metric for the two personal tax rates and let

$$m_{d,t} = \mathbf{h}_t^\top M_t \mathbf{h}_t$$

denote the squared length of the progressivity direction. The construction of M_t is described in Appendix B.2. The response matrices above use the computational order personal, corporate, and tariff. When constructing alignment, I permute the broad entries into the order used in the main text, corporate, personal, tariff, and progressivity. Under equal incidence, the expanded gradients are therefore

$$\mathbf{w}_t = \begin{pmatrix} \widetilde{B}_{C,t} \\ \widetilde{B}_{P,t} \\ \widetilde{B}_{T,t} \\ 0 \end{pmatrix}, \quad \mathbf{r}_{sA,sQ,t} = \begin{pmatrix} r_{C,sQ,t} \\ r_{P,sQ,t} \\ r_{T,sQ,t} \\ r_{d,sA,t} / \sqrt{m_{d,t}} \end{pmatrix}. \quad (\text{A.51})$$

C.4 Combining the Two Blocks

The quarterly and annual blocks are estimated separately. I retain 1,000 posterior draws from each and evaluate the expanded gradient, cosine, and sine for every annual–quarterly draw pair at each accounting date. The reported intervals therefore use the complete $1,000 \times 1,000$ Cartesian product of the two sets of draws. This preserves posterior uncertainty within each empirical block but treats their marginal posteriors as independent; it does not estimate posterior dependence between the annual and quarterly systems.

D Robustness of the Broad Tax Elasticity Estimates

This appendix reports robustness of the present-value elasticities from for the broad estimates to alternative estimation methods, accounting for S-corporations, adding in macro controls, and accounting for if the VAR may not be invertible.

1. **Frequentist local projections.** Standard local projections using the Mertens-Ravn structural shocks identified from the BVAR as regressors. Uncertainty about A_0 is propagated using the wild bootstrap proxy SVAR of Mertens and Ravn (2013a).

2. **Smooth local projections.** Frequentist local projections with a penalty on deviations across adjacent horizons following Barnichon and Brownlees (2019), which reduces the lumpiness of standard LP estimates at the cost of imposing smoothness.
3. **Bias-corrected local projections.** Local projections with the bias correction of Herbst and Johansson (2024), which adjusts for the finite-sample bias that arises from projecting on persistent regressors.
4. **Bayesian VAR.** A six-variable BVAR with four lags and hierarchical Minnesota priors following Giannone, Lenza, and Primiceri (2015). Impulse responses are computed by propagating shocks through the VAR companion form rather than estimating separate regressions at each horizon.
5. **Proxy SVAR.** A frequentist structural VAR using the narrative measures as external instruments following Mertens and Ravn (2013a), extended to 2019Q4. Confidence intervals are constructed using the wild bootstrap.
6. **Linear trends.** Each of the above methods (where applicable) is re-estimated with a linear trend added to the specification. Results are reported in the lower panel of Table D.1.
7. **LP-IV.** At each horizon h , the log retention rate is instrumented directly with the narrative measure, controlling for four lags of all variables. This specification is robust to non-invertibility of the VAR representation by construction. Three variants are reported in Table D.5: plain LP-IV, smoothed LP-IV following Barnichon and Brownlees (2019), and bias-corrected LP-IV following Herbst and Johansson (2024). All produce elasticities quantitatively similar to the baseline.
8. **TVP-BVAR.** A four-variable time-varying parameter VAR with stochastic volatility following Primiceri (2005), which allows the full transmission mechanism to drift over time.

Table D.1: Present-Value Elasticities of Tax Bases under Alternative Methods

Method	τ_P			τ_C			τ_T		
	ε_{PP}	ε_{CP}	ε_{TP}	ε_{PC}	ε_{CC}	ε_{TC}	ε_{PT}	ε_{CT}	ε_{TT}
BLP	-0.37	-10.55	0.16	0.41	-3.40	-0.58	0.83	-4.99	-14.26
	(-0.99, 0.28)	(-13.36, -8.20)	(-1.86, 2.07)	(0.14, 0.71)	(-4.45, -2.45)	(-1.41, 0.19)	(-0.09, 1.86)	(-8.75, -1.72)	(-17.32, -11.91)
LP	-0.15	-11.70	0.31	0.57	-3.80	-1.50	1.51	-5.69	-18.11
	(-1.69, 1.16)	(-21.79, -6.88)	(-5.05, 3.01)	(-0.06, 1.30)	(-7.54, -1.63)	(-3.10, -0.33)	(-0.60, 3.55)	(-13.93, 2.25)	(-23.05, -13.26)
SLP	-0.15	-11.70	0.31	0.57	-3.80	-1.51	1.51	-5.70	-18.11
	(-1.69, 1.16)	(-21.81, -6.89)	(-5.06, 3.01)	(-0.06, 1.30)	(-7.55, -1.63)	(-3.10, -0.33)	(-0.60, 3.55)	(-13.94, 2.24)	(-23.05, -13.26)
Bias-corrected LP	0.08	-10.42	2.55	0.89	-3.07	-0.82	2.37	3.15	-20.30
	(-2.15, 1.95)	(-24.99, -3.50)	(-2.16, 5.23)	(-0.20, 2.45)	(-13.91, 0.63)	(-2.73, 2.08)	(-3.47, 10.28)	(-16.20, 49.77)	(-32.62, -12.11)
BVAR	-0.43	-4.81	-0.91	0.61	-3.98	-0.08	-0.31	-7.01	-14.98
	(-1.06, 0.30)	(-11.72, -0.00)	(-5.50, 1.98)	(0.37, 0.92)	(-6.87, -2.33)	(-1.99, 1.12)	(-1.94, 1.42)	(-20.16, 0.93)	(-25.73, -9.19)
SVAR-IV	-0.84	-17.88	0.00	0.80	-4.90	0.97	1.37	-15.11	-10.72
	(-9.96, 5.51)	(-61.68, 4.79)	(-19.03, 13.72)	(-6.95, 9.62)	(-44.22, 18.90)	(-16.86, 17.56)	(-22.82, 31.00)	(-125.65, 61.70)	(-55.75, 40.38)

Note: Each semi-elasticity ε_{ij} is the cumulative discounted percent change in base i per percentage-point increase in rate j , holding the other two rates fixed. Within each draw the rate-response matrix D and base-response matrix Y are mapped into $E = YD^{-1}$, an elasticity with respect to the net-of-tax coordinate, and converted to the tax rate by $-E_{ij}/(1 - \tau_j)$ for the income taxes and $-E_{ij}/(1 + \tau_T)$ for tariffs, evaluated at sample-average tax rates. Parentheses report 90% intervals. LP, SLP, and Bias-corrected LP retain the BLP contemporaneous rate map and alter only propagation; BVAR and SVAR-IV re-estimate both maps. Bayesian rows report credible intervals, the remaining rows hybrid or bootstrap intervals.

Table D.2: Present-Value Elasticities of Tax Bases under Macro Controls

	τ_P	τ_C	τ_T
ε_P . (Personal Base)	-0.12 (-1.92, 0.82)	-0.71 (-1.77, -0.14)	-0.27 (-3.67, 1.53)
ε_C . (Corporate Base)	-10.74 (-16.70, -6.98)	-5.61 (-9.77, -3.59)	-2.94 (-14.96, 3.89)
ε_T . (Tariff Base)	-1.22 (-5.09, 1.23)	-1.61 (-4.04, -0.36)	-14.12 (-21.83, -10.02)

Note: Each semi-elasticity ε_{ij} is the cumulative discounted percent change in base i per percentage-point increase in rate j . Within each draw the rate-response matrix D and base-response matrix Y are mapped into $E = YD^{-1}$, so each column holds the other rates fixed, and the result is converted from the net-of-tax coordinate to the tax rate at sample-average tax rates. Macro Controls adds government spending, output, and debt to the three rates and bases. Estimated by Bayesian local projection; parentheses report 90% credible intervals.

Table D.3: Present-Value Elasticities under S-Corporation Reallocation

	τ_P	τ_C	τ_T
ε_P . (Personal Base)	-0.78 (-1.59, -0.06)	0.55 (0.32, 0.81)	0.89 (0.15, 1.80)
ε_C . (Corporate Base)	-10.21 (-14.16, -7.38)	-4.20 (-5.39, -3.27)	-6.25 (-9.78, -3.34)
ε_T . (Tariff Base)	-0.71 (-3.46, 1.54)	-0.94 (-1.77, -0.19)	-15.62 (-18.25, -13.30)

Note: Each semi-elasticity ε_{ij} is the cumulative discounted percent change in base i per percentage-point increase in rate j , holding the other two rates fixed. Within each draw the rate-response matrix D and base-response matrix Y are mapped into $E = YD^{-1}$, an elasticity with respect to the net-of-tax coordinate, and converted to the tax rate by $-E_{ij}/(1 - \tau_j)$ for the income taxes and $-E_{ij}/(1 + \tau_T)$ for tariffs, evaluated at sample-average tax rates. Parentheses report 90% intervals. S-corporation income is reallocated between the personal and corporate bases, and the rates and narrative measures are reconstructed to match.

Table D.4: Present-Value Elasticities of Tax Bases under Alternative Methods (PIT–CIT System)

Method	τ_C		τ_P	
	ε_{CC}	ε_{PC}	ε_{CP}	ε_{PP}
BLP	-3.28 (-4.49, -2.48)	0.21 (-0.04, 0.44)	-11.77 (-14.73, -9.59)	-0.87 (-1.54, -0.33)
LP	-3.66 (-8.74, -1.66)	0.31 (-0.33, 1.06)	-15.82 (-27.76, -10.73)	-0.70 (-2.79, 0.89)
SLP	-3.67 (-8.76, -1.67)	0.31 (-0.33, 1.06)	-15.86 (-27.83, -10.74)	-0.70 (-2.79, 0.89)
Bias-corrected LP	-3.27 (-11.20, -1.00)	0.53 (-0.40, 1.58)	-17.04 (-33.92, -11.36)	-0.61 (-3.33, 1.86)
BVAR	-5.22 (-7.94, -3.15)	0.64 (0.37, 0.98)	-12.81 (-21.24, -5.45)	-0.65 (-1.59, 0.22)
SVAR-IV	-3.90 (-24.47, 1.94)	0.66 (-2.13, 4.98)	-21.43 (-64.31, -11.11)	-1.01 (-7.58, 4.03)


Note: Each semi-elasticity is the cumulative discounted percent change in base i per percentage-point increase in the named rate, holding the other broad rate fixed, computed within each draw as an element of $E = YD^{-1}$ for the PIT–CIT system and converted from the net-of-tax coordinate at sample-average tax rates. LP, SLP, and Bias-corrected LP retain the BLP contemporaneous impact basis and re-estimate only propagation; BVAR and SVAR-IV re-estimate both maps. Parentheses report 90% intervals.

Table D.5: Present-Value Elasticities under Direct LP-IV (Three-Instrument System)

Method	τ_P			τ_C			τ_T		
	ε_{PP}	ε_{CP}	ε_{TP}	ε_{PC}	ε_{CC}	ε_{TC}	ε_{PT}	ε_{CT}	ε_{TT}
LP-IV	0.94	-7.68	-2.19	0.60	-3.18	-1.66	2.25	-10.28	-7.79
	(-5.20, 8.78)	(-33.79, 9.20)	(-22.47, 9.99)	(-2.32, 5.21)	(-16.18, 5.83)	(-10.71, 6.88)	(-11.23, 21.50)	(-73.74, 45.76)	(-51.11, 42.26)
Smooth LP-IV	0.94	-7.68	-2.19	0.60	-3.18	-1.66	2.25	-10.28	-7.79
	(-5.20, 8.78)	(-33.79, 9.20)	(-22.48, 9.99)	(-2.32, 5.21)	(-16.18, 5.83)	(-10.70, 6.87)	(-11.23, 21.50)	(-73.73, 45.76)	(-51.10, 42.26)
Bias-corrected LP-IV	0.81	-6.72	-1.75	0.61	-3.05	-1.38	2.28	-9.07	-6.92
	(-4.23, 7.80)	(-28.01, 8.66)	(-19.11, 10.10)	(-2.33, 4.77)	(-15.59, 5.27)	(-9.26, 6.98)	(-12.77, 22.43)	(-72.14, 42.13)	(-50.26, 48.79)


Note: Each semi-elasticity ε_{ij} is the cumulative discounted percent change in base i per percentage-point increase in rate j , holding the other two rates fixed. Within each draw the rate-response matrix D and base-response matrix Y are mapped into $E = YD^{-1}$, an elasticity with respect to the net-of-tax coordinate, and converted to the tax rate by $-E_{ij}/(1 - \tau_j)$ for the income taxes and $-E_{ij}/(1 + \tau_T)$ for tariffs, evaluated at sample-average tax rates. Parentheses report 90% intervals. The LP-IV rows identify the rate and base responses directly with the narrative measures before forming YD^{-1} , unlike the LP rows of Table D.1, which retain the BLP contemporaneous rate map. Intervals are bootstrap intervals.

Figure D.1: Time-Varying PIT-CIT-Tariff Elasticities



Note: The figure reports time-varying present-value elasticities in the PIT-CIT-tariff system.

Figure D.2: Tax Elasticity Matrix Leave-One-Out



Note: This figure shows the how the elasticity matrix changes leaving out one reform at a time.

D.1 Rationalizing the Magnitude of ε_{CL}

I find across a variety of methods that the elasticity of the corporate tax base with respect to personal taxes is the largest entry in the elasticity matrix. The main text rationalized that with operating leverage, and this appendix expands on that.

Corporate output—or corporate gross value added—decomposes as

$$Y = wL + \pi + F,$$

where wL is the wage bill, π is pre-tax corporate profits, and F is the remainder of costs, including interest payments, taxes, depreciation, and so on. For convenience, define the labor share as $\gamma = wL/Y$ and the profit share as $s_\pi = \pi/Y$. When output changes by dY and F does not adjust, the change in profits is $d\pi = dY - d(wL)$. If wages are fully flexible and the labor share is roughly constant at the margin, then $d(wL) = \gamma dY$ and

$$d \log \pi = \frac{1 - \gamma}{s_\pi} d \log Y.$$

The ratio $(1 - \gamma)/s_\pi$ is the *frictionless* operating leverage. It bounds leverage from below because it assumes that wages fully absorb their full share of output movements within the relevant horizon. If wages are sticky, then labor compensation adjusts by less than γdY in the short run. As a result, a larger share of any output gain initially accrues

to profits. This raises effective leverage above $(1 - \gamma)/s_\pi$ and simultaneously attenuates the passthrough of output to the personal tax base below one. The same friction that makes the cross-elasticity larger than the frictionless prediction makes the personal base elasticities smaller. It exceeds one whenever $F > 0$ since non-labor income must cover both fixed costs F and profits π , so the percentage change in profits exceeds the percentage change in output by a factor of $(1 - \gamma)/s_\pi$.¹⁷ This is a lower bound on effective leverage because it assumes the labor bill freely adjusts. Unlike profits, wages have no leverage. When compensation adjusts freely, $d \log(wL) = d \log Y$, so a percent increase in output fully passes through to the personal tax base. When wages are sticky, the passthrough is attenuated further below one. In other words, the same friction that raises leverage on the corporate tax base lowers passthrough on the personal income tax base.

Tax Cuts and the Response of Output. The leverage result shows that the elasticity of the corporate base with respect to any shock ε_{CX} decomposes into the response of output to that shock, amplified by the factor $(1 - \gamma)/s_\pi$:

$$\varepsilon_{CL} = \frac{1 - \gamma}{s_\pi} \times \varepsilon_{YL} \quad \text{and} \quad \varepsilon_{CC} = \frac{1 - \gamma}{s_\pi} \times \varepsilon_{YC}$$

where ε_{YL} is the elasticity of output with respect to the log personal retention rate. To evaluate the contribution of leverage to the response of the corporate base, we need to know both ε_{YL} and ε_{YC} . These are not directly estimated in the baseline specification, but I extend them to include GDP, government dept, and government spending. That allows direct estimation of both ε_{YL} and ε_{YC} . Across estimation methodologies, I find that both ε_{YL} and ε_{YC} are robustly positive, and that the personal response is the larger of the two. Under the Bayesian local projection, $\varepsilon_{YL} = 0.96$ and $\varepsilon_{YC} = 0.85$. The ordering is intuitive. Personal tax cuts raise labor supply through a higher after-tax wage, which immediately raises output. On the other hand, corporate tax cuts should have gradual effects. They operate through the user cost of capital, and so any corporate tax cut would presumably have dynamic effects through capital accumulation. As a result, the output response to a corporate tax cut should be smaller and more back-loaded.

Since leverage applies equally to both shocks, the full elasticity matrix is determined by two objects: the leverage factor $(1 - \gamma)/s_\pi$ and the output elasticities. For the corporate base, the predictions are given above. For the personal base, there is no leverage since wages should scale appropriately with output, so

$$\varepsilon_{LL} \approx \varepsilon_{YL} \quad \text{and} \quad \varepsilon_{LC} \approx \varepsilon_{YC}.$$


Both γ and s_π are directly observable from NIPA Table 1.14 (Gross Value Added of Domestic Corporate Business): γ is compensation of employees divided by GVA, and s_π is corporate profits with inventory valuation and capital consumption adjustments divided by GVA. Figure D.3 plots the leverage series over time. It ranges from approximately 1.9 to 3.4, with a postwar mean of 2.59; the empirical leverage from the extended specification ($\varepsilon_{CL}/\varepsilon_{YL} = 9.3$) lies above it at every date, consistent with wage stickiness raising effective leverage above the frictionless benchmark. Leverage was lowest in the early postwar period and peaked in the 1980s. Table 2 compares the frictionless predictions to the data, using the postwar mean leverage of 2.59.

The frictionless model matches the qualitative ordering $\varepsilon_{CL} > \varepsilon_{CC} > \varepsilon_{LL} > \varepsilon_{LC}$ and delivers a cross-elasticity of the right sign and order of magnitude, $2.59 \times 0.96 \approx 2.5$ against a baseline estimate of 8.89, using only national ac-

17. All that we require for the argument to work is that some portion of the cost structure is predetermined. This is almost trivially satisfied. Capital is subject to fixed adjustment costs, with the majority of firms holding their capital stock fixed in any given quarter (Winberry 2021; Koby and Wolf 2020). Employment adjusts in infrequent, discrete changes rather than continuously (Caballero, Engel, and Haltiwanger 1997). Leases, debt service, and administrative overhead are fixed by construction.

counts arithmetic and the output elasticities from the extended specification. In principle, the frictionless prediction is a floor on this channel rather than a point prediction because, as already discussed earlier, it assumes wages are not sticky.

Figure D.3: Leverage in the National Income and Product Accounts



Note: Leverage is computed using Table 1.14 from the National Income and Product Accounts. I compute the compensation share γ as compensation of employees (line 7) divided by gross value added of corporate business (line 1), and the profit share s_π as corporate profits with inventory valuation and capital consumption adjustments (line 14) divided by gross value added of corporate business. Raw leverage is $(1 - \gamma)/s_\pi$. I remove fluctuations at business cycle frequencies (2–40 quarters) using the Christiano and Fitzgerald (2003) filter and plot the resulting trend.

E Robustness of the Progressivity Semi-Elasticities

Table E.1: Semi-Elasticities under Alternative System Methods

	Bottom 99	Top 1	Corporate base	Tariff base
BLP	-0.15 (-0.65, 0.42)	-2.28 (-3.80, -1.08)	4.12 (1.83, 12.63)	7.97 (4.64, 21.37)
LP	-0.08 (-0.67, 1.17)	-1.65 (-3.10, 2.45)	3.75 (1.69, 10.49)	7.78 (3.53, 25.16)
Smooth LP	-0.08 (-0.67, 1.17)	-1.65 (-3.10, 2.45)	3.75 (1.69, 10.49)	7.78 (3.53, 25.16)
Bias-corrected LP	0.99 (-10.07, 10.19)	1.03 (-29.02, 29.80)	4.61 (-28.41, 38.49)	10.03 (-53.84, 74.01)
BVAR	-0.07 (-0.29, 0.20)	-2.60 (-3.31, -2.11)	-0.21 (-1.15, 0.54)	1.98 (1.09, 3.03)
SVAR-IV	-0.11 (-9.10, 9.90)	-2.16 (-24.17, 18.82)	5.88 (-64.08, 66.41)	6.19 (-88.67, 84.47)

Note: Posterior medians with 90-percent credible intervals in parentheses. Entries are the percentage change in the discounted five-year response of each tax base to a one-percentage-point increase in the top-1-minus-bottom-99 average marginal tax-rate gap, holding initial personal tax payments fixed, evaluated draw by draw at each year's observed rates and exposure weights and averaged across accounting years. The personal block has two groups, identified by two narrative instruments over nine reform years, conditioning on the annual corporate and tariff narrative measures. The LP rows hold the first-stage impact basis fixed and vary only propagation. SVAR-IV is a separate VARX identified by the Mertens–Montiel Olea triangular restriction, with a wild bootstrap.

Table E.2: Semi-Elasticities under Direct LP-IV

	Bottom 99	Top 1	Corporate base	Tariff base
BLP	-0.15 (-0.65, 0.42)	-2.28 (-3.80, -1.08)	4.12 (1.83, 12.63)	7.97 (4.64, 21.37)
LP-IV	0.80 (-4.67, 7.54)	-1.47 (-14.60, 14.69)	-2.01 (-45.68, 77.77)	-5.65 (-52.46, 60.86)
Smooth LP-IV	0.80 (-4.67, 7.54)	-1.47 (-14.60, 14.69)	-2.01 (-45.68, 77.77)	-5.65 (-52.47, 60.86)
Bias-corrected LP-IV	0.91 (-4.07, 6.18)	-1.15 (-15.67, 10.54)	-1.63 (-54.61, 54.81)	-5.66 (-66.63, 57.79)

Note: Posterior medians with 90-percent credible intervals in parentheses. Entries are the percentage change in the discounted five-year response of each tax base to a one-percentage-point increase in the top-1-minus-bottom-99 average marginal tax-rate gap, holding initial personal tax payments fixed, evaluated draw by draw at each year's observed rates and exposure weights and averaged across accounting years. The personal block has two groups, identified by two narrative instruments over nine reform years, conditioning on the annual corporate and tariff narrative measures. Each row instruments the two personal marginal-rate equations directly with the narrative measures; no first-stage impact basis is reused.

Table E.3: Semi-Elasticities under Alternative System and Deterministic Specifications

	Bottom 99	Top 1	Corporate base	Tariff base
BLP	-0.15 (-0.65, 0.42)	-2.28 (-3.80, -1.08)	4.12 (1.83, 12.63)	7.97 (4.64, 21.37)
Expanded macro system	-0.07 (-0.36, 0.21)	-2.07 (-2.84, -1.51)	2.22 (1.28, 3.73)	4.37 (3.06, 6.81)
MMO trends and dummies	-0.54 (-3.37, 1.69)	-2.30 (-7.45, 2.93)	2.31 (-10.47, 21.50)	9.30 (-38.34, 70.11)
Macro + MMO trends	-0.24 (-1.75, 0.56)	-0.65 (-2.61, 1.82)	2.13 (-0.23, 8.24)	7.86 (3.96, 28.89)
Two annual lags	-0.22 (-6.45, 6.76)	-2.45 (-12.95, 8.10)	2.90 (-34.59, 39.37)	9.66 (-96.41, 126.31)


Note: Posterior medians with 90-percent credible intervals in parentheses. Entries are the percentage change in the discounted five-year response of each tax base to a one-percentage-point increase in the top-1-minus-bottom-99 average marginal tax-rate gap, holding initial personal tax payments fixed, evaluated draw by draw at each year's observed rates and exposure weights and averaged across accounting years. The personal block has two groups, identified by two narrative instruments over nine reform years, conditioning on the annual corporate and tariff narrative measures. The macro system adds government spending, GDP, and federal debt to the endogenous state. The MMO rows add linear and quadratic trends and 1949 and 2008 dummies in both stages.

Table E.4: Semi-Elasticities under S-Corporation Reallocation

	Bottom 99	Top 1	Corporate base	Tariff base
BLP	-0.15 (-0.65, 0.42)	-2.28 (-3.80, -1.08)	4.12 (1.83, 12.63)	7.97 (4.64, 21.37)
S-corporation	0.62 (0.09, 2.02)	-1.37 (-2.82, 0.75)	7.47 (4.04, 19.96)	7.56 (3.98, 20.65)

Note: Posterior medians with 90-percent credible intervals in parentheses. Entries are the percentage change in the discounted five-year response of each tax base to a one-percentage-point increase in the top-1-minus-bottom-99 average marginal tax-rate gap, holding initial personal tax payments fixed, evaluated draw by draw at each year's observed rates and exposure weights and averaged across accounting years. The personal block has two groups, identified by two narrative instruments over nine reform years, conditioning on the annual corporate and tariff narrative measures. The second row reallocates S-corporation income between the personal and corporate bases before estimation.


Figure E.1: Progressivity Semi-Elasticity Leave-One-Out



Note: This figure shows the semi-elasticity of tax bases to the progressivity gradient leaving out one reform at a time.

F Identification Diagnostics

Figure F.1: Future-Shock Placebos



Note: The top row reports quarterly own-base placebos for the personal, corporate, and tariff directions. The bottom row reports annual placebos for the retained personal progressivity direction. Lines are point estimates. Shaded regions are 90-percent blockwise simultaneous max- t bands from 1,000 Gaussian wild-bootstrap draws using common multipliers across the displayed coefficients, conditional on the observed joint first stage.


Table F.1: First-Stage Strength of Tax Instruments

	Effective F
<i>Panel A. Quarterly broad-tax structural shocks</i>	
Personal income tax	120.5 [37.5, 406.9]
Corporate income tax	9.3 [5.8, 15.8]
Tariff	11.0 [8.0, 16.7]
<i>Panel B. Annual personal-tax rate innovations</i>	
Bottom 99 percent	12.1 [8.8, 16.9]
Top 1 percent	131.1 [88.2, 155.1]

Note: Posterior medians, with 5th–95th posterior percentiles in brackets. Panel A reports the Montiel-Olea–Pflueger effective F for each recovered Mertens–Ravn structural shock, using one lag of the six quarterly states and Bartlett HAC(20). Panel B reports two-proxy effective F statistics for the annual BLP first-stage personal-rate innovations, whose rate coordinates are the bottom-99 and top-1 average marginal tax rates. It conditions the innovations and the two raw MMO measures on the annual corporate and tariff narrative measures and uses Bartlett HAC(5).


G Alignment Metric Robustness

Figure G.1: Posterior Probability that Alignment Exceeds Its 1948 Level




Note: At each date, the figure reports the share of posterior draws for which full alignment exceeds its value in 1948. The comparison is made draw by draw, retaining the joint uncertainty in the broad and progressivity components of the revenue gradient.

Figure G.2: Alignment under Alternative Dynamic Estimators




Note: The figure compares the baseline Bayesian local projection estimator with conventional, smooth, and bias-corrected local projections, a Bayesian VAR, and a proxy SVAR. The shaded region around the baseline is its 90-percent posterior credible interval; other series are posterior medians. Panel A reports $\cos \theta$ and Panel B reports $\sin \theta$.

Figure G.3: Alignment under Direct IV Estimation




Note: The figure compares the baseline Bayesian local projection estimates with direct LP-IV, bias-corrected LP-IV, smooth LP-IV, and proxy-SVAR estimates. These alternatives address the possibility that the tax shocks are not invertible from the reduced-form innovations. The shaded region is the baseline 90-percent posterior credible interval; alternative series are posterior medians.

Figure G.4: Alignment under a Progressivity-Orthogonal Personal Instrument




Note: The figure plots posterior medians and 90-percent credible intervals of $\cos \theta$ under the baseline quarterly instrument and under a personal narrative purged of its progressivity loading. The alternative instrument is the residual from projecting the quarterly personal narrative on the top-minus-lower bracket gap of the annual narrative record, without an intercept, which removes 36 percent of its variation and lowers the first-stage effective F for the personal coordinate from 120.5 to 50.4. Nine of the seventeen quarterly narrative events have an annual bracket record and receive their measured gap; the remaining eight have none and are assigned a loading of zero. The orthogonalized path lies above the baseline in all 76 years, with the paired posterior difference positive with probability at least 0.95 in every year.

Figure G.5: Alignment under Alternative Annual Specifications




Note: The figure varies the specification of the annual progressivity block by adding aggregate controls, including the deterministic terms used by Mertens and Montiel Olea (2018), and adding a second annual lag. The baseline shaded region is its 90-percent posterior credible interval; alternative series are posterior medians.

Figure G.6: Alignment under Matched Present-Value Horizons




Note: The annual and quarterly response windows are varied together, so that each series uses the same present-value horizon in both empirical blocks. The baseline uses five annual observations and twenty quarterly observations. The shaded region is the baseline 90-percent posterior credible interval; alternative series are posterior medians.

Figure G.7: Alignment under Alternative Treatment of S-Corporation Income



Note: The alternative series reallocates S-corporation income from the measured corporate base to the personal base before constructing the revenue gradient. The comparison assesses whether the baseline treatment of pass-through income affects the alignment calculation. The shaded region is the baseline 90-percent posterior credible interval.

Figure G.8: Alignment across Alternative Income-Group Partitions



Note: The figure compares the preferred native two-group representation of the personal tax schedule with a constructed two-group representation and with three- and four-group partitions. Each specification defines progressivity as a revenue-neutral redistribution of initial personal tax payments within its respective partition. The shaded region is the preferred specification's 90-percent posterior credible interval.

Figure G.9: Alignment under Alternative Social Welfare Weights



Note: Each panel fixes the social welfare weight on the top 1 percent relative to the lower 99 percent. Within each panel, the lines vary the corporate-income weight relative to the personal-income weight. The figure reports posterior medians; intervals are omitted to keep the full welfare-weight grid legible.

H Fiscal Constraints over the Observed Horizon

The baseline alignment calculation focuses on whether a revenue-neutral change in the composition of taxes can improve welfare. It therefore restricts attention to reforms which satisfy $\mathbf{r}_t^\top d\boldsymbol{\tau}_t = 0$, where $\mathbf{r}_t$ is the present-value revenue gradient. This is the natural restriction when the government can accommodate a tax reform through its remaining fiscal choices. However, many reforms are not revenue-neutral, and the reforms studied empirically may have effects on government spending and debt, which contradicts the intention of focusing solely on revenue-neutral reforms. Similarly, the reforms may be joint insofar as a tax cut today may lead to a tax increase in the future. This section studies a narrower slice of reforms by asking how much scope for welfare-improving reform remains if fiscal outcomes like debt and government spending must also remain fixed.

Toward that end, let $\mathbf{c}_{j,t}$ denote the response of fiscal outcome j to a local change in the tax vector, expressed in the same policy metric as the welfare gradient $\mathbf{w}_t$. Collect the relevant constraints in the matrix

$$\mathbf{C}_t = \begin{pmatrix} \mathbf{c}_{1,t} & \cdots & \mathbf{c}_{J,t} \end{pmatrix}.$$

The feasible reform space is

$$\mathcal{F}_t(C_t) = \{d\tau_t : C_t^\top d\tau_t = 0\}. \quad (\text{A.52})$$

Let $Q_{C,t}$ be a matrix whose columns form an orthonormal basis for the column space of C_t . The projections onto the constraint space and its orthogonal complement are

$$P_{C,t} = Q_{C,t}Q_{C,t}^\top, \quad P_{\mathcal{F},t} = I - P_{C,t}.$$

The best constrained reform of policy-metric size α is therefore

$$d\tau_{C,t}^* = -\alpha \frac{P_{\mathcal{F},t} \mathbf{w}_t}{\|P_{\mathcal{F},t} \mathbf{w}_t\|}, \quad (\text{A.53})$$

provided that $P_{\mathcal{F},t} \mathbf{w}_t \neq 0$. Its first-order welfare gain is

$$-\mathbf{w}_t^\top d\tau_{C,t}^* = \alpha \|P_{\mathcal{F},t} \mathbf{w}_t\|.$$

The constrained analogue of the sine statistic is thus

$$\sin \theta_{C,t} = \frac{\|P_{\mathcal{F},t} \mathbf{w}_t\|}{\|\mathbf{w}_t\|}. \quad (\text{A.54})$$

It measures the share of the unconstrained local welfare gain that remains feasible after imposing the fiscal restrictions in equation (A.52). The complementary quantity

$$\cos \theta_{C,t} = \frac{\|P_{C,t} \mathbf{w}_t\|}{\|\mathbf{w}_t\|} \quad (\text{A.55})$$

measures the share of the welfare gradient spanned by the constraints. By construction,

$$\cos^2 \theta_{C,t} + \sin^2 \theta_{C,t} = 1.$$

Unlike the baseline cosine, equation (A.55) is nonnegative. If tax revenue is the only constraint, then it yields the same absolute value as the baseline alignment metric. But since it is constrained to be nonnegative, it does not preserve the baseline alignment statistic's information about whether the welfare and revenue gradients even point in the same direction. However, since the welfare metric $\sin \theta$ is constrained to be positive anyway, it does nest the baseline for that metric. Importantly, adding constraints can only shrink the feasible policy space and thereby weakly reduce the constrained sine metric. Moreover, if the constraint matrix has a rank equal to the number of available tax instruments, then then it is not possible to constraint a local reform to satisfy every restriction.

H.1 Empirical Implementation

To empirically implement this, I use the expanded version of the system with macro controls and consider three distinct fiscal outcomes: tax revenue (the baseline), government spending, and debt. Let cumulative government spending over horizon H be denoted as

$$\mathbf{g}_{t,H} = \nabla_{\tau_t} \left(\sum_{h=0}^H \beta^h G_{t+h} \right), \quad (\text{A.56})$$

where G_t denotes government spending. Mirroring that construction, let the debt gradient be

$$\mathbf{b}_{t,H} = \nabla_{\mathbf{r}_t} \mathcal{B}_{t+H}, \quad (\text{A.57})$$

The resulting constraint sets are

$$C_{R,t} = \begin{pmatrix} \mathbf{r}_t \end{pmatrix}, \quad C_{RG,t} = \begin{pmatrix} \mathbf{r}_t & \mathbf{g}_{t,H} \end{pmatrix}, \quad C_{RB,t} = \begin{pmatrix} \mathbf{r}_t & \mathbf{b}_{t,H} \end{pmatrix},$$

and


$$C_{RGB,t} = \begin{pmatrix} \mathbf{r}_t & \mathbf{g}_{t,H} & \mathbf{b}_{t,H} \end{pmatrix}.$$

Since each restriction is required to equal zero, rescaling any column of the constraint set C_t leaves the feasible space unchanged. Consequently, I construct the projectors using a QR decomposition.

I evaluate these constraints in only the expanded reform space, which includes progressivity. In the three-instrument space, imposing tax revenue, government spending, and debt restrictions generically exhausts all three dimensions and there is no feasible reform. Consequently, I only show the results for the expanded space.

Figure H.1 compares the baseline the more restrictive fiscal problems in the complete reform space. Holding debt or government spending fixed generally shifts up the level of the cosine, but does not change the time trend toward improvement. Holding all three fixed shows a modest reduction in the gains remaining from reform until 1990, but is considerably flatter than the baseline.


Figure H.1: Alignment under Additional Fiscal Constraints



Note: The figure reports the constrained cosine and sine defined in equations (A.55) and (A.54) in the complete reform space. The revenue series imposes only tax-revenue neutrality. The remaining series add restrictions that government spending, terminal debt, or both remain unchanged over the stated horizon. The constrained cosine is an unsigned projection statistic and is not the signed baseline Ramsey cosine.


I Additional Robustness Figures

Figure I.1: Alignment under Alternative Assumptions about Tariff Incidence




Note: The figure varies the fraction of the marginal tariff burden borne by domestic groups included in welfare. The three series correspond to zero, one-half, and full domestic incidence. The shaded region is the full-domestic-incidence baseline's 90-percent posterior credible interval.

Figure I.2: Alignment under an Adjustment for Foreign Corporate Ownership




Note: The alternative series adjusts the welfare gradient for the share of corporate income accruing to foreign owners. The comparison isolates this incidence adjustment while holding the estimated revenue responses fixed. The shaded region is the equal-incidence baseline's 90-percent posterior credible interval.

Figure I.3: Alignment under Alternative Welfare Weights on Tariff Burdens




Note: The figure varies the social welfare weight on the tariff burden relative to the equal-weight benchmark, holding the personal and corporate weights fixed. The shaded region is the equal-weight baseline's 90-percent posterior credible interval; alternative series are posterior medians.

Figure I.4: Sensitivity to the Annual Extension and Recovered TCJA Shock




Note: The figure compares the preferred annual progressivity design with one restricted to the original Mertens–Montiel Olea sample through 2012 and with one that omits the recovered TCJA narrative measure. The shaded region is the preferred specification's 90-percent posterior credible interval.

Figure I.5: Welfare-Weight Threshold for the Progressivity Component




Note: Panel A reports the top-1-to-lower-99 welfare weight at which the locally optimal reform has a zero progressivity component, where such a threshold exists. The horizontal line marks equal welfare weights. Panel B reports the share of posterior draws with a finite threshold and the share for which the reform is not progressive under equal weights. The shaded regions in Panel A are 68-percent and 90-percent posterior credible intervals.

Figure I.6: Sensitivity to the Measurement of the TCJA Progressivity Shock



Note: The figure compares the preferred Statistics of Income construction of the 2018 TCJA measure with alternative SOI and Current Population Survey constructions, as well as a specification that omits the recovered TCJA measure. The shaded region is the preferred specification's 90-percent posterior credible interval.

Figure I.7: Revenue Channels of Personal Progressivity



Note: The figure decomposes r_d into personal, corporate, and tariff base contributions. Lines are posterior medians.